

Bots

Profound Sheets Sitemap Importer

The Profound Sheets Sitemap Importer is a user-initiated HTTP agent that fetches a website's sitemap for its URLs to be imported into a Profound Sheet. This bot only runs when an authenticated Profound user clicks **+ Create new → Sitemap** inside the Sheets section of the Profound platform and supplies a site URL.

Summary

User-Agent string	ProfoundSheetsSitemapImporter/1.0
Trigger	User-initiated
Schedule	None
Resources fetched	sitemap.xml , sitemap-index variants, robots.txt , and any sitemap files those documents reference
Network origin	AWS, no fixed IP allow-list available

What it does

When initiated with a website origin URL, the Sitemap Importer makes a small, bounded set of HTTP GET requests to locate and read a site's sitemap. It tries the following URLs in sequence:

- https://{origin}/sitemap.xml
- https://{origin}/sitemap_index.xml
- https://{origin}/sitemap-index.xml
- https://{origin}/robots.txt (parsed only for Sitemap: directives, which are added to the candidate list)
- any child sitemaps referenced by a sitemap index

If you enter a URL ending in .xml , the importer also tries that URL directly.

The first response that parses as a valid XML sitemap is used. The URLs it contains are returned to you for import. The importer does not store, cache, redistribute, or republish sitemap content.

What it doesn't do

- Fetch HTML pages, images, scripts, stylesheets, or any other assets
- Submit forms, follow login flows, or access authenticated areas
- Run JavaScript or use a headless browser
- Crawl recursively beyond sitemap-index to child-sitemap links

- **Persist** a long-running index of your site

Request behavior

Each import action triggers a small burst of requests scoped to a single origin.

Per request timeout	10 seconds
Redirects	Followed manually, up to 5 hops
Concurrency	Sequential (not in parallel)
Caching	None. Each user invocation issues fresh requests

Headers sent

```
User-Agent: ProfoundSheetsSitemapImporter/1.0
Accept: application/xml,text/xml,application/xhtml+xml,text/plain;q=0.9,
```

Volume and frequency

Traffic from the Importer is driven entirely by customer activity. There's no continuous crawl. A typical interaction involves a handful of requests, after which the Importer doesn't contact an origin again until another import action targets the same site.

robots.txt handling

The importer reads `robots.txt` to discover sitemap URLs declared via `Sitemap:` directives. It doesn't evaluate `Disallow` rules against sitemap files, because the sitemap protocol treats those files as publicly discoverable.

To prevent the importer from accessing your site, see [How to block](#) section.

How to identify the bot

Use the exact User-Agent header:

```
User-Agent: ProfoundSheetsSitemapImporter/1.0
```

💡 All Profound bots use User-Agent strings that start with `Profound`, so a prefix match on `Profound` is a good way to identify all Profound traffic, current and future.

How to block

Block the importer using any of the following methods:

- A WAF or CDN rule matching User-Agent equal to `ProfoundSheetsSitemapImporter/1.0`

- A WAF or CDN rule matching User-Agent containing the prefix `Profound`
- Returning a 4xx response to the bot's User-Agent on `/sitemap.xml` , `/sitemap_index.xml` , `/sitemap-index.xml` , and `/robots.txt`

Reporting abuse

If you observe behavior that doesn't match this documentation, such as requests for resources other than sitemap and `robots.txt` files, request rates that look like a crawl, or activity from an unrecognized `Profound*` User-Agent, report it to security@tryprofound.com with example log lines.

Profound treats these reports as security issues and responds to all of them.

Was this page helpful?

 Yes

 No

[< ProfoundBot](#)

[Enterprise Single Sign-On \(SSO\) >](#)