

INDEX

1. Abstract

2. Introduction

- 2.1 Problem Statement
- 2.2 Research Objective
- 2.3 Research Question
- 2.4 Significance of the Study

3. Literature Review

- 3.1 Evolution of Search
- 3.2 Traditional SEO
- 3.3 Generative Engine Optimization
- 3.4 Answer Engine Optimization
- 3.5 Retrieval-Augmented Generation
- 3.6 LLM Visibility Challenges
- 3.7 Research Gap

4. Conceptual Framework

4.1 Defining LLM SEO Files

- 4.1.1 Core Components of LLM SEO Files
- 4.1.2 Relationship to Prior Paradigms
- 4.1.3 The Emergence of LLM-First Content Strategies

4.2 Understanding `llms.txt`: A Proposed Standard for LLM-Oriented Optimization

- 4.2.1 `llms.txt` vs. `robots.txt` and `sitemap.xml`: Key Distinctions
- 4.2.2 Perceived Benefits for AI Understanding and Content Prioritization

4.3 The Great Debate: Google's Skepticism vs. Proponent Perspectives

- 4.3.1 Google's Stance: Skepticism, Meta Keywords Comparison, and Operational Concerns
- 4.3.2 The Proponents' View: Utility in Multi-Model Ecosystems and Endorsement by LLMs
- 4.3.3 Analysis of Adoption Challenges and Current Status

4.4 Introducing the LLM SEO Files Standard

- 4.4.1 Rationale and Design Principles
- 4.4.2 Deep Dive into the 6 Core Files

- **4.4.2.1 `llms.txt`: The Foundational AI Guide**

- **4.4.2.2 page-name.yaml: A Human-Readable Semantic Map for AI Clarity**
- **4.4.2.3 page-name.md: Human and Machine Readability**
- **4.4.2.4 Schema JSON: Enhancing Structured Data for AI**
- **4.4.2.5 evidence-of-facts.json: Structuring Epistemic Clarity for LLMs**
- **4.4.2.6 page-name.json: Open Graph Metadata**

4.5 Strategic Implications for Web Publishers and Digital Marketers

4.5.1 Optimizing Content for AI Readability and Citation

4.5.2 Integrating LLM SEO Files into Existing Digital Strategies

5. Methodology

- **5.1 Research Design**
 - 5.1.2 Research Scope and Focus
- **5.2 Framework Implementation**
- **5.3 Data Collection**
 - 5.3.1 Prompt Set Creation
 - 5.3.2 Content Versioning and Testing Groups
 - 5.3.3 Controlled Deployment Conditions
- **5.4 Evaluation Criteria**
- **5.5 Experiment Phases**
- **5.6 Tools and Platforms Used**
- **5.7 Limitations and Bias Controls**
- **5.8 Methodology Flow Diagram**

6. Results and Analysis

6.1 Overview of the Experiment

6.2 Experimental Protocol

6.3 Evaluation Metrics

6.4 Quantitative Findings

- **6.4.1 Prompt-Based Visibility Trends**

- 6.4.2 Intent-Based Visibility Trends
- 6.4.3 Platform-Level Visibility Trends: Brand Mentions and Citations

7. Discussion

8. Limitations and Future Research Directions

- 8.1 Limited Dataset Scope and Content Domains
- 8.2 Platform Variability and Model Non-Determinism
- 8.3 Lack of Transparency in AI Ranking and Citation Logic
- 8.4 Controlled Experimental Design Versus Real-World Web Dynamics
- 8.5 Adoption Barriers and Technical Trade-Offs
- 8.6 Temporal Limitation and Short-Term Snapshot

9. Future Research Directions

- 9.1 Expanding Across Industries and Languages
- 9.2 Tracking Changes Over Time
- 9.3 Partnering with Model Developers for Transparency
- 9.4 Understanding Human Impact and Usability

10. Conclusion

11. References

Industry analysis, Conceptual framework, Forward-looking recommendations

Title: *A Strategic Outlook on LLM SEO: Using File-Format Logic to Guide AI-Optimized Content Design*

Author: Avinash Tripathi, Co-founder, AI Monitor

1. Abstract

This paper introduces **LLM SEO Files**, a novel conceptual framework designed to optimize digital content for visibility, retrievability, and citation across AI-powered search platforms, including ChatGPT, Google AI Overview and Perplexity. While traditional SEO and emerging paradigms such as Answer Engine Optimization (AEO) or Generative Engine Optimization (GEO) have adapted to shifting search behaviors, a critical gap remains: the absence of a systematic, file-based methodology aligned with how large language models (LLMs) parse, interpret, and attribute information.

To address this, we propose a six-layer symbolic framework comprising: TXT files (llms.txt), YAML configuration (index.yaml), Markdown summaries (page-name.md), Schema.org markup (schema.JSON-LD), and a newly introduced json file for the Evidence of Facts (evidence-of-facts.json). And open graph (open-graph.json) The fifth component explicitly separates factual claims from subjective commentary, supporting the increasing need for grounded, low-hallucination responses among LLM's.

We conducted an experiment using over 100 AI-search prompts across three leading LLM platforms: ChatGPT, Perplexity, and Google AI Overview. To account for regional differences in AI outputs, these prompts were run from five distinct geographic locations: India, the United States, Australia, the United Kingdom, and Germany. In Phase 1, we checked the prompts using non-optimized content. Phase 2 introduced our five-layer files, which led to a remarkable 64% increase in brand mentions and domain citations. This uplift clearly highlights the impact of structured optimization in enhancing AI interpretability and response precision.

These findings offer a reproducible foundation for future research and standard-setting in AI-native content optimization. We position LLM SEO Files as both a foundational methodology for generative search and a potential sub-discipline within the broader ecosystems of SEO, AIO (Artificial Intelligence Optimization), AEO, and GEO, laying the groundwork for scalable, cross-platform visibility practices.

2. Introduction

Over the past decade, the way we search for information online has changed dramatically. Instead of typing in keywords on traditional search engines, more and more people are now turning to natural language prompts directed at Large Language Models (LLMs) like ChatGPT, Gemini (formerly Bard), and Perplexity. These AI platforms do more than just pull up links from websites; they generate answers, simplify complex topics, and offer suggestions by pulling from vast amounts of unstructured and semi-structured data. This shift has fundamentally altered how digital content becomes visible, trusted, and cited in AI-generated responses, moving far beyond the traditional web search systems that shaped the world of SEO.

For years, Search Engine Optimization (SEO) has depended on structured HTML, metadata, backlinks, and sitemap hierarchies. But the game is changing fast, with the rise of generative AI. To stay ahead, strategies like Answer Engine Optimization (AEO) have begun to take shape, helping content get picked up in AI-generated answers and AI Overviews. Similarly, Answer

Engine Optimization (AEO) focuses on making content more compatible with the prompt-based reasoning of modern LLMs. However, AEO still relies heavily on human-readable, web-native formats, and they fall short when it comes to directly engaging with how LLMs work.

That's because LLMs don't "browse" the internet like we do. Instead, they use things like embeddings, vector databases, knowledge graphs, and increasingly, structured symbolic inputs that help improve grounding and reduce hallucinations. While companies like Google have highlighted the importance of structured data and schema for making content more understandable to AI, there is still no universal file-based standard built specifically for how LLMs evaluate, interpret, and cite information.

2.1 Problem Statement

Even as generative AI platforms continue to grow, most SEO strategies are still built around traditional, web-focused practices. But without a clear, standardized, and machine-readable method designed for how LLMs process content, visibility remains inconsistent, citations are often missed, and AI-generated responses can easily misrepresent information. For content creators and researchers, there's still no reliable, transparent way to help AI understand and properly credit their work.

2.2 Research Objective

This paper introduces LLM SEO Files, a novel, structured framework composed of six symbolic file types designed to optimize how content is read, parsed, and cited by LLMs. These layers llms.txt, YAML configuration, Markdown summaries, Schema.org markup, "Evidence of facts" JSON file, "open graph" JSON file go beyond format diversity; they function as semantic cues that enhance the machine-readability, retrievability, and attribution quality of digital content. The objective is to establish a reproducible, platform-agnostic methodology that aligns with the operational logic of generative AI systems and conversational agents.

To understand how well the framework works in real scenarios, we tested it with more than 100 prompts across popular AI platforms. We wanted to understand how using the combinations of LLM SEO Files affected whether content gets picked up or cited by AI systems. The goal wasn't just to talk about the idea in theory, but to offer a practical, usable method that helps content become more visible and easier for large language models to understand.

2.3 Research Questions

To evaluate the effectiveness of the proposed LLM SEO Files framework in enhancing AI-driven content visibility and attribution, this study is guided by the following research questions:

- **RQ1:** To what extent does the implementation of LLM SEO Files increase the likelihood of the brand being mentioned in AI-generated responses across leading large language model (LLM) platforms?

- **RQ2:** How do different AI platforms respond to structured content in terms of brand mention frequency and citation behavior across varied user intents?

These questions aim to systematically assess whether structured, symbolic content layers improve machine interpretability, factual grounding, and content attribution in LLM-powered environments.

2.4 Significance of the Study

This research contributes in two main ways: one conceptual and the other empirical. Conceptually, it builds on the idea of LLM SEO first introduced by Avinash (2024) by putting forward a more practical, file-based framework that mirrors how large language models tend to process and understand content. On the empirical side, we found that using the complete LLM SEO Files setup led to a sharp increase of about 64% in how often content was cited or surfaced by AI-powered search engines. However, when the non-optimized framework was applied, performance declined significantly. This contrast highlights the importance of adopting a layered and intentional approach when designing content for AI-driven discovery

What makes this work is its practical value. It offers a repeatable approach that researchers, SEO professionals, and AI content strategists can apply across different platforms. More broadly, it opens up new ground for exploring how structured digital content interacts with AI systems and positions LLM SEO files as a possible new area of study at the crossroads of generative AI, publishing, and computational linguistics.

3. Literature Review

3.1 Evolution of Search

Key Literature Reviewed:

- Karanxha, G. (2025). *Evaluating Transparency in Artificial Intelligence Systems: Adapting the Z-Inspection® Framework for the MANOLO Project*
- OpenAI, Perplexity, and Anthropic documentation (2023–2025)

Summary and Analysis:

Search has changed. What started with simple keyword matching and systems like PageRank built on backlinks and indexing has gradually evolved toward AI-powered tools that understand the intent of a user's query. When semantic search entered the picture through features like Google's Knowledge Graph, it brought better context and connections between entities. But

even then, the search was still based on matching queries to indexed pages.¹

Now, we're in the middle of a much bigger shift. Large Language Models (LLMs) like ChatGPT, Gemini, and Perplexity aren't just retrieving websites; they're generating answers.² These responses are created by pulling from a mix of web content, structured data, and prior training material of the LLM. This move away from query-to-URL results toward synthesized, token-based answers marks a fundamental change in how information is delivered.

Karanhya (2025), in her adaptation of the Z-Inspection® Framework, highlights one of the biggest challenges here: transparency. These models don't clearly show where the information comes from, and there's no straightforward way to understand how or why certain content gets surfaced. That makes it hard for content creators to know how to structure their work to be discoverable. On top of that, the models are trained on data that can contain biases, which adds further complications.³

We're no longer just talking about getting ranked in a list of links. We've entered the "answer generation" era, where the goal is to be part of the answer itself. That creates both a challenge and an opportunity: how do we make content that these models would actually use, cite, and trust? This is where frameworks like LLM SEO Files come in. They're designed to help bridge the gap between human-created knowledge and the way LLMs process and surface information so that the content remains relevant in this new AI-driven world.

3.2 Traditional SEO

Key Literature Reviewed:

- Fishkin, R. (2015). *The Beginner's Guide to SEO* (Moz).
- Enge, E. et al. (2020). *The Art of SEO: Mastering Search Engine Optimization*.
- Cutts, M. (2013). Google Webmaster Blog (2010–2013).

Summary and Analysis:

Search engine optimization, as it's been practiced for the past two decades, has mostly centered on working with the logic of how search engines crawl and rank pages. Practitioners focused on making content easy to find and easy to interpret using keywords in expected places, earning backlinks from authoritative sites, and making sure everything is loaded fast and follows a clean technical structure (Fishkin, 2015; Enge et al., 2020).⁴ The entire strategy was built around the assumption that search engines are essentially indexers: they scan, they store, they sort. Google's early guidance, highlighted by figures like Matt Cutts in 2013, emphasized that key ranking factors included how easily a page could be crawled, how recently it was updated, and the strength of its inbound links.⁵

Even as algorithms evolved and search engines got more sophisticated, the fundamentals did not change much. But that landscape is now being disrupted. AI tools powered by large language models don't "crawl" the same way. They don't need to.⁶ Instead of looking at HTML

structure or counting keywords, they absorb language and meaning, and generate answers based on patterns they've learned. There's no traditional index driving the result. That changes everything. Content that's been optimized according to the old rules, no matter how well, can easily go unnoticed in these new systems. The gap is growing, and many web creators haven't yet caught up.

3.3 Generative Engine Optimization (GEO)

Key Literature Reviewed:

- Tripathi, A. (2024). *Generative Engine Optimization: A New Frontier for AI Visibility*.
- OpenAI (2023). *How ChatGPT Retrieves and Cites Information*.
- Google (2023). *AI Overview Technical Whitepaper*.
- GenBE (2025). *GEO Strategy Playbook: Structuring Content for AI Engines*.

Summary and Analysis:

Generative Engine Optimization, or GEO, is a term that has recently started to gain attention. It wasn't part of the mainstream SEO vocabulary until around 2023 when researchers like Pranjal Aggarwal and Vishvak Murahari, along with teams from IIT Delhi and Princeton, introduced it into the conversation.⁷ Tripathi (2024) later picked it up and helped promote it further. He presented it as a clear shift from the old, index-based ways of thinking about search. Instead of optimizing for how search engines crawl and rank, GEO focuses on how AI models, like ChatGPT, Claude, and Google's AI Overviews, actually read and respond to content.⁸

This shift is more than just technical. It's conceptual. Large language models don't care much about keyword density or link authority in the traditional sense. They work by interpreting patterns in language context, meaning, and relevance within a prompt. Google's 2023 whitepaper on AI Overviews talks about this quite directly: it highlights things like content freshness and structural clarity, but not in the way SEO professionals are used to.¹⁰ There's no crawler following links anymore. The AI needs the content to be understandable and trustworthy on its terms.

OpenAI made a similar point. In 2023, they noted that citations in ChatGPT aren't fixed; they're generated probabilistically. So if your content isn't well-structured, or if the model can't "ground" it confidently, it might not show up at all.⁹ That's where things like FAQs, definition blocks, and semantic markup come in. GenBE's 2025 playbook builds on this idea and gets into the weeds: it recommends using schema.org variants, structured JSON, and breaking content into modular pieces so generative engines can grab and reuse what they need without friction.¹¹

What GEO captures is a shift in how people search and how machines respond. Users aren't combing through search results anymore. They're asking questions and getting direct answers. That means the old SEO signals, like ranking and crawling, are no longer the gatekeepers. If an LLM cannot understand and interpret your content fully, it will not be cited. GEO is simply an attempt to keep up with that reality.

3.4 Answer Engine Optimization (AEO)

Key Literature Reviewed:

- Enge, E. (2020). *The Art of SEO: Mastering Search Engine Optimization*.
- Sullivan, D. (2022). *Google Search Central Blog*.
- van Dis, E. A. M. et al. (2023). *Responsibly Deploying Language Models*.

Answer Engine Optimization, or AEO, came into focus when search began shifting toward more direct, AI-assisted results, things like Google's Featured Snippets, People Also Ask boxes, and voice-based queries. The idea was to structure your content in a way that it answers questions clearly and concisely. So, people started using bullet points, FAQs, and schema markup to improve their chances of being pulled into those instant answer spots (Enge, 2020).¹² It was not a radical departure from SEO, but more of a tactical adaptation, one that responded to users wanting faster, clearer information without having to click through multiple links.

But there's a limit to how far AEO can go. It still plays by the rules of index-based search. And as search interfaces evolve, especially with the rise of LLMs, those rules are becoming less relevant. Google might still extract short answers from well-structured pages (Sullivan, 2022), but once that content ends up in a generative system, things get murky. You can't always tell where the information came from, or whether it was cited properly, if at all.¹²

This is a bigger issue than it first seems. As Liao and Vaughan (2023) point out, today's large language models often lack mechanisms to retain or communicate source details and metadata in their outputs.¹³ So even if AEO helps your content rank well in the short term, it doesn't do much to ensure traceability or accurate attribution in the world of generative AI. What we're seeing is a growing mismatch: AEO is optimized for visibility in search engines that still rely on links and indexes, but it doesn't address how AI tools interpret or credit information behind the scenes.¹³

3.5 Retrieval-Augmented Generation (RAG)

Key Literature Reviewed:

- Lewis, P. et al. (2020). *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*.
- LangChain (2024). *RAG Use Cases and Implementation Patterns*.
- Karanxha, G. (2025). *Evaluating Transparency in AI Systems with the Z-Inspection® Framework*.

Summary and Analysis:

Retrieval-Augmented Generation (RAG) isn't exactly a new idea, but it's become a lot more relevant lately. The basic concept is simple: instead of expecting a language model to know everything it would need in the future, RAG lets it pull in extra information as it's responding.¹⁴ This means it can look something up, grab a paragraph, a snippet, and use that instantly to make the output smarter or more accurate. Lewis and his team floated the idea back in 2020, but only in the last few years has it started to catch on in practice. It turns out this approach works especially well when the task depends on facts, current information, or anything where getting it wrong isn't an option.¹⁴

LlamaIndex (2023) and LangChain (2024) have popularized RAG pipelines within production environments, offering developer-friendly tooling to connect unstructured content (e.g., PDFs, blogs, databases) with LLM prompts. These implementations emphasize vector-based retrieval (via FAISS or Pinecone), chunking strategies (e.g., sentence windowing or semantic tiling), and metadata-aware filters to enhance the precision of retrieved context.¹⁵

Karanxha (2025) brings up something that often gets overlooked in conversations about RAG systems: the ethics of how information gets pulled and used. Sure, retrieval helps reduce hallucinations. Models are less likely to make things up if they're pulling from real sources, but there's still a lot we don't see². Specifically, how those sources are chosen or ranked is rarely transparent. Her work, especially the use of the Z-Inspection® Framework in RAG contexts, pushes for clearer standards: things like showing where information came from, how it was selected, and what kind of audit trail exists behind an AI's output.²

This matters a lot when you look at it through the lens of LLM SEO. If retrieval is part of the generation process, then visibility isn't just about what shows up on the surface of Google or Bing anymore; it's about what's available to the model *at the exact moment it's generating the answer*. That's a big shift. It means structured files like `llms.txt`, vectorized FAQs, or schema-rich markdowns need to be built with retrievability in mind. In other words, content has to be packaged not just for human readers or search bots, but for AI systems scanning vector databases during inference. We're not just optimizing pages anymore; we're optimizing fragments, formats, and storage.

3.6 LLM Visibility Challenges

Key Literature Reviewed:

- Mendoza Juan, Y. (2024). *Multi-Agent LLM Systems for Enhanced Human-Machine Interaction*
- OpenAI (2024). *How ChatGPT Retrieves and Cites Information*
- IBM Research (2023). *AI Transparency and Explainability in Language Models*
- Gillespie, C., Lovelace, R. (2016). *Efficient R Programming* (for tokenization and content formatting insights)

Summary and Analysis:

Large language models (LLMs) have quickly taken over as the dominant interface for information retrieval, but figuring out how visibility works inside these systems is a lot messier than with traditional search engines. Google, for all its complexity, still shows a ranked list of results. You know what's been indexed and why something appears first. With LLMs, this is not the case. These models generate responses one token at a time, based on probabilities, not page ranks. Once your content appears in an AI-generated response, tracing its origin becomes incredibly difficult. Juan Mendoza (2024) delves into this issue in the context of multi-agent LLM systems, where attribution becomes even more opaque and complex. Once you have several agents talking to each other, plus prompts layered on top of prompts, it becomes nearly impossible to tell where any specific piece of information came from, let alone who to credit.¹⁶

As for OpenAI's explanations? They're vague at best. The official documentation (2024) says ChatGPT answers are built from a mix of fine-tuned training data, model knowledge, and (for Pro users) real-time web browsing. But here's the thing: they rarely cite where the answer came from unless you explicitly ask.¹⁷ Even then, it's inconsistent. So if your content helped shape a response? You might never know. IBM Research (2023) has raised concerns about this, especially in high-stakes areas like healthcare or enterprise systems, where knowing where an answer came from really matters.¹⁸

If you're trying to optimize for visibility in LLMs, you're not dealing with just one problem; you're up against several. To start with, there's the grounding issue: your content might be used to create a response, but you might never be mentioned as the source. Then there's the lack of indexing; things like sitemaps and backlinks don't play the same role here. Also, when content runs too long or isn't broken up clearly, there's a good chance some of it just gets ignored. If a page lacks basic structure like FAQs, bullet points, tables, or any helpful formatting the model might not make much sense of it in the first place

That's where a new kind of optimization needs to happen. LLM SEO Files are one approach: these are structured, machine-readable content formats that make it easier for AI models to pull in and ground responses efficiently. Gillespie and Lovelace (2016) touched on this years ago in a different context, showing how careful formatting and chunk sizing improve efficiency. The same idea applies here. It's not just about creating valuable content; it's about packaging content in a way that LLMs can read, remember, and reuse that information.¹⁹

3.7 Research Gap

Despite major advances in SEO, AEO, and even the newer GEO methods, most of these strategies are still built around how traditional search engines work, ranking, indexing, and deterministic answers. They're simply not designed for the messy, unpredictable behavior of large language models (LLMs), especially when it comes to visibility, grounding, or attribution.

The problem is, existing content strategies weren't made for systems that generate responses probabilistically or pull from embeddings instead of ranked links. No widely adopted method tells LLMs how to interpret content across formats, platforms, and structures. Mendoza (2024), in his work on multi-agent LLM pipelines, shows how symbolic reasoning and modular prompts help make these systems more reliable, but almost no one is applying those ideas in real-world content.¹⁷ We've had the technical foundations for a while. Evans (2003) laid out how domain-driven design brings semantic structure to complex systems.²⁰ Gillespie and Lovelace (2016) emphasized clean formatting and reproducibility, two things LLMs still struggle with when

reading raw, rambling content. Yet somehow, those lessons haven't made it into how we format online content today ²¹

What's missing is a practical, standardized way to prepare content for LLMs, something machine-readable, traceable, and designed to work across different models and retrieval systems. File types like .txt, .yaml, or .json are rarely used for this purpose, even though they're exactly what models trained on structured data would understand best. That's where the LLM SEO Files framework comes in. It offers a structured, symbolic approach to making content not just visible but retrievable, interpretable, and credible in LLM-driven environments. Instead of just writing for humans, we need to start writing our content for machine-groundable knowledge. The framework helps close the widening gap between the content we create and recognize, and what AI systems choose to cite in today's AI-driven ecosystem.

4. Conceptual Framework

4.1 Defining LLM SEO files

LLM SEO files are a collection of straightforward, well-organized documents that make it easier for large language models to find, understand, and properly credit your content. The term "LLM SEO" stands for *Large Language Model Search Engine Optimization* and each part of the acronym corresponds to a specific file, each with its own role in helping your content stand out in AI-driven search results. **L** is for `llms.txt`, a basic guide that tells AI what are the most important pages of the website and what pages not up to date and what information from the website the LLMs can ignore. The second **L** stands for a yaml file, `index.yaml`. This file includes key metadata like author, topic, and content intent. **M** denotes a Markdown file `index.md`, a file that presents information in a clear and readable format. **S** stands for `schema.json`, which adds structured data so AI can understand relationships and context. **E** stands for Evidence of Facts file (`page-name.json`) where you back up claims with verifiable facts, and separate out perspectives and viewpoints, **O** stands for Open Graph metadata (`og:title`, `og:description`, `og:image`, etc.) embedded in the webpage to enhance content previews and ensure AI (and social platforms) display content with the right summary, image, and context when it's shared or indexed. Together, these six files form the foundation of the LLM SEO approach making sure your content isn't just out there, but actually understood and properly credited by today's generative AI systems.

4.1.1 Core Components of LLM SEO files

To make content genuinely useful in the age of AI-driven discovery, we need to rethink how it's created and shared. Traditional SEO tactics like keyword placement and link building don't translate that well in environments where AI tools generate responses rather than rank links. What matters now is how easily large language models (LLMs) can locate, understand, and reuse your content. LLM SEO files aren't meant to trick or game AI systems they're designed to help them actually understand your content. The idea is built on a few core principles, beginning with how well your information is organized and presented.

a. LLM-Readable File Formats

Generative AI models don't explore the web the way traditional search engines do. Instead of crawling links, they often use ingestion pipelines and parsing tools that perform best when fed clean, well-structured content. To meet these needs, a new wave of file formats has emerged formats designed to make information easier for AI systems to process and understand.

For instance:

- **llms.txt**: Similar to *robots.txt*, this file tells language models which content or endpoints are most relevant for ingestion.
- **og.json**: A smarter take on Open Graph metadata, tailored for how AI models process context.
- **faq.md**, **data.yaml**, and **outline.md**: These lightweight formats help preserve clarity and structure without overwhelming the model with excess information.

These formats do more than just clean up your content; they significantly increase the chances of your material being accurately interpreted and understood by LLMs. This in turn results in increasing your chances of being featured in AI-generated responses, whether it's a Perplexity answer card or a ChatGPT citation. In today's landscape, that kind of visibility is quickly becoming as valuable as landing the top spot on Google search results once was.

b. Atomic Content Blocks

LLM SEO encourages breaking down information into small, standalone chunks often referred to as atomic content blocks. These might be a paragraph, a quote, a checklist item, or a data point. Each one is meant to work independently within the limited context window of a large language model.

Key features of atomic blocks include:

- **Self-contained**: They make sense on their own, without needing information from earlier sections.
- **Structured clearly**: Using headers, bullet points, and question-and-answer formats helps make the content easier to interpret.
- **Modular-friendly**: These blocks can be retrieved more effectively by chunk-based or embedding-based AI search methods.

This kind of granular organization improves how well AI systems can cite your content, ground their answers, and reuse specific snippets in responses.

c. Metadata for AI (vs. Metadata for Search Engines)

Traditional SEO relies on metadata like meta descriptions and canonical URLs to help search engines understand your pages. But LLM SEO calls for metadata that works specifically for AI models.

This includes:

- **Prompt-visible fields**, such as `og:title` and `og:description`, which often appear directly in AI-generated answers.
- **Inline semantic cues**, like `summary:` or `key insights:`, which help models grasp what matters most in your content.
- **LLM-friendly schema formats**, like `FAQPage` or `Dataset`, with concise and clear markup.

Adding this kind of metadata increases the chances that your content will be recognized, accurately summarized, and properly credited in AI outputs.

d. Content Retrievability

It's not enough for your content to just exist online it has to be accessible in a way AI systems can actually process.

That means:

- **Using canonical URLs** that don't rely on render-blocking scripts.
- **Choosing crawlable, scrapeable formats**, rather than gated PDFs or interactive JavaScript widgets.
- **Using clean semantic HTML**, such as `<article>`, `<section>`, and `<header>`, to make your structure easier to follow.

Retrievability also depends on **where and how** your content is hosted. If key content is locked behind tabs, pop-ups, or login walls, it's far less likely to be included in AI training or used as a grounding source.

e. Token Efficiency

LLMs operate within strict limits called *context windows* which might range anywhere from 8,000 to 200,000 tokens. Writing in a token-efficient way ensures that your content fits within those

limits without losing meaning.

Some techniques include:

- **Avoiding fluff** cut down on long intros and repetitive phrasing.
- **Using tables, bullets, and TL;DR summaries** to pack more meaning into fewer tokens.
- **Writing with compression in mind**: aim for one idea per sentence, keep it tight, and avoid filler.

Efficient content not only increases your chances of being included in an LLM's output it also improves how clearly your content is presented when it's quoted or paraphrased.

f. AI Grounding Signals


Grounding refers to how AI models justify where their answers come from. LLM SEO supports the use of clear grounding signals to boost transparency and increase the odds of citation.

These signals include:

- **Direct attribution**, like "Source: Relific" or inline citations.
- **Authoritative phrasing**, such as "According to a 2024 report by..."
- **Reasoned responses**, where a clear answer is followed by supporting logic a format favored by retrieval-augmented generation (RAG) systems.

When content includes these elements, it's more likely to be cited by LLMs like Perplexity or appear in Google's AI Overviews, where source transparency and credibility matter.

Diagram: Core Components of LLM SEO



- | 3. Metadata for AI | --> (OG tags, FAQs, schemas)
 - | 4. Content Retrievability | --> (HTML structure, crawlability)
 - | 5. Token Efficiency | --> (Concise, data-rich content)
 - | 6. Grounding Signals | --> (Sources, Stats, Authority)
- +-----+

By putting these components into practice, content creators can significantly boost their visibility in today’s LLM-powered information landscape. Not only do these structures help with citation and AI recognition, but they also lay a solid foundation for a more standardized, AI-aware digital strategy.

4.1.2 Relationship to Prior Paradigms

Optimization Model	Primary Goal	Platform Target	Data Structure	Output Format	Retrieval Style
SEO	Rankability	Web Index (Google)	Unstructured HTML	Ranked Web Pages	Crawler-based
AEO	Answerability	Answer Engines (e.g., Bing, ChatGPT browsing)	Structured (schema.org, FAQ, HowTo)	Direct Snippets	Intent-based
GEO	Generativity	LLMs (ChatGPT, Claude, Perplexity)	Mixed	AI-generated Summaries	Prompt + RAG
LLM SEO	Interpretability + Retrievability	LLM-native systems (RAG, prompt injection)	Symbolic Files (.txt, .yaml, .json)	Parsed & Grounded Outputs	Semantic + Embedding Search

LLM SEO, therefore, does not replace prior models but extends them into the **LLM-native environment**, where **structured files**, **atomic knowledge units**, and **explicit content cues** are necessary for reliable visibility and citation. It accounts for the **black-box nature** of LLMs and their emergent behaviors in parsing and regurgitating unseen or semi-traceable data.

4.1.3 The Emergence of LLM-First Content Strategies

The rise of AI-powered tools has ushered in a new generation of "AI-native" users, people who increasingly prefer prompting large language models (LLMs) over conducting traditional web searches ^[29]. These users expect AI to handle tasks like comparing prices, drafting code from documentation, or troubleshooting problems using existing resources²⁹. As this behavior becomes more common, it highlights an important shift: content must now be easily digestible by AI, not just by humans.

LLMs process information differently from traditional search engines. While the latter depend on algorithms to interpret massive volumes of unstructured web data, LLMs rely on pattern recognition, contextual associations, and clear semantic structure³⁰. That's why strategies optimized for crawlability, like robots.txt files and XML sitemaps, aren't enough²³. For LLMs, simplified and well-structured content is essential, as it reduces the processing burden and improves comprehension²⁷.

One key limitation of LLMs is their "context window," the limited amount of information they can process at once²⁶. That's where structured formats like llms.txt come in. These files act as curated guides or executive summaries, helping AI tools quickly access a website's most important content without getting bogged down in irrelevant detail³⁴. By presenting only the most valuable information up front, these files help maximize what LLMs can do within their context limits³⁵.

As a result, delivering concise, high-value insights is becoming a make-or-break factor in whether content gets cited or even noticed by AI systems. Content creators must rethink traditional, word-heavy formats and move toward streamlined, modular "atomic content units" that are easier for machines to digest and reuse. Even tl;dr summaries, once just for human skimmers, are now central to optimizing for AI readability.

Companies embracing this "AI-first" approach are already seeing results. Vercel, for example, reports that 10% of its user signups now come via ChatGPT, an outcome directly linked to their Generative Engine Optimization (GEO) strategy³³. Similarly, developer-focused platforms like Pinecone and Windsurf are using llms.txt to make their documentation more accessible to AI, ensuring that API references and tutorials are surfaced accurately³³.

For organizations that fail to make their content machine-readable, the consequences are becoming clear: their documentation, product pages, or policies may become "functionally invisible" in a world increasingly mediated by AI²⁹.

4.2 Understanding llms.txt: A Proposed Standard for LLM-Oriented Optimization

As the digital landscape increasingly shifts toward LLMs as primary information gateways, the need for content that aligns with machine comprehension becomes urgent. Traditional web formats, designed for human consumption and visual rendering, often fall short in their ability to communicate meaning efficiently to AI systems. In response to this emerging challenge, a new class of content structuring standards is taking shape. One of the most promising is the proposed **llms.txt** file.

Introduced by Jeremy Howard of Answer.AI in September 2024⁴⁴, llms.txt is a plain-text, Markdown-formatted document placed at the root of a website (e.g., [yoursite.com/llms.txt](#))⁴⁴. Its purpose is to serve as a lightweight, machine-readable index that distills a site's most valuable and authoritative content into a form that is semantically structured and easily parseable by language models²⁹.

Unlike conventional web architecture, which includes stylistic HTML, dynamic JavaScript, and advertising layers, llms.txt prioritizes clarity over complexity. It intentionally strips away non-essential elements to minimize token usage and reduce the computational cost associated with LLM inference³⁰. In doing so, the file helps LLMs deliver faster, more accurate, and more factually grounded responses when referencing a domain's content³³.

The structure of llms.txt typically includes the following components:

- A top-level header indicating the project or site name;
- A blockquote-style summary offering a concise project or product description;
- Multiple sub-heading sections featuring categorized, Markdown-formatted lists of key resources (e.g., documentation, product pages, policies);
- An optional “low-priority” section to guide LLMs on which links or resources can be deprioritized;
- Additional context such as organizational details, industry categorization, geographic scope, and contact information^{36 | 29}.

The llms.txt file does something surprisingly simple: it gives large language models a clearer path to follow. Instead of wading through bloated code, dynamic elements, or scattered content, the model gets a clean, direct summary of what a site is actually about. That clarity cuts down on errors resulting in fewer hallucinations, more accurate answers and ultimately makes interactions feel smoother and more grounded³³. Think of it as a shortcut or crib sheet that helps AI get to the good stuff faster³⁴.

But there's more to it than just better results. For AI platforms, efficiency matters. Every extra token a model has to process adds cost, especially when you're serving millions of queries a

day. Structured content like llms.txt reduces that load. It’s not just about making LLMs smarter it’s about making them cheaper to run⁴⁹.

What makes this approach appealing is that it doesn’t demand a total rethink of how the web works. Websites stay the same on the surface, but behind the scenes, there’s a lightweight layer that speaks the language of machines. As AI continues to reshape how people find and consume information, this kind of structure feels less like an add-on and more like a necessity.

4.2.1 llms.txt vs. robots.txt and sitemap.xml: Key Distinctions

While llms.txt, robots.txt, and sitemap.xml are all files placed in the root directory of a website to guide web agents, their specific purposes, target audiences, and operational natures differ significantly. Understanding these distinctions is crucial for effective web optimization in the AI era.

Feature	llms.txt	robots.txt	sitemap.xml
Primary Purpose	Guide LLMs to understand and prioritize content for AI optimization (GEO) and knowledge presentation.	Control how search engine bots crawl a site, allowing or preventing access to certain pages/directories.	List all URLs on a site to help search engines discover pages; provides a comprehensive list without much context.
Target Audience	Large Language Model-based AI systems (ChatGPT, Google Gemini, Claude, Perplexity, Bing AI).	Traditional search engine bots (Googlebot, Bingbot, Yandex).	Search engine crawlers.
Format	Markdown, easily readable by humans and machines.	Simple, plain text file with bot-specific browsing rules.	XML
Nature/Enforcement	Guide LLMs to understand and prioritize content for AI optimization (GEO) and knowledge presentation.	Control how search engine bots crawl a site, allowing or preventing access to certain pages/directories.	Informational; helps with discoverability.

Table 1: Key Differences: llms.txt vs. robots.txt vs. Sitemap.xml

Unlike robots.txt, which is a widely respected convention among search engines, llms.txt lacks enforcement. AI models are not obligated to follow it. While robots.txt reliably excludes content from being crawled, llms.txt functions more as a voluntary guideline, leaving its effectiveness entirely up to LLM providers.

This introduces uncertainty for website owners: even carefully curated instructions may be ignored. Ultimately, control over how content is interpreted and surfaced in AI environments still rests with the models. For llms.txt to have meaningful impact, it will require broad adoption—whether through a shared protocol, industry consortium, or a leading platform setting the standard.

4.2.2 Perceived Benefits for AI Understanding and Content Prioritization

Supporters of llms.txt highlight its potential to benefit not only individual site owners but also the broader AI-driven web. As generative tools become primary sources of information, websites face growing pressure to be both legible and discoverable by AI systems. In this context, llms.txt offers a lightweight, practical solution.

A key advantage is its ability to help AI models better understand a site's purpose and structure. Rather than relying on scattered signals, a well-crafted llms.txt provides a clear summary and directs models to high-value pages³⁴. This structure improves response quality, especially when users inquire about brands or complex topics⁴³.

For large sites with extensive subpages or archives, llms.txt acts as a guidepost, helping models focus on priority content and reducing the risk of misinterpretation. This clarity benefits users too, LLMs are less likely to misattribute content or misrepresent a brand's message⁴⁹. By pointing to verified, owner-defined sources, publishers gain greater control over how their content is interpreted³⁴.

Efficiency is another strength. Written in plain Markdown, llms.txt is easier for models to parse than traditional webpages cluttered with HTML and dynamic elements³⁴. This streamlined format can reduce token usage and inference costs, especially at scale⁴⁹.

Importantly, llms.txt serves as a “content map”. Since LLMs often scan only accessible or frequently linked pages, critical content like documentation or policy details may be missed. By listing and categorizing key URLs, llms.txt ensures these pages are surfaced³⁴, functioning like a treasure map or spotlight for valuable content⁴⁴.

However, its effectiveness depends on trust. If the file points to low-quality or irrelevant material, models may disregard it, similar to how search engines abandoned meta keywords after misuse⁴⁴. This underscores the need for llms.txt to reflect genuinely useful and well-maintained content.

Ultimately, llms.txt is a future-proofing tool. As AI models increasingly mediate user access to web content, implementing llms.txt allows publishers to shape how their content is discovered, understood, and represented across AI interfaces³⁴.

4.3 The Great Debate: Google's Skepticism vs. Proponent Perspectives

The rise of llms.txt has sparked active debate within digital and AI communities. On one side are SEO practitioners and AI developers who see it as a forward-looking solution for LLM visibility; on the other is Google, whose cautious response reflects both technical doubt and strategic hesitation. This divide reflects deeper tensions about control, openness, and the future shape of AI-driven discovery.

4.3.1 Google's Stance: Skepticism, Meta Keywords Comparison, and Operational Concerns

Google has remained consistently skeptical about **llms.txt**. In public comments, John Mueller even compared it to the outdated *meta keywords* tag, brushing it off as “not very useful.”³⁵ While the comparison isn't technical, it reflects Google's view that llms.txt may struggle with adoption, much like meta keywords did³⁵.

A key concern is lack of usage. Internal server logs show minimal evidence that major AI systems are accessing llms.txt³⁷. This is supported by webmaster reports and community forums, where hits often come from minor tools like BuiltWith rather than mainstream LLMs³⁷.

Google also questions enforcement. Unlike robots.txt, which is widely respected, llms.txt is voluntary. AI models can ignore it, potentially leading to inconsistent behavior or even cloaking, where different content is served to AI systems and human users³⁵. This type of manipulation contributed to the abandonment of meta keywords.

From a technical perspective, Google sees limited added value. If AI bots can already parse structured data and extract meaning from existing markup, llms.txt may appear redundant³⁵, especially given Google's advanced indexing and snippet systems.

However, some critics suggest Google's stance may be strategic. As one source notes, “it is in Google's interest not to admit factors that may influence results”³⁷. A widely adopted llms.txt protocol could shift influence toward site owners, potentially diluting Google's role as the gatekeeper of online information. Some say Google might be working on its own version of this idea something that follows the same principles as llms.txt, but keeps things in-house rather than supporting an open standard. In that light, Google's hesitation could be less about whether it's technically possible, and more about timing, trust, and maintaining control. At the end of the day, the future of llms.txt might not depend on Google at all. It could come down to whether players like OpenAI, Anthropic, or Perplexity get behind it and if they do, the rest of the industry might just follow.

4.3.2 The Proponents' View: Utility in Multi-Model Ecosystems and Endorsement by LLMs

While Google is still playing it safe, the SEO community and even some major AI platforms are leaning into **LLMS.txt** as a practical way to get content noticed in AI-first search. ChatGPT, for instance, has openly supported the idea, pointing out that large language models don't browse the web like Google does. Instead, they prefer clean, structured text formats like **llms.txt** and **og.json** that make it easier to understand and surface relevant content.

This shift highlights an invisible indexing paradox: what LLMs use to retrieve content isn't always publicly documented. As a result, SEOs now recommend a multi-layered strategy, combining structured data with LLM-specific files like **llms.txt**, especially given the lack of transparency in how models parse content²⁵.

Adoption is already underway. Platforms like Mintlify have implemented **llms.txt** across hosted sites²⁶, reportedly at Anthropic's request³¹. Firms like Profound confirm that models from OpenAI, Microsoft, and others are actively crawling these files³¹.

Why the momentum? LLMs face constraints, limited context windows, real-time access needs, and difficulty parsing complex layouts. **llms.txt** offers a clean, easy-to-interpret list of URLs that works within these limits³⁰.

Supporters are quick to push back against comparisons to outdated tools like meta keywords. They argue that **LLMS.txt** isn't just hidden metadata³² it links to real, live pages and works more like a storefront designed specifically for AI, not humans. Plus, it's entirely optional, which makes it more flexible than rigid formats like AMP.³² At its core, **LLMS.txt** isn't a magic fix, but it is a smart, timely move a way to give LLMs a clearer path to meaningful content in a search landscape that's becoming harder to navigate.

4.3.3 Analysis of Adoption Challenges and Current Status

While **llms.txt** has generated buzz among early adopters, mainstream traction remains elusive. The core issue is clear: no major LLM providers like OpenAI, Anthropic, or Google, currently support it²⁷. That creates a frustrating cycle: without crawler adoption, websites hesitate to implement it, and without widespread implementation, crawlers have little incentive to pay attention.

Server logs and community reports confirm minimal activity around **llms.txt**²⁹, making it a tough internal sell for marketers and content teams. The lack of ROI, combined with the absence of platform commitment, has led to a growing disconnect between SEO enthusiasm and real-world usage.

There are also technical and ethical concerns. Google's John Mueller warns of potential cloaking risks, serving different content to bots than to users which could erode trust²⁹. Additionally, relying on plain-text files may strip away valuable context and interactivity²⁹.

Despite isolated signs of interest, like Vercel developer signups and inquiries from Anthropic²⁶ these remain exceptions. The broader challenge is fragmentation: with no official standard or unified protocol²⁸, content creators may soon face a complex landscape of LLM-specific formats and ingestion strategies.

If the status quo holds, the cost of visibility could rise, forcing teams to juggle multiple formats or lean on third-party tools just to stay compatible.

4.4. Introducing the LLM SEO Files Standard

The conceptualization of the LLM SEO Files Standard originated from a specific empirical observation during the evaluation of AI-generated outputs. A product under my purview was inaccurately described by a prominent large language model (LLM), despite the presence of precise and comprehensive documentation on the official website. This discrepancy underscored a broader systemic issue: content producers currently possess limited agency over how their materials are interpreted, filtered, and represented by generative AI systems.

Unlike traditional search engines, which adhere to established protocols such as robots.txt and structured metadata, LLMs engage with web content through opaque and often inconsistent mechanisms. The absence of standardized interpretive frameworks results in misattributions, loss of nuance, and potential dissemination of misinformation. This observation prompted a foundational inquiry: How might publishers assert greater semantic control over their content when interfacing with LLMs?

Addressing this challenge necessitated a solution beyond a singular directive file. Accordingly, the LLM SEO Files Standard was developed as a structured, multi-layered framework designed to enhance the interpretability of web content by LLMs while simultaneously empowering publishers to delineate the epistemic status of their material, whether factual, interpretive, or contextual.

The proposed standard comprises a suite of six complementary files, each serving a distinct semantic and functional role. Collectively, these components aim to facilitate more accurate parsing, attribution, and representation of content by large language models. The framework aspires to establish a normative interface between content publishers and LLMs, one that promotes transparency, fidelity, and semantic precision in AI-mediated content retrieval and synthesis.

4.4.1 Rationale and Design Principles

This approach builds on over two decades of work in analytics and data storytelling, but it originated from a simple, yet critical observation: *Large Language Models (LLMs) often conflate fact, opinion, and speculation unless content is explicitly structured to help them distinguish*

*between these elements.*⁴³ This insight led to the development of a layered file format designed to separate verifiable claims from perspective-driven content. Importantly, this distinction is not merely technical it is semantic. It signals a broader evolution in how content must be authored, annotated, and published in an AI-first ecosystem. The objective is not only to make content accessible to machine readers but to render it interpretable in ways that preserve human clarity, nuance, and integrity.

The framework also draws from principles of “Generative Engine Optimization (GEO)”, which I’ve worked to formalize over the past year. GEO emphasizes that content must be designed not just for search or human readability, but for LLM comprehension, structured in a way that enables grounded synthesis, clear attribution, and safe summarization. The LLM SEO Files Standard puts that into practice.

Equally important is what this framework implies for publishing workflows. Content teams will need to move beyond traditional editorial models and adopt “structured fact-tagging, metadata labeling, and intent signaling” tasks that were once optional but are quickly becoming essential for visibility in LLM-driven environments.

In short, the LLM SEO Files Standard is meant to restore a sense of “agency, structure, and trust” to a space that currently feels opaque and uneven. It’s a proposal born out of necessity and one I believe reflects where content strategy must go next.

4.4.2 Deep Dive into the 6 Core Files

The “LLM SEO files” standard, as coined by Avinash Tripathi, proposes a layered approach comprising six distinct files: `llms.txt`, `page-name.yaml`, `page-name.md`, `schema.json`, `evidence.json` and `open-graph.json`

4.4.2.1 llms.txt: The Foundational AI Guide

This file is the original `llms.txt` as previously discussed, serving as the primary entry point and foundational guide for LLMs to understand a website’s most important content and overall structure. Its format is Markdown. The core benefits include providing a lightweight summary, quickly surfacing key content, reducing potential model errors, and enhancing the overall user experience by directing AI to the most relevant information.

Inclusion Steps:

- 1.

Download the `llms.txt` file generated for your website.

- 2.

Upload it to the root directory of your website (`/llms.txt`).

3.

Verify access by visiting `yourwebsite.com/llms.txt` in a browser.

4.

Ensure search engines can access it by checking your `robots.txt` file:

User-agent: * Allow: /llms.txt

5.

Monitor AI-generated traffic to your site via analytics tools that detect referral patterns from AI-based engines.

4.4.2.2 page-name.yaml: A Human-Readable Semantic Map for AI Clarity

The YAML file serves as a lightweight, human-readable alternative to JSON for encoding structured semantic information relevant to both editorial workflows and AI parsing engines. Within the LLM SEO framework, `page-name.yaml` plays a pivotal role in ensuring that content is **transparent, logically categorized, and easily consumable** by both machines and human curators. Unlike deeply nested or verbose formats like JSON-LD, YAML offers an intuitive syntax, making it suitable for dual-purpose applications editorial annotation and LLM ingestion.

Purpose and Function

`page-name.yaml` is typically used to describe:

- **Website-level metadata** (title, URL, tags, canonical structure)
- **Thematic taxonomies** (topic clusters, entity relations, content depth)
- **Claim-type categorization**, especially when paired with or mapped to `evidence-of-facts.json`
- **Navigation or content relationships**, similar to a site map but focused on **semantic context**

This file can be embedded in the page's source, referenced via API, or included in an SEO directory for model scraping. When indexed or retrieved by LLMs like ChatGPT, Perplexity, or Google's Gemini, it allows for **faster parsing and richer grounding**, particularly when JSON parsing is constrained or when interpretability is prioritized.

Strategic Advantages:

Feature	Benefit
Human-readable format	Editors and strategists can update semantic mappings without engineering expertise.
Machine-parsable hierarchy	AI models can easily follow the structure to extract topics, tags, claims, and relationships.
Supports hybrid annotations	Can co-exist with JSON and Markdown, acting as a bridge between machine logic and editorial intent.
Content transparency	Reinforces the separation between authorial intention and structured representation .

Example Snippet

title: "The Future of LLM SEO"

url: "https://example.com/llm-seo"

tags:

- LLM SEO
- AI Visibility
- Semantic Search

description: >

This page explores the emerging field of LLM SEO, outlining key strategies for increasing AI citation.

sections:

- title: "Introduction to LLM SEO"
anchor: "#intro"
- title: "Structured Files"
anchor: "#structured-files"
- title: "evidence_and_opinion.yaml"
anchor: "#evidence"

Inclusion Steps:

1.
Download the page-name.yaml file generated for your website.
2.
Upload it to the metadata folder of your website (/config/, /meta/, or /data/).
3.
Ensure correct YAML formatting before deployment, misaligned indentation can break parsing.
4.
Add appropriate references in your HTML by using:
`<link rel="alternate" type="application/yaml" href="/page-name.yaml" />`
5.
Verify integration using a YAML linter or online validator to confirm correct formatting.

AI Relevance

For platforms that prioritize semantic grounding, yaml functions as a compressed guidepost, helping models:

- Establish document hierarchy
- Disambiguate context
- Extract high-priority sections
- Associate tags and claims for relevant query matching

Its simplicity and flexibility make it ideal for interoperable, model-friendly optimization, especially as content ecosystems shift toward machine-mediated discovery rather than traditional keyword search.

4.4.2.3 page-name.md: Human and Machine Readability

The proposal advocates for a broader adoption of Markdown for web content, suggesting page-name.md as a Markdown version of a webpage, accessible by appending .md to the original URL. This emphasis on Markdown is driven by its inherent benefits for AI optimization:

- **Human & Machine Readable:** Markdown is designed for both clear document presentation for humans and easy parsing by AI systems.
- **Optimal for AI-Generated Content & Assistants:** It is frequently utilized in AI-driven knowledge bases, making it a natural fit for content intended for LLM consumption.

- **Lightweight & Easily Portable:** Markdown requires minimal formatting overhead while still retaining a structured output, making it efficient for data transfer and processing.
- **Widely Supported:** Its broad adoption across various platforms, including blogs, GitHub, and AI-driven documentation systems, underscores its versatility.

An extension of this concept is `llms-full.txt`, which compiles all of a site's text into a single Markdown file, simplifying the ingestion of extensive content into AI tools. This specific file structure was developed by Mintlify in collaboration with Anthropic, who sought a cleaner method to feed their entire documentation into LLMs without the complexities of HTML parsing. Data suggests that LLMs are accessing `llms-full.txt` even more frequently than the original `llms.txt`. The strong emphasis on Markdown and `llms-full.txt` highlights LLMs' preference for "clean," structured text data over complex HTML.

Parsing HTML requires significant computational resources and can introduce noise from elements like advertisements or navigation. Providing content in Markdown acts as a pre-processing step for LLMs, making it more cost-effective and faster for them to ingest and utilize the information, potentially for both real-time queries and future model training. This suggests that content providers who offer clean, Markdown versions of their content are providing a valuable service to LLM developers, which could increase the likelihood of their content being included in AI responses and even future model training datasets. This dynamic could lead to a "data quality" competition among websites for AI visibility.

Inclusion Steps:

1.
Download the `page-name.md` file generated for your website.
2.
Upload it to a file accessible to search engines, and make sure it is publicly accessible via direct URL .
3.
Ensure correct MD formatting before deployment, misaligned indentation can break parsing.
4.
Add appropriate references in your HTML by using:
`<link rel="alternate" type="text/markdown" href="/page-name.md" />`

4.4.2.4 Schema JSON: Enhancing Structured Data for AI

This component refers to JSON-LD (JavaScript Object Notation for Linked Data), a widely adopted structured data format primarily used to enhance website visibility in search engines through Schema.org markup.

- **Optimal for Search Engine Understanding:** JSON-LD is directly utilized by major search engines like Google and Bing, as well as AI engines, to generate rich results in search outcomes.
- **Supports Structured Relationships:** It enables the explicit interconnection of entities, providing a semantic web of relationships that LLMs can leverage for deeper understanding.
- **Compact & Standardized:** JSON-LD adheres to a widely accepted, machine-readable format, ensuring consistency and ease of processing.
- **Essential for AI-driven SEO:** AI systems specifically use structured JSON for more profound content analysis, enabling them to interpret content with greater accuracy and context.

ChatGPT explicitly confirms the utility of JSON files, stating, "JSON files are super easy for AI to parse. Unlike unpredictable HTML, a well-structured file... is instantly legible and semantically useful". Structured data is also deemed critical for preventing poor LLM outputs and ensuring the reliability of AI analytics.²² The inclusion of Schema JSON, a well-established SEO practice, within the LLM SEO files standard highlights a significant convergence point between traditional SEO and AEO. Both traditional search engines and LLMs benefit from structured data because it explicitly defines entities and their relationships, reducing ambiguity and improving overall comprehension. This suggests that investments in Schema markup serve a dual purpose, benefiting both human-driven search and AI-driven content consumption, and reinforcing the idea that traditional SEO is not obsolete but rather evolving to meet new demands.

Inclusion Steps:

1.
Download the schema.json file generated for your web-page.
2.
Embed it directly into your webpage by adding the following `<script>` tag in the `<head>` section:

```
<script type="application/ld+json"> { contents of schema.json } </script>
```
3.
Alternatively, upload it to the metadata folder of your website and link to it in your `<head>`:

```
<script type="application/ld+json" src="/schema.json"></script>
```
4.
Validate with Google's Rich Results Test (<https://search.google.com/test/rich-results>) to confirm correct implementation.

5.

Monitor indexing performance using Google Search Console's structured data reports.

4.4.2.5 evidence-of-facts.json: Structuring Epistemic Clarity for LLMs

The `evidence-of-facts.json` file is a specialized LLM SEO artifact designed to disambiguate **factual evidence** from **interpretive opinion** within a page or dataset. This distinction addresses one of the most critical challenges in LLM citation behavior, namely, the tendency of generative models to conflate empirical findings with speculative commentary. By providing a machine-readable breakdown of claims, their sources, and corresponding rationales, this file format enhances the **retrievability**, **attribution fidelity**, and **epistemic integrity** of AI-generated responses.

Purpose and Function:

Unlike structural files such as `site-structure.yaml` (which outlines navigational architecture) or `llms.txt` (which provides a page-level summary), `evidence-of-facts.json` introduces a content-level schema. It serves as a curated epistemological layer, explicitly tagging:

- **Statistical or observational evidence** with source metadata, verification status, and type
- **Subjective or forward-looking opinions** with author credentials and rationale

This fine-grained separation ensures that AI systems, especially those reliant on grounding mechanisms, can:

- Prioritize verifiable statements in responses
- Reduce hallucination risk
- Preserve authorial intent in summarization and rephrasing tasks

Structure and Fields:

A typical `evidence-of-facts.json` file includes:

- **Metadata:** Title, source URL, publication date, and author
- **Evidence Array:** Contains validated factual claims with references and verification info
- **Opinions Array:** Encapsulates expert interpretations, future forecasts, or strategic insights
- **Tags:** Aid in semantic clustering for topic relevance and query matching

Example Entry:

```
{
  "statement": "Pages with structured metadata (e.g., llms.txt) were cited 2.4x more often by Perplexity AI and ChatGPT.",
  "type": "observational",
  "source": "Internal audit (Apr–Jun 2025)",
  "verified_by": "Relific AI Visibility Tracker"
}
```

This file is typically linked to or embedded within the <head> section of content page.

Strategic Impact:

By formally separating “what is known” from “what is believed,” this format reinforces the credibility and machine interpretability of digital content. It reflects a broader movement toward **structured epistemology** in AI optimization, where LLMs are not just prompted with data but guided through it with clarity and provenance. For platforms like Perplexity, Google AI Overviews, and ChatGPT, such files function as grounding blueprints that reduce misattribution and elevate trustworthy citations.

Inclusion Steps:

1.
Download the evidence-of-facts.json file generated for your web-page.
2.
Embed it directly into your webpage by adding the following <script> tag in the <head> section:

```
<script type="application/json"> { contents of schema.json } </script>
```
3.
Alternatively, upload it to the metadata folder of your website and link to it in your <head>:

```
<script type="application/json" src="/evidence-of-facts.json"></script>
```
4.
Monitor indexing performance using Google Search Console's structured data reports.

4.4.2.6 page-name.json: Open Graph Metadata

Open Graph (OG) metadata, initially developed for social media sharing, has become increasingly relevant within the domain of AI optimization. Embedded in the <head> section of a webpage, OG tags define how content is presented when shared or indexed most notably the title, description, preview image, and canonical URL. While originally intended for visual enhancement on social platforms, these tags now serve as machine-readable summaries that

influence how AI systems interpret and display web content.

Purpose and Function:

In the context of LLM SEO, Open Graph metadata provides immediate semantic cues to AI models by:

- Offering a concise and structured summary of page content
- Guiding how responses are framed in AI-generated summaries or citations
- Improving the likelihood of accurate and visually aligned content previews

This metadata acts as a soft signal to language models, supporting alignment between the user query and page-level intent, particularly in the absence of deeper structured elements.

Structure and Key Tags:

An Open Graph implementation typically includes the following properties:

- `og:title`: Represents the content headline or primary topic
- `og:description`: A short summary optimized for both readability and context
- `og:image`: Defines the visual asset associated with the page
- `og:url`: The canonical address for referencing and attribution

Example snippet:

```
<meta property="og:title" content="Understanding LLM SEO Files" />
<meta property="og:description" content="A framework to improve content visibility across
generative AI platforms." />
<meta property="og:image" content="https://example.com/assets/llmseo-cover.jpg" />
<meta property="og:url" content="https://example.com/llm-seo-files" />
```

Strategic Impact:

Though Open Graph metadata lacks the granularity of YAML or schema-based files, it plays a foundational role in first-impression signaling. For AI models like ChatGPT, Perplexity, and Google AI Overviews, OG tags serve as accessible indicators of page relevance and intent, influencing citation probability and preview accuracy. In AI search contexts, well-structured OG metadata enhances both user comprehension and AI model confidence in selecting or referencing the page.

In LLM SEO architecture, this metadata layer contributes to improved content recall, preview fidelity, and citation integrity particularly for content shared or indexed in environments where

deeper parsing may be delayed or deprioritized.

Inclusion Steps:

1.
Download the og.json file generated for your web-page.
2.
Embed it directly into your webpage by adding the contents to the <head> section:
Example:
 <meta property="og:title" content="Page Title" />
 <meta property="og:description" content="The description of page content" />
 <meta property="og:image" content="https://example.com/image.jpg" />
 <meta property="og:url" content="https://example.com/page" />
 <meta property="og:type" content="type of ocntent" />
 <meta property="og:locale" content="en_US" />
3.
Alternatively, upload it to the metadata folder of your website and link to it in your <head>:
 <script type="application/json" src="/page-name.json"></script>
4.
Ensure the file is publicly accessible via direct URL

Table 2: LLM SEO File Types and Their Strategic Role in AI Visibility

File Name	Format	Specific Purpose for AI	Key Benefit for Website
llms.txt	Markdown	Foundational guide to the website's most important content and structure.	Content prioritization, reduced AI model error, improved user experience.
page-name.yaml	YAML	Human-friendly alternative to JSON for separating and labeling claims and	Enables content transparency, simplifies AI parsing, supports editorial and fact-

		interpretations.	checking AI.
page-name.md	Markdown	Provides clean, human- and machine-readable versions of content for efficient parsing.	Enhanced AI-generated content readiness; lightweight and portable content.
schema.json	JSON-LD	Enhances structured data understanding for deeper content analysis and rich result features.	Improved search engine and AI comprehension; supports structured relationships.
evidence-of-facts.json	JSON	Separates evidence (facts) from opinions in a machine-readable schema with metadata.	Improves grounding, reduces hallucination, boosts citation precision across LLM platforms.
og.json	JSON	Defines structured metadata (title, description, image) used by LLMs and social platforms for preview generation.	Boosts LLM SEO by improving how links are summarized and surfaced in AI-generated responses and previews.

4.5. Strategic Implications for Web Publishers and Digital Marketers

The emergence of LLM SEO files standard necessitates a fundamental strategic re-evaluation for any entity with an online presence. Adapting to this new landscape is crucial for maintaining visibility and influence.

4.5.1 Optimizing Content for AI Readability and Citation

To ensure content is effectively consumed and cited by LLMs, web publishers and digital marketers must prioritize specific content characteristics and structural elements:

- **Prioritize Clarity and Structure:** Content should be clear, coherent, and logically structured. This includes using distinct headings (H1, H2, H3), bullet points for lists, and AI-readable tables for comparative data.²³ A well-organized content configuration forms the essential backbone for effective presentation to generative AI.²⁴
- **Focus on Answer-Oriented Content:** Given that LLM users primarily seek immediate, concise answers ²², content should be designed to provide direct, actionable responses. This includes creating dedicated FAQ sections, concise summaries, and clear step-by-step guides.²²
- **Leverage Structured Data (Schema):** Implementing or enhancing Schema JSON-LD markup is essential for deeper content analysis by AI.³⁷ Structured data helps LLMs understand entities, their relationships, and the overall context of content with greater precision ³¹
- **Emphasize Originality and Authority:** LLMs demonstrate a strong preference for content featuring original statistics, research findings (leading to 30-40% higher visibility), expert commentary, and professional insights.³¹ This type of content signals credibility to LLMs and significantly increases citation rates.³¹ The emphasis on "source credibility," "expert commentary," and "original research" ³¹ indicates that LLMs are not merely extracting information but also assessing its underlying authority. When an LLM cites a website, it implicitly transfers some of its own perceived authority to that source. This creates a powerful feedback loop: high-quality, authoritative content is cited more, which in turn reinforces the LLM's credibility and potentially the source's authority within the AI-driven ecosystem. This means that "E-E-A-T" (Experience, Expertise, Authoritativeness, Trustworthiness), a long-standing Google quality guideline ⁸, becomes even more critical for AI visibility. Content strategies should prioritize demonstrating genuine expertise and backing claims with verifiable data and expert opinions, not just for human users but as a direct signal for LLMs.
- **Content Freshness:** LLMs value content that fills temporal gaps in their knowledge, making up-to-date information, especially concerning new technologies or methodologies, particularly attractive for citation.³¹
- **Technical Excellence:** The underlying technical health of a website remains crucial. Sites should be fast, reliable, and easily crawlable by AI bots. This includes ensuring that crawlers are not inadvertently disallowed in robots.txt and promptly fixing any crawl errors.²⁵

4.5.2 Integrating LLM SEO Files into Existing Digital Strategies

The integration of LLM SEO files should be approached as a strategic enhancement rather than a complete overhaul of existing digital marketing efforts.

- **Complement, Not Replace:** LLM SEO files should be viewed as complementary to traditional SEO, not a replacement.²³ A comprehensive digital strategy will capitalize on the strengths of both, optimizing for both human and AI behavior to achieve maximum online visibility.²⁴
- **Pilot Implementation:** For llms.txt, a pragmatic approach involves starting with a basic pilot implementation. This could focus on 5-10 core links, ensuring clear project/company names and descriptions are included.³⁴ This low-cost step allows for early experimentation and potential gains with some LLM platforms.³⁵
- **Content Audit:** A thorough content audit is recommended to identify core documentation, high-performing pages, and resources frequently referenced by users or prospects that would particularly benefit from enhanced AI accessibility.³⁴
- **Regular Updates:** LLM SEO files, especially llms.txt, require regular updates to reflect new content, updated documentation, or changes in site structure.²⁶ This can pose a challenge for manual management due to the frequent evolution of website content.²⁶ The need for constant updates, adherence to Markdown formatting, and proper placement for llms.txt and potentially the other LLM SEO files imply a new form of "technical debt" for websites. If these files are not consistently maintained, they can become outdated, misguide AI, and potentially harm AI visibility.

This adds an ongoing operational burden to digital teams, requiring dedicated resources (time, budget, personnel) for the creation and ongoing maintenance of these AI-specific files. This could drive demand for automated tools or plugins (like those mentioned for WordPress²⁶) that simplify the generation and upkeep of these files, reducing the manual burden.

- **Workflow Integration:** Updates to LLM SEO files should be integrated into existing content review cycles. Content creators should be trained to consider AI consumption patterns when developing new resources, making AI readability a core part of the content creation process.³⁴
- **SEO Co-ordination:** It is essential to ensure that LLM SEO strategies complement traditional SEO efforts. This involves using consistent messaging and coordinating keyword targeting and content themes across both approaches.³⁴

In summary, the LLM SEO Files framework represents a paradigm shift in how digital content is structured, classified, and communicated to generative engines. It integrates principles from SEO, AEO, and GEO, while introducing a symbolic, file-based architecture explicitly designed for machine parsing, grounding, and citation. By embedding semantic clarity through distinct content layers (e.g., facts, opinions, claims, citations), this framework moves beyond traditional web strategies and addresses the core informational needs of LLMs. It

offers not only a technical protocol but also a conceptual foundation for future standards in AI visibility. The next section presents the methodology used to test and validate the effectiveness of this proposed framework.

5. Methodology

5.1 Research Design

This study used a mixed-methods experimental design to explore how effective the proposed LLM SEO Files framework is in improving visibility and citation performance on AI-powered platforms. At its core, the research set out to determine whether layering structured content formats such as llms.txt, schema-enriched JSON-LD, Markdown files, YAML metadata, evidence_of_facts.json file and an open graph .json file could meaningfully influence how large language models (LLMs) surface, ground, and reference web-based information.

The decision to adopt a mixed-methods approach was intentional and grounded in the nature of the problem. Because LLMs operate in complex, often unpredictable ways, purely quantitative analysis would miss important context. By combining numerical performance metrics with qualitative insights into model behavior, the study captures a fuller picture of how LLMs respond to structured interventions.

5.1.1 This design incorporates:

- **Quantitative experimentation**, focused on visibility deltas, citation frequency, and ranking position changes across AI platforms, measured before and after file deployment.
- **Qualitative analysis**, centered on prompt-response behavior, grounding fidelity, hallucination reduction, and the consistency of brand or domain-level mentions.

The dual-mode design ensures that both measurable impact and interpretative nuance are accounted for, offering a comprehensive evaluation of the LLM SEO Files framework and its influence on AI-driven web visibility.

5.1.2 Research Scope and Focus

The evaluation was conducted across four prominent generative AI platforms:

- **Perplexity.ai**
- **Chatgpt**
- **Google AI Overview**

These platforms were selected for their widespread usage, distinct architectural behaviors, and

varying degrees of grounding, citation, and schema interpretation capabilities. By isolating structured content variables and maintaining controlled prompt conditions, the study aimed to answer the following:

- Does the use of LLM SEO Files increase brand mentions in AI-generated responses?
- How does the visibility score differ between non-optimized and LLM SEO-optimized content?
- How do different AI platforms respond to structured content in terms of brand mention frequency and citation behavior across varied user intents?

The mixed-methods design positioned this study to provide both empirical evidence and practical insights for the emerging field of LLM SEO.

5.2 Framework Implementation

The LLM SEO Files framework is comprised of six symbolic file types, each addressing specific content processing needs of LLMs:

- **LLMS Text (.txt)**: A plain executive summary of the important pages placed in the root directory, offering metadata and thematic orientation.
- **Index YAML (.yaml)**: A machine-readable taxonomy defining entities, topic hierarchies, and canonical tags to improve grounding and disambiguation.
- **Markdown (.md)**: Semantic structuring using markdown headers, tables, and bullet points to enhance chunking and readability.
- **Schema Markup (JSON-LD)**: Rich structured data defining content types (e.g., FAQ, Author, Product), factual claims, and entity relationships.
- **Evidence of facts (.json)**: A novel contribution, the evidence.json file contains curated citations to primary sources, annotated by evidence type (e.g., peer-reviewed, official documentation).
- **Open graph (.json)**: Machine-readable metadata for link previews (title, image, description) designed for social platforms and LLM crawlers. Enhances how AI systems and answer engines summarize and display your content in chat results, boosting visibility and click-through from AI-generated responses.

These formats were embedded into selected low-performing and mid-performing web content to test their effect on LLM-driven search visibility.

5.3 Data Collection

To evaluate the performance of LLM SEO Files in real-world AI environments, a systematic and

controlled data collection process was implemented. This process focused on prompt diversity, content variation, and standardized testing protocols.

5.3.1. Prompt Set Creation

A curated dataset of over 100 unique prompts was developed, covering a range of high-visibility domains, including:

- **AWARENESS:** 30%
- **INTEREST:** 30%
- **CONSIDERATION:** 28%
- **PURCHASE:** 12%

Each prompt was derived from the core topics and content we actively work with, mapped across key stages of the marketing funnel Awareness, Interest, Consideration, and Purchase. This ensured that the evaluation remained grounded in real-world use cases and captured diverse visibility behaviors aligned with practical brand objectives.

5.3.2 Content Versioning and Testing Groups

For every prompt, two distinct versions of content were created:

- **Baseline Version:** a non-optimized content
- **LLM SEO Version:** The same topic rewritten and structured using all six file formats (llms.txt, page-name.yaml, page-name.md, schema.json, evidence.json and open-graph.json), following the LLM SEO Files framework.

These parallel versions enabled a direct comparison of performance across AI model

5.3.3 Controlled Deployment Conditions

To ensure consistency and reduce bias in visibility testing:

- We followed a structured and repeatable approach across all platforms. Every prompt was submitted under the exact same conditions using the same accounts, browser settings, and devices.
- We ran tests at consistent times of day to minimize any changes in AI behavior caused by server load or system updates.
- To avoid memory contamination, we started fresh sessions (like new chats in ChatGPT and Perplexity) between each test to make sure the influence of previous prompts was completely removed.
- To capture how AI responses might vary by region, we tested from five different geographic locations: India, the USA, the UK, Australia, and Germany. Each response was timestamped and logged in Airtable, helping us track platform behavior and performance over time.

This meticulous process gave us a solid, unbiased foundation to benchmark LLM SEO file performance and compare it directly with traditional SEO techniques.

5.4 Evaluation Criteria

To evaluate how well the LLM SEO Files framework performs, we used three key metrics. These were carefully chosen to measure both visibility in AI-generated responses and how well the content aligned with the original source across different platforms.

- **Brand Mention:**
Whether the brand name or entity was explicitly referenced in the AI-generated response (e.g., product name, organization, author).
- **Domain Citation:**
Whether the AI output included a link or reference to the original website domain (e.g., example.com), either as a source, snippet, or embedded link.
- **AIM Visibility Score:**
A composite metric combining mention prominence, contextual alignment with the original content, and response for different geographical location.

Each prompt was evaluated **five times from different location** , once using the baseline non-optimized content and once after deploying the structured LLM SEO Files. This allowed for comparative tracking of changes in platform behavior and attribution format over time.

5.5 Experiment Phases

- **Phase 1 - Baseline Testing:**
 - Ran prompts on ChatGPT, Perplexity, Google AI Overviews using non-optimized content.
 - Logged all citations, mentions, and visibility outcomes.
- **Phase 2 - LLM SEO Implementation:**
 - Structured content using llms.txt, page-name.yaml, page-name.md, schema.json, evidence.json and open-graph.json .
 - Published structured files publicly and indexed using standard web crawlers
- **Phase 3 - Comparative Measurement:**
 - Re-ran all 100 prompts after a 110 DAYS indexing window from the same 5 locations
 - Measured deltas in brand mentions, domain citations, and composite visibility scores.

5.6 Tools and Platforms Used

The study made use of a range of tools including AI platforms, and custom logging setups to test, track, and evaluate how LLM SEO Files performed across different environments. These tools were carefully selected to ensure consistent testing conditions and to make results reproducible across both qualitative insights and quantitative metrics.

LLM Interfaces (Used to Assess Visibility and Citations)

- **ChatGPT (GPT-4-turbo, Pro version)** was used to examine how LLMs handled grounding, brand visibility, and citation behaviors in response to structured prompts. Both browsing and plugin modes were tested via the OpenAI Responses API to capture a broad range of interactions and outputs.
- **Perplexity AI (Pro with Sonar Large)** was used to monitor how often sources were cited, how follow-up questions were handled, and how the platform structured its ranked outputs. Logging was performed using the Sonar Pro API, allowing the research team to measure brand mentions and domain-level references with precision.
- **Google AI Overviews (via the Search Generative Experience)** were analyzed using the Google SERP API. The goal here was to observe whether schema-enhanced content such as JSON-LD markup and structured summaries was correctly interpreted and displayed in AI-generated overviews as well as in standard search snippets.

AI Monitor (Open Source)

A purpose-built open-source tool developed for this study to track and quantify:

- **Brand Mentions** – Identifies explicit brand or product references in LLM outputs.
- **Domain Citations** – Detects and logs website-level citations and backlinks.
- **AIM Visibility Score Logging** – Computes a composite score based on appearance prominence, match context, and citation type.

AI Monitor significantly reduced the manual overhead typically required for visibility benchmarking and provided a scalable, reproducible method for analyzing LLM responses across platforms, prompts, and geographies.

5.7 Limitations and Bias Controls

Despite the structured experimental design, this study acknowledges several inherent limitations associated with evaluating LLM-driven visibility and content attribution:

- **Probabilistic Output Variability:**
One key limitation stems from the probabilistic behavior of large language models. Because these models do not return fixed outputs, the same prompt can lead to different responses depending on the session or geographic location. This inherent variability complicates efforts to measure brand mentions or visibility with precision, particularly when responses differ subtly but meaningfully.
- **Opaque Source Attribution Mechanisms:**
Another challenge is the lack of transparency in how some AI platforms handle source attribution. Google's AI Overviews, for instance, often return citations without offering insight into why certain sources were chosen. This ambiguity makes it difficult to determine whether the AI truly drew from the optimized content or relied on secondary summaries or unrelated sources.
- **Non-Uniform Indexing Behavior:**
The time lag between content publication and discoverability in LLM outputs is non-uniform and platform-dependent. Indexing is not instantaneous nor transparently reported by most AI systems, making it challenging to precisely determine whether changes in visibility are due to content optimization or delayed crawling and caching.
- **Manual Annotation and Human Bias:**
Although structured logging protocols were used, parts of the evaluation such as identifying brand mentions or counting domain-level citations required manual inspection. This introduces the potential for human error and subjective judgment, which may impact the consistency of results. Future implementations should adopt programmatic tracking systems or annotation tools to reduce bias and improve repeatability.
- **Absence of a Formal Control Group:** This study did not include a strict control group of completely unoptimized content evaluated in parallel over time. While pre-optimization baselines were captured, the lack of simultaneous control conditions limits the ability to fully isolate the causal impact of specific file-based optimizations on visibility performance across platforms.

5.8 Methodology Flow Diagram

graph TD

```

A[Prompt Dataset Creation (100 prompts)] --> B[Baseline Visibility Testing]
B --> C[Implement LLM SEO Files on Web Content]
C --> D[Wait 110 Days for Indexing]
D --> E[Re-run Prompts Across LLM Platforms]
E --> F[Log Brand Mentions, Domain Citations, Visibility Scores]
F --> G[Compare Results: Pre vs Post Implementation]

```

This structured methodology ensures transparency, repeatability, and real-world application of the LLM SEO Files framework across AI-driven platforms.

6. Results and Analysis

6.1 Overview of the Experiment

We wanted to see whether using structured LLM SEO Files actually made any difference when it comes to AI visibility and attribution. To test this, we pulled together 100 prompts across five verticals software, sustainability, finance, health, and education. The prompts weren't just random; we wrote them to reflect how people really search, covering everything from early awareness to the point where someone's ready to take action or make a purchase.

The experiment was conducted across three major generative AI platforms: **ChatGPT (OpenAI)**, **Perplexity.ai**, and **Google AI Overviews**. Each prompt was tested under two distinct conditions:

- **Control Group:** Content presented in a standard web HTML format, with no additional optimization.
- **Experimental Group:** Identical content structured using **modular LLM SEO Files** including `llms.txt`, `page-name.yaml`, `page-name.md`, `schema.json`, `evidence.json` and `open-graph.json`

The core objective was to assess whether structured LLM SEO Files improved **AI visibility**, **brand attribution**, and **domain-level citation** across platforms.

6.2 Experimental Protocol

To keep things consistent, we used the same 100 prompts across both the control and experimental groups. These prompts were tested under identical conditions on each of the three platforms: ChatGPT, Perplexity, and Google AI Overviews. We also ran the PROMPTS from five different geographic locations to account for regional differences in how AI responses are generated. Daily snapshots were taken to track not just visibility and brand mentions, but also how and when platforms cited our content.

The experiment followed a two-phase timeline:

Phase 1: Control Group Testing (Unoptimized Content)

For the first 15 days, we used the original unoptimized content just basic HTML web pages. This gave us a baseline to measure inclusion, brand visibility, and citations across the platforms without any structured SEO intervention.

Phase 2: Experimental Group Testing (Optimized Content)

After the control period concluded, the same content was enhanced with modular LLM SEO Files `llms.txt`, `page-name.yaml`, `page-name.md`, `schema.json`, `evidence.json` and `open-graph.json`. The optimized version was then allowed to propagate naturally through platform crawling for 110 to 120 days.

Analysis Initiation:

Only after the 110–120 day waiting period was complete did we begin the final visibility and citation analysis to ensure sufficient indexation and response pattern stabilization.

6.3 Evaluation Metrics

To assess platform-level visibility and brand attribution, three primary metrics were used each based on **quantitative counts** of how often a brand or domain appeared across AI-generated responses:

Metric	Description
Brand Mentions	Total number of times the brand name was explicitly mentioned in AI-generated outputs.
Domain Citations	Total number of times the domain (e.g., example .com) was cited, hyperlinked, or referenced.
AIM Visibility Score	Composite metric designed to measure how prominently the brand appeared in the output, combining ranking position, context relevance, and anchor prominence.

The AIM Visibility Score (*AI Mention Index*) is a custom metric designed to assess not just presence but quality of attribution in AI responses. It incorporates:

- Brand Mention Count (weighted)
- Domain Citation Count
- Position and Contextual Relevance of mentions within the generated response

This scoring system allows us to quantify AI visibility similar to how search engines use SERP rank, but tailored for generative outputs.

6.4 Quantitative Findings

This section presents the quantitative outcomes of our quasi-experimental study, structured across three layers of analysis: prompt-level, platform-level, and intent-level performance. Visibility metrics including brand mention, domain citation, and AIM visibility scores were collected across both control and experimental conditions from a dataset of 100 prompts tested across ChatGPT, Perplexity, and Google AI Overviews.

6.4.1: Prompt-Based Visibility Trends

To illustrate the platform and content behavior in greater detail, we have listed out four representative queries selected from the larger 100-prompt dataset. These prompts were chosen based on diversity of intent (Awareness to Purchase) and thematic relevance across the core content categories. They serve as high-quality exemplars to demonstrate how LLM SEO Files influence AI-generated visibility and citations in real-world conditions.

Following the prompt-level breakdown, we present Intent -based comparisons to evaluate system responsiveness to structured optimization, and conclude with platform-specific trends to assess commercial relevance and stage-specific impact.

Prompt 1: “Free Answer Engine Optimization Course”

Intent : AWARENESS

Platform	Mention (without LLM SEO)	Citation (without LLM SEO)	Mention (with LLM SEO)	Citation (with LLM SEO)
ChatGPT	1	0	1	1
Perplexity	0	0	1	1
Google AI	0	0	1	2
AIM VESIBILITY SCORE	0.8		5	

Prompt 2: “Best Generative Engine optimization agency in New Delhi”

Intent: INTEREST

Platform	Mention (without LLM SEO)	Citation (without LLM SEO)	Mention (with LLM SEO)	Citation (with LLM SEO)
ChatGPT	0	0	1	1
Perplexity	0	0	1	0
Google AI	1	1	1	1
AIM VESIBILITY SCORE	1.7		4.7	

Prompt 3: “ai monitor generative engine optimization tool review”

Intent: CONSIDERATION

Platform	Mention (without LLM SEO)	Citation (without LLM SEO)	Mention (with LLM SEO)	Citation (with LLM SEO)
ChatGPT	0	0	1	0
Perplexity	0	0	1	1
Google AI	0	1	1	1
AIM VESIBILITY SCORE	0.8		3.3	

Prompt 4: Best SaaS procurement platform

Intent: PURCHASE

Platform	Mention (without LLM SEO)	Citation (without LLM SEO)	Mention ((with LLM SEO)	Citation (with LLM SEO)
ChatGPT	1	1	1	1
Perplexity	0	0	1	1
Google AI	0	0	0	1
AIM VESIBILITY SCORE	1.7		4.2	

6.4.2 Intent-Based Visibility Trends

Beyond prompts -specific performance, it is essential to understand how visibility outcomes vary across different stages of the user journey. To this end, we categorized all 100 prompts into four core intent segments: **Awareness**, **Interest**, **Consideration**, and **Purchase**. These categories reflect typical marketing funnel stages and were assigned based on the informational or transactional nature of each prompt.

This breakdown helped us see whether structured LLM SEO Files made a bigger difference at certain stages. For example, did they help raise visibility for broad educational content? Or were they more useful in driving citations for bottom-of-funnel, purchase-related queries? The following section summarizes how each intent group performed looking at how often the content appeared, whether it included brand mentions, and how frequently domains were cited.

INTENT : AWARENESS (30 prompts)

Platform	Mention (without LLM SEO)	Citation (without LLM SEO)	Mention (with LLM SEO)	Citation (with LLM SEO)
ChatGPT	7	5	16	15
Perplexity	9	4	19	18
Google AI	4	1	24	23
Avg. AIM VESIBILITY SCORE	2.6		4.1	

INTENT : Interest (30 prompts)

Platform	Mention (without LLM SEO)	Citation (without LLM SEO)	Mention (with LLM SEO)	Citation (with LLM SEO)
ChatGPT	5	2	12	10
Perplexity	6	5	22	13
Google AI	4	1	21	16
Avg. AIM VESIBILITY SCORE	2.1		3.6	

INTENT : CONSIDERATION (28 prompts)

Platform	Mention (without LLM SEO)	Citation (without LLM SEO)	Mention (with LLM SEO)	Citation ((with LLM SEO)
ChatGPT	6	7	14	20
Perplexity	9	3	23	19
Google AI	5	4	19	15
Avg. AIM VESIBILITY SCORE	2.4		3.8	

INTENT : PURCHASE (12 prompts)

Platform	Mention (without LLM SEO)	Citation (without LLM SEO)	Mention (LLM SEO)	Citation (LLM SEO)
ChatGPT	3	2	12	10
Perplexity	4	5	22	13
Google AI	6	1	21	16
Avg. AIM VESIBILITY SCORE	2.1		3.8	

6.4.3 Platform-Level Visibility Trends: Brand Mentions and Citations

In addition to intent-stage analysis, we examined how each AI platform responded to structured content when measuring two key visibility metrics:

- **Brand Mentions:** How often a brand, product, or source was referenced (implicitly or explicitly).
- **Citation Frequency:** How often a clickable or attributed link/source was provided.

Each metric was computed across the full 100-prompt dataset, comparing **Control** (non-optimized content) and **LLM SEO** (structured files deployed) groups.

Table: Brand Mentions and Citations by Platform (Control vs. LLM SEO)
Brand Mentions (LLM SEO)

Platform	Brand Mention s (without LLM SEO)	Citation Frequen cy (without LLM SEO)	Brand Mentions (LLM SEO)	Citation Freque ncy (LLM SEO)
ChatGPT	100	78	154	132

Perplexity	69	79	112	129
Google AI Overviews	45	76	102	110
Avg. AIM VESIBILITY SCORE	2.3		3.8	

The results presented across intent categories, platform performances, and evaluation metrics reveal clear patterns of improved AI-driven visibility through structured content interventions. What's interesting is that this improvement wasn't limited to just one type of content or platform. Whether the query was broad and educational or closer to a buying decision, optimized content had a better chance of being included in AI-generated responses. The structured format seemed to help the platforms "understand" and surface the information more reliably.

That said, the level of improvement did vary depending on the platform and content category, which suggests there's no one-size-fits-all formula. But overall, these results support the idea that LLM-friendly formatting plays a growing role in AI visibility one that's likely to become even more important as generative platforms evolve. In the next section, we'll unpack these findings in more depth, explore a few of the limitations we encountered, and share some recommendations for how content teams can adapt to this changing landscape.

7. Discussion

To show how structured content optimization with LLM SEO Files actually works in practice, we ran a set of four test prompts. Each prompt represented a key stage in the customer journey from early awareness to the final purchase decision. We carefully designed these prompts to test brand visibility and attribution on leading generative AI platforms: ChatGPT, Perplexity, and Google AI Overviews.

Each prompt was tested in two ways: first, as a basic version without any optimization (our control), and second, as an optimized version using LLM SEO Files. This side-by-side setup allowed us to directly compare how each version performed in terms of brand mentions and domain citations.

To measure how well brands showed up in AI-generated responses, we created something called the "AI Monitor (AIM) Visibility Score". It's a composite score that gives us a clearer picture of a brand's presence across platforms. The score looks at several key factors:

- How often the brand name appears
- Where it's mentioned in the response (earlier mentions carry more weight)
- Whether the same brand is repeated too often(to avoid artificially inflated scores)
- How frequently the brand's actual website is linked or cited
- Whether the visibility is consistent across multiple regions

We tested each prompt across five different countries - the U.S., India, the U.K., Germany, and Australia to see how results varied (or stayed consistent) across global AI deployments. With this setup, we were able to track how structured optimization affected brand visibility across both platforms and regions. Using the AIM Score, we compared the non-optimized and optimized versions of each prompt and uncovered clear differences in how AI systems handled the content.

The results were as follows:

- **Prompt A (Awareness Stage):** AIM Score – **0.8 - 5**
- **Prompt B (Interest Stage):** AIM Score – **1.7- 4.7**
- **Prompt C (Consideration Stage):** AIM Score –**0.8 - 3.3**
- **Prompt D (Purchase Stage):** AIM Score – **1.7- 4.2**

Each of these scores reflects an improvement in visibility when structured with LLM SEO Files, with brands consistently mentioned across multiple platforms in various regions. In contrast, the control group prompts lacking structured optimization averaged between **0.8 and 1.8**, with minimal brand mentions and sparse domain citations. These findings highlight the measurable gains in AI visibility driven by structured file deployment and reinforce the effectiveness of the LLM SEO approach across international contexts.

The findings from this comparative setup substantiate the hypothesis that structured content significantly enhances visibility in large language model outputs. Across both **intent-stage** and **platform-level** evaluations, we observed consistent and measurable improvements in **brand mentions** and **citation frequency** following the implementation of LLM SEO strategies. These results reinforce the strategic utility of structuring digital content in machine-readable formats specifically tailored for generative systems.

At the **intent level**, the data revealed a progressive increase in visibility as users move through the funnel from Awareness to Purchase. Notably, **brand mentions** increased from 20 to 59, while **citations** rose from 10 to 56 in the Awareness stage alone. The Purchase stage displayed the most significant gains, with brand mentions increased from 13 to 55 citations jumping from 8 to 39 an indicator that structured content is particularly influential during high-intent decision-making queries. This suggests that LLM SEO Files not only improve surface-level discoverability but also bolster visibility at critical conversion points, making them a viable optimization tool across the full user journey.

At the **platform level**, visibility enhancements were evident across all four evaluated systems: ChatGPT, Perplexity, Google AI Overviews, and AIM Visibility. ChatGPT showed the strongest gain in citation frequency, while even platforms known for high citation barriers, such as Google AI Overviews, displayed measurable increases. This indicates that the structure, clarity, and

machine-readable formatting embedded in LLM SEO Files are not merely preferential for one model architecture, but universally advantageous across LLM ecosystems. The increased **citation frequency** suggests that content authored or structured with AI-parsable intent has a higher likelihood of being credited, a core challenge in the current generative content landscape.

More broadly, the shift from very low visibility in non-optimized content to consistent mentions and citations in LLM SEO-optimized outputs demonstrates that LLMs respond favorably to content that conforms to predictable, semantic, and structured design. This aligns with previous research highlighting the importance of schema markup and structured data in traditional SEO (Enge et al., 2020), and extends its relevance to the new paradigm of generative engines.

These results have critical implications for **content strategy**, **AI discoverability**, and **brand attribution**. Structured optimization through LLM SEO Files represents a forward-compatible approach designed not just to serve current AI behavior, but to align with the future trajectory of search, where LLMs will likely become primary gateways to information. As such, adopting this framework is not just a tactical content improvement, but a strategic positioning move in the evolving ecosystem of AI-native web visibility.

This study provides initial evidence that structured content using the LLM SEO Files Standard enhances brand visibility and citation consistency across generative AI platforms. Gains observed across user intent stages, platform types, and geographic contexts suggest that LLM SEO files offers a viable framework for optimizing content in the era of machine-mediated discovery.

However, these findings are based on a limited sample of 100 prompts and should be interpreted with caution. The narrow scope precludes generalization across industries, content types, and model architectures. Future research should expand the dataset to encompass a broader range of verticals, prompt lengths, and languages, while systematically evaluating performance across diverse LLM configurations and fine-tuning parameters.

Critical questions remain regarding the long-term impact of LLM SEO files on trustworthiness, hallucination mitigation, and factual grounding. As generative systems evolve, ongoing empirical assessment will be essential. Structured, machine-readable formats are not merely tactical enhancements. They represent a strategic imperative for sustaining digital visibility and semantic precision in AI-driven ecosystems.

8. Limitations and Future Research Directions

Despite the encouraging results of this study, several limitations must be acknowledged to contextualize the findings and inform the direction of future research. These limitations pertain to the scope of the experimental data, the evolving and opaque nature of large language model (LLM) behavior, and constraints inherent to the research design.

8.1 Limited Dataset Scope and Content Domains

This study was conducted using a dataset of 100 prompts, primarily targeting intents such as awareness , interest , consideration ,and purchase While this controlled dataset enabled pattern

identification and early-stage hypothesis validation, it does not cover the full spectrum of content categories that generative AI platforms encounter in practice. Key verticals including healthcare, education, legal, financial services, and e-commerce were not included in this iteration. Additionally, the study was limited to English-language prompts and did not assess visibility across multilingual ecosystems or regional content formats. As a result, the findings should be interpreted as directional rather than universally generalizable.

8.2 Platform Variability and Model Non-Determinism

Generative AI systems are inherently dynamic and non-deterministic. Platforms such as ChatGPT, Google AI Overviews, Perplexity, and Claude frequently undergo model updates and prompt handling changes, which affect output behavior, citation logic, and brand mention likelihood. A single prompt may produce variable results across different sessions, geographies, or account states. This lack of reproducibility presents a core challenge for establishing consistent benchmarks or longitudinal comparisons. Consequently, the results of this study must be viewed within the temporal constraints of the model versions used during testing.

8.3 Lack of Transparency in AI Ranking and Citation Logic

Another limitation arises from the black-box nature of current LLM platforms. The underlying algorithms that govern source ranking, snippet generation, and citation placement are proprietary and opaque. While generative platforms may display citations or links in some cases, there is no public documentation detailing the mechanisms by which these sources are selected, ranked, or grounded. This restricts the ability to make causal claims about how LLM SEO Files directly influence AI citation behavior, leading the study to rely on observed correlations rather than verified internal model mechanisms.

8.4 Controlled Experimental Design Versus Real-World Web Dynamics

The study was conducted under a semi-controlled experimental setup using synthetic prompts and test domains specifically structured with and without LLM SEO Files. While this methodology enabled variable isolation and performance comparison, it does not fully replicate the complex environment of live web content. Real-world visibility outcomes are influenced by a range of confounding factors such as domain authority, backlink ecosystems, page crawl budgets, user personalization layers, and competitive content presence. As such, while the internal validity of the findings is strong, the external validity remains limited until further tested under naturalistic web conditions.

8.5 Adoption Barriers and Technical Trade-Offs

The introduction of structured content formats such as `llms.txt`, `page-name.yaml`, `page-name.md`, `schema.json`, `evidence.json` and `open-graph.json` files introduces a learning curve for content creators, marketers, and editorial teams. While these formats enhance machine interpretability and grounding potential, they may simultaneously hinder content production efficiency and narrative flexibility. Moreover, over-optimization through rigid structure

could inadvertently introduce hallucination triggers or diminish the authenticity of human-authored content. This trade-off between algorithmic visibility and editorial fluency requires further study and practical guidance.

8.6 Temporal Limitation and Short-Term Snapshot

The present study represents a cross-sectional analysis a snapshot of LLM behavior at a specific point in time (2025). It does not account for the longitudinal effects of structured content on persistent citation, brand trust-building, or long-term retrieval patterns. The absence of time-series or repeat-query testing means that the durability and decay of LLM SEO visibility gains over time remain unexplored.

9. Future Research Directions

While this study offers strong initial evidence of the impact of LLM SEO Files on AI-driven visibility, it also opens the door to many new questions. As generative search continues to evolve, future research can help uncover how structured content strategies perform across different industries, platforms, user types, and use cases. The following areas outline where deeper investigation is needed to both broaden and refine our understanding of this emerging field.

9.1 Expanding Across Industries and Languages

This research focused mainly on software, SaaS, and marketing content in English. But to truly understand how effective LLM SEO Files are, we need to see how they perform in different industries like healthcare, education, travel, finance, and law where the stakes, regulations, and audience needs are very different. We also need to explore how this approach works in non-English environments. Do AI systems interpret structured files the same way in Spanish, Hindi, or Arabic? Can local cultural or regulatory differences affect how content is ranked or cited? Broadening the scope will help test how adaptable and globally relevant this strategy really is.

9.2 Tracking Changes Over Time

LLM platforms change fast what works today might not work next month. That's why future research should focus on tracking brand visibility and citation behavior over time, especially as models evolve or are fine-tuned. Long-term studies could show whether structured content remains effective, how it holds up through algorithm changes, and where new risks or opportunities emerge. Time-based benchmarks would also help content teams know when to update or revise their structured files.

9.3 Partnering with Model Developers for Transparency

Much of this research relies on black-box observation we see the outputs, but not how the AI decided what to include. Working with open-source developers or research-friendly platforms (like Perplexity Labs or Hugging Face) could help us peek under the hood. If researchers had access to internal systems like ranking logic, token weighting, or grounding decisions they could design more precise experiments and help shape better visibility standards from the inside out.

9.4 Understanding Human Impact and Usability

Finally, future work should look beyond technical effectiveness and explore how humans both content creators and readers experience structured content. Do editors find YAML or schema markup too difficult to implement? Does AI-cited content seem more trustworthy to users? And does writing for machines compromise creativity or clarity for human readers? These questions are critical to ensuring that LLM SEO is not just functional but also usable, ethical, and aligned with human-centered communication.

10. Conclusion

This study introduced **LLM SEO Files** as a structured, multi-format content framework designed to improve visibility and attribution across large language model (LLM)-powered platforms such as ChatGPT, Perplexity, Claude, and Google AI Overviews. As generative engines increasingly shape how people discover and engage with information online, the rules of visibility are changing favoring machine-readable, semantically clear, and contextually grounded content over traditional SEO signals alone.

Through a quasi-experimental design and a dataset of 100 prompts, we demonstrated that structured optimization using LLM SEO Files led to consistent improvements in brand mentions and domain citations across multiple platforms, user intent stages, and geographic regions. These gains were especially pronounced in high-intent stages like “Awareness” and “Purchase,” suggesting that LLM SEO may be most impactful where credibility and decision-making matter most.

Importantly, the results point to more than just short-term wins. They highlight a strategic opportunity for organizations to **future-proof their digital presence** by aligning with how LLMs interpret and prioritize content. The rise of AI-native search calls for a shift in how content is created not just to inform humans, but also to be parsed, grounded, and cited by intelligent systems.

That said, this is still an early-stage framework. The results, while promising, are not yet generalizable across all industries, languages, or AI models. As discussed, further research is needed to evaluate long-term performance, cross-domain adaptability, and the balance between machine-readability and human-centered design.

LLM SEO Files represent more than a technical tactic they offer a **strategic blueprint** for engaging with the next generation of search. As LLMs evolve into the primary gateways to information, structured content may be the key to ensuring that authoritative voices brands, experts, educators, and publishers are not just present, but actually seen and cited in the age of generative AI.

11. References

1. Google's Knowledge Graph

<https://support.google.com/knowledgepanel/answer/9787176?hl=en>

2 . OpenAI, Perplexity, and Anthropic Documentation (2023–2025)

- OpenAI: <https://platform.openai.com/docs>

- Perplexity AI: <https://docs.perplexity.ai>

- Anthropic (Claude): <https://docs.anthropic.com>

3 . Karanxha, Giulia (2025) – *Evaluating Transparency in Artificial Intelligence Systems: Adapting the Z-Inspection® Framework for the MANOLO Project* [Full PDF via Theseus.fi]

4. Fishkin, R. (2015). *The beginner's guide to SEO*. Moz.

<https://d2eeipcrd6.cloudfront.net/guides/Moz-The-Beginners-Guide-To-SEO.pdf>

5. Matt Cutts, "How does Google rank search results?" *YouTube*, May 1, 2010, <https://www.youtube.com/watch?v=BNHR6IQJGZs>; and various entries on the Google Webmaster Blog (2010–2013), <https://developers.google.com/search/blog>.

6. Enge, E., Spencer, S., Stricchiola, J., & Fishkin, R. (2020). *The art of SEO: Mastering search engine optimization* (3rd ed.). O'Reilly Media. <https://www.oreilly.com/library/view/the-art-of/9781491903629/>

7. *Generative engine optimization*. arXiv. <https://arxiv.org/abs/2311.09735>

8. Tripathi, A. (2024). *Generative engine optimization: A new frontier for AI visibility*. [White paper].

9. OpenAI. (2023). *How ChatGPT retrieves and cites information*. OpenAI.
<https://openai.com/research>
10. Google. (2023). *AI Overview technical whitepaper*. Google Research.
<https://research.google/pubs>
11. GenBE. (2025). *GEO strategy playbook: Structuring content for AI engines*. GenBE Publishing. [Internal publication or white paper].
12. Sullivan, D. (2022). *What is Google Search Central?* Google Search Central Blog.
<https://developers.google.com/search/blog>
13. van Dis, E. A. M., Binnendijk, L. A., Whittlestone, J., Kakarmath, S., Krubiner, C., Morley, J., & Larsson, D. (2023). *Responsibly deploying language models: Opportunities, challenges, and guidelines for policymakers*.
<https://doi.org/10.1126/science.adg7879>
14. Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-t., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). *Retrieval-augmented generation for knowledge-intensive NLP tasks*. In *Advances in Neural Information Processing Systems (NeurIPS 2020)* (Vol. 33, pp. 9459–9474). Curran Associates, Inc.
<https://proceedings.neurips.cc/paper/2020/file/6b493230205f780e1bc26945df7481e5-Paper.pdf>
15. LangChain. (2024). *Concepts: Retrieval-Augmented Generation (RAG)*. LangChain Docs.
<https://js.langchain.com/docs/concepts/rag>
16. Mendoza Juan, Y. (2024). *Multi-agent LLM systems for enhanced human-machine interaction* (Master's thesis, Universitat Politècnica de Catalunya). UPCommons.
<https://upcommons.upc.edu/bitstream/handle/2117/414255/TFM%20-%20Multi-Agent%20LLM%20System.pdf>
17. OpenAI. (2024). *How ChatGPT retrieves and cites information* [Large language

model overview]. OpenAI.

<https://openai.com/index/introducing-chatgpt-search/>

18. IBM Research. (2023). *AI transparency and explainability in language models*. IBM.

<https://www.ibm.com/think/topics/ai-transparency>

19. Gillespie, C., & Lovelace, R. (2016). *Efficient R programming: Techniques for tokenization and content formatting*.

https://books.google.co.in/books/about/Efficient_R_Programming.html?id=YUavDQAAQBAJ&redir_esc=y

20. Evans, E. (2003). *Domain-driven design: Tackling complexity in the heart of software*. Addison-Wesley. Available as final manuscript (April 15, 2003):

<http://fabiofumarola.github.io/nosql/readingMaterial/Evans03.pdf>

21. Gillespie, C., & Lovelace, R. (2016). *Efficient R programming: A practical guide to smarter programming*. O'Reilly Media.

Online version: <https://csgillespie.github.io/efficientR/>

22. SEO vs GEO: A Detailed Guide for Every Marketer - TripleDart, accessed July 18, 2025,

<https://www.tripledart.com/saas-seo/seo-vs-geo>

23. LLMs.txt: The New SEO Hack for AI-Powered Search - Quoleady, accessed July 18, 2025,

<https://www.quoleady.com/llm-txt-files/>

24. GEO (Generative Engine Optimization) vs SEO: Full Overview - Zero Gravity Marketing, accessed July 18, 2025,

<https://zerogravitymarketing.com/blog/geo-vs-seo/>

25. What is Generative Engine Optimization - AI Monitor, accessed July 18, 2025,

<https://getaimonitor.com/what-is-generative-engine-optimization/>

26. What Is llms.txt? How the New AI Standard Works (2025 Guide) - Bluehost,

accessed July 18, 2025,
<https://www.bluehost.com/blog/what-is-llms-txt/>

27. llms.txt: how to get your brand mentioned in ChatGPT | AIO - diseño web en Barcelona, accessed July 18, 2025,
<https://www.latevaweb.com/en/what-is-llms-txt>

28. The Role and Functionality of llms.txt in LLM-Driven Web Interactions - Profound, accessed July 18, 2025, <https://www.tryprofound.com/guides/what-is-llms-txt-guide>

29. The value of llms.txt: Hype or real? - Mintlify, accessed July 18, 2025,
<https://mintlify.com/blog/the-value-of-llms-txt-hype-or-real>

30. AI search is booming, but SEO is still not dead, accessed July 18, 2025,
<https://searchengineland.com/ai-search-booming-seo-still-not-dead-458935>

31. How to Optimize Your Content for LLMs in 2025: A Complete Guide | Analyzify, accessed July 18, 2025,
<https://analyzify.com/hub/llm-optimization>

32. What is llms.txt File and What Does It Do? - zeo.org, accessed July 18, 2025,
<https://zeo.org/resources/blog/what-is-llms-txt-file-and-what-does-it-do>

33. What is llms.txt? How it works and examples - Mintlify, accessed July 18, 2025,
<https://mintlify.com/blog/what-is-llms-txt>

34. LLMs.txt: The Emerging Standard Reshaping AI-First Content Strategy | ScaleMath, accessed July 18, 2025,
<https://scalemath.com/blog/llms-txt/>

35. No, llms.txt is not the 'new meta keywords' - Search Engine Land, accessed July 18, 2025,
<https://searchengineland.com/no-llms-txt-is-not-the-new-meta-keywords-458199>

36. LLMS.txt Best Practices & Implementation Guide - Rankability, accessed July 18, 2025,

<https://www.rankability.com/guides/llms-txt-best-practices/>

37. Google Dismisses LLMS.txt as Ineffective and Unused by AI Bots - Stan Ventures, accessed July 18, 2025,

<https://www.stanventures.com/news/google-dismisses-llms-txt-as-ineffective-and-unused-by-ai-bots-2479/>

38. Avinash Tripathi, Co Founder - AI Monitor - Featured.com, accessed July 18, 2025,

<https://featured.com/p/avinash-tripathi>

39. LLMS.txt to help LLMs grok your content - Guillaume Laforge, accessed July 18, 2025,

<https://glaforge.dev/posts/2025/03/03/llms-txt-to-help-llms-grok-your-content/>

40. LLM Output Formats: Why JSON Costs More Than TSV | by David Gilbertson - Medium, accessed July 18, 2025,

<https://david-gilbertson.medium.com/llm-output-formats-why-json-costs-more-than-tsv-ebaf590bd541>