

Brand Erasure

A Research Agenda for Brand Theory in Agentic AI Markets

Marcos Guimarães Figueira

Fundação Getulio Vargas (FGV)

Rio de Janeiro, Brazil

Wyse Brand Intelligence

Rua do Russell, 804/4º andar, Rio de Janeiro, 22210-010 RJ, Brazil

`marcos.figueira@fgv.edu.br`

April 2026

Abstract

Brand theory was built for an era in which consumers found brands through search, evaluated them through marketing communications, and chose them through deliberate decision processes. That era is ending. AI assistants—generative chat interfaces, voice agents, and the autonomous systems that act on users’ behalf—are becoming the dominant intermediary between consumers and the products, services, and content that satisfy their needs. In this layer, the brand can be omitted from the answer, paraphrased into the synthesis, or substituted at the moment of action without the consumer ever noticing. This article advances *brand erasure* as a research construct and lays out an agenda for the work that needs to follow. It distinguishes erasure from adjacent constructs (dilution, disintermediation, algorithmic invisibility), argues that customer-based brand equity (Keller, 1993) requires extension into the latent-space substrate of foundation models, and proposes five research questions that map the empirical and theoretical work the field has not yet done. The agenda is written for marketing scholars, but its consequences extend to information retrieval, platform governance, and the regulation of attribution in AI systems.

Keywords: brand erasure; brand equity; generative AI; large language models; agentic AI; consumer–AI delegation; research agenda; brand theory.

1 Introduction

In March 2025, Pew Research Center documented a shift in user behavior that the marketing literature has not yet absorbed. Tracking 68,879 Google searches across 900 U.S. adults, the study found that when a search produced an AI Overview, users clicked a traditional search result 8% of the time, against 15% for searches without one. They clicked a link inside the AI summary 1% of the time. The summaries themselves cited three or more sources in 88% of cases (Pew Research Center, 2025). The numbers tell a single story. Attention no longer travels from the search engine to the brand’s website by default. It stops at the synthesized answer.

This is the milder version of the change. A second shift, less visible to most consumers but more consequential for brands, is already happening downstream. Tool-using language models (Schick et al., 2023), reasoning-and-acting agents (Yao et al., 2023), and computer-use APIs (Anthropic, 2024; OpenAI, 2025) are producing autonomous systems that complete tasks on the user’s behalf—booking, ordering, querying, configuring, transacting. The brand the agent transacts with is the brand the consumer is actually using, regardless of whether the consumer knew the brand existed. The unit of competition is no longer the click. It is the executed action.

Marketing theory has accumulated around the click. Brand awareness presumes the consumer encounters the brand. Brand consideration presumes a deliberate evaluation. Brand choice presumes a moment of decision in which the consumer’s preferences are operative. The AI-mediated layer breaks each presumption in turn. The consumer may not encounter the brand, may not consider it deliberately, and may not be the active party at the moment of choice. The intermediary is.

This article makes a simple argument with broad consequences. The construct that captures what is at stake in this transition—the systematic elimination of brand identity from synthesized answers and completed agentic actions—is not yet part of the brand-theory vocabulary. We propose to call it **brand erasure**, and we argue that establishing it as a formal research construct opens a research agenda the field is not currently equipped to address. The article distinguishes brand erasure from adjacent constructs, argues that the dominant brand-equity frameworks require extension into the latent-space substrate that foundation models operate in, and lays out five research questions intended to organize the empirical and theoretical work that should follow. The agenda is written for marketing scholars but takes seriously the disciplines whose work brand theory will need to draw on more deeply: information retrieval, computational linguistics, and the emerging literature on AI governance.

The argument unfolds in five sections after this introduction. Section 2 develops the construct of brand erasure and distinguishes it from related ideas. Section 3 argues that existing brand theory, while necessary, is insufficient to explain the new dynamics. Section 4 lays out the research agenda. Section 5 discusses methodological foundations. Section 6 concludes with the implications for marketing science as a discipline.

2 The Construct of Brand Erasure

We define brand erasure as follows.

Brand erasure is the condition under which an AI assistant satisfies a user’s underlying need without surfacing, citing, or transacting with the brand whose content, product, or service the response depends on.

The definition is deliberately operational. It does not require the assistant to actively misrepresent the brand or to harbor any disposition toward it. Erasure is what happens when the brand is *absent* from the user’s experience even though the brand is present in the substrate the assistant draws

on. The conditions are everyday and increasingly automatic. When a user asks a chatbot for a recommendation in a category and receives a synthesized answer that draws on multiple brand sources but names none of them, that is erasure in its first form. When the same user asks an agent to complete a transaction in the category and the agent selects a substitutable option without considering the brand at all, that is erasure in its terminal form.

2.1 Three Pathways

Brand erasure operates through three distinct pathways, each tied to a stage of the assistant’s procedure.

The first pathway is **extractive erasure**. A user issues a closed-domain query—a factual question—and the assistant lifts a fact from the brand’s content without identifying the source. The consumer receives the information but never encounters the brand. This is the smallest form of erasure and the one most continuous with classical disintermediation in the search literature.

The second pathway is **synthetic erasure**. A user issues an open-domain query and the assistant constructs a multi-source synthesis. Brand sources are present in the retrieval set, but their content is paraphrased into prose that omits attribution, partly because the underlying content was paraphrase-friendly in the first place. Liu, Zhang, and Liang (2023) audited four production generative search engines and found that only 51.5% of generated sentences were fully supported by their citations and only 74.5% of citations supported their associated sentence. The gap is not random. Distinctive content survives the synthesis pass with attribution intact; commodity content is recombined and reattributed.

The third pathway is **agentic erasure**, the most consequential. The user delegates a task to an autonomous agent, and the agent fulfills the category-level need by selecting a substitutable platform that may have no relation to the brand the user might have chosen if the user had been the active party. The brand is not delisted from a search results page. It simply never enters the action surface.

The three pathways correspond to the three optimization layers—Answer Engine Optimization, Generative Engine Optimization, and Agentic Optimization—discussed in our companion conceptual paper (Figueira, 2026). For the purposes of this agenda, the relevant point is that erasure has more than one mechanism, and any research program that addresses it must distinguish among them.

2.2 Distinction from Related Constructs

Brand erasure is adjacent to several established constructs in marketing and related disciplines, but it is not reducible to any of them.

Brand dilution (Aaker, 1996; Loken & John, 1993) describes the weakening of brand associations through poorly fitting extensions or off-brand communications. Erasure is the inverse problem: the brand’s associations may be intact in the consumer’s mind, but the consumer never reaches the moment of choice. Dilution is about what the consumer thinks of the brand. Erasure is about whether the consumer thinks of the brand at all.

Disintermediation (Sarkar, Butler, & Steinfield, 1995) describes the removal of intermediaries between producer and consumer, originally framed as a benefit of digital channels. Erasure is, in one sense, a re-intermediation problem: a new intermediary (the AI assistant) inserts itself between the brand and the consumer, with effects that are sometimes the *opposite* of classical disintermediation, since the assistant can hide the brand from view rather than connect the consumer to it directly.

Algorithmic invisibility (Caplan & boyd, 2018; Pasquale, 2015) describes the demotion or suppression of content by ranking systems. Erasure is related but distinguishable. Algorithmic invisibility is about whether content surfaces in a results list. Erasure is about whether the brand survives the *post-retrieval* synthesis or action stage. A brand can be visible to the retriever and still be erased by the synthesizer.

Brand neglect in algorithmic recommender systems (Chung, Wedel, & Rust, 2016) describes the systematic underweighting of brands by personalization algorithms. Erasure is a stronger condition: it includes neglect but extends to active *omission* during synthesis and *substitution* during action.

The distinctions matter because each adjacent construct has its own measurement tradition, intervention literature, and theoretical commitments. Treating erasure as merely the latest term for any of them imports the wrong assumptions and obscures what is genuinely new about the AI-mediated layer.

2.3 The Asymmetry Property

A property worth establishing up front, because it shapes the agenda that follows: brand erasure is *asymmetric*. Established brands with high entity-graph presence—canonical entries in public knowledge graphs, structured data coverage, stable named-entity recognition across foundation models—experience erasure at substantially lower rates than challenger brands whose distinctiveness is carried in copywriting and visual identity rather than in machine-readable structure. The legacy of brand-equity investment shows up in the AI substrate, but only if the investment was indexed. The asymmetry inverts the conventional brand-management orthodoxy in interesting ways and produces several of the research questions we propose below.

3 Why Existing Brand Theory Is Insufficient

The argument of this section is not that existing brand theory is wrong. It is that the dominant frameworks were built to explain phenomena in a substrate—human associative memory—that is no longer the only substrate that matters for brand outcomes.

3.1 The Customer-Based Brand Equity Inheritance

Keller’s (1993) Customer-Based Brand Equity (CBBE) framework defined brand knowledge as a network of associations in consumer memory, organized along dimensions of awareness and image, and varying in favorability, strength, and uniqueness. The framework has been the most generative

idea in brand theory of the last three decades. It produced measurement instruments (Keller, 2001; Yoo & Donthu, 2001), informed the brand-asset valuation work of Aaker (1991, 1996) and Ailawadi, Lehmann, and Neslin (2003), and sustained an empirical literature on association networks (Krishnan, 1996; John, Loken, Kim, & Monga, 2006).

The framework’s foundational assumption is that brand outcomes are mediated through consumer cognition: the consumer encounters the brand, retrieves associations, evaluates options, and chooses. The mediation chain runs through human memory.

In AI-mediated markets, an additional mediation layer exists. The assistant retrieves, evaluates, and chooses on the user’s behalf, drawing on a different substrate—the latent representations of foundation models, the entity graphs that ground retrieval, the structured data in machine-readable manifests. The consumer’s memory matters less if the assistant’s representation matters more. CBBE theory has nothing to say about the second substrate, because the second substrate did not exist when the theory was developed.

3.2 The Brand-as-Vector Reframing

The translation from CBBE to the AI substrate is, at the level of structure, surprisingly direct. A brand association in consumer memory has a node, a connection, a weight. A brand association in the latent space of a language model has a vector position, a distance to attribute vectors, a magnitude. The mathematical structure is the same—a graph, plus weights, plus a similarity measure. What changes is whose graph it is.

The implication is twofold. First, brand-equity theory does not need to be discarded, but it does need to be extended. The dimensions Keller specified—awareness, favorability, uniqueness—have direct analogs in vector space. Awareness corresponds to whether the brand has a stable, well-localized vector representation. Favorability corresponds to the proximity of the brand vector to positively valenced attribute vectors. Uniqueness corresponds to the distance from competitor vectors in dimensions that matter for the category. Each of these is empirically tractable using methods from the embedding-bias literature (Bolukbasi et al., 2016; Caliskan, Bryson, & Narayanan, 2017).

Second, the new substrate has its own dynamics that do not reduce to consumer cognition. The latent representation of a brand is shaped by what the model was trained on (corpus presence), how the brand is connected to category attributes in that training corpus (definitional anchoring), and how the model’s reranker and synthesizer behave at inference time (computational utility). None of these are properties of consumer memory. They are properties of the AI system. Brand theory has to engage with them on their own terms.

3.3 What Brand Theory Has to Absorb

To move forward, brand research will have to absorb conceptual and methodological vocabulary from three adjacent disciplines. From information retrieval: the structure of dense retrieval, late-interaction reranking (Khattab & Zaharia, 2020), retrieval-augmented generation (Lewis et al.,

2020; Gao et al., 2024), and the citation-behavior literature on generative search (Liu et al., 2023; Menick et al., 2022). From computational linguistics: word- and sentence-embedding methods, the bias-probe tradition, and the more recent work on attributed question answering. From the emerging literature on agentic AI: tool-using models, large-action models, and the early evaluations of agent reliability and substitutability.

The absorption is not optional. A brand-theory paper that tries to address the AI-mediated layer without engaging the technical literature reads, to readers in adjacent disciplines, as a paper that does not understand its own object. The discipline cost of staying inside the marketing canon will be exclusion from the broader conversation about how AI systems shape consumer behavior—a conversation that is currently being held primarily in computer-science venues with limited brand-theory input.

4 The Research Agenda

We propose five research questions, each grounded in the construct of brand erasure and each designed to be tractable with methods that already exist in the marketing or adjacent literatures.

4.1 RQ1: How Should Brand Erasure Be Measured?

The first task is measurement. Brand erasure is currently a verbal construct without a standard instrument. The marketing literature has not yet converged on a measurement protocol that is reliable, valid, and comparable across studies.

Two measurement traditions can be combined. The first is the **embedding probe**, adapted from word-embedding bias research. For a target brand, attribute vectors are defined for category-relevant dimensions, and cosine proximity between the brand vector and each attribute vector is computed across multiple foundation models. The output is a matrix that quantifies the brand’s positioning in each model’s latent space and that can be tracked longitudinally. The second is the **response audit**, in which a representative set of category-relevant prompts is submitted to a population of AI assistants and the responses are coded for brand mention frequency, citation, position, sentiment, and substitutability. The aggregate yields share-of-citation and erasure-rate measures.

The agenda items here are concrete. Standardizing the attribute sets used in embedding probes. Developing inter-rater reliability for response-audit coding schemes. Establishing convergent validity between latent-space measures and survey-based brand-equity instruments (Aaker, 1996; Yoo & Donthu, 2001). Developing change-point detection for longitudinal designs that account for the non-stationarity of foundation models. None of these tasks is particularly novel methodologically, but none has been done specifically for brand outcomes.

4.2 RQ2: What Predicts Erasure Asymmetry?

If erasure is asymmetric—if some brands are erased more than others, controlling for category and corpus presence—then the determinants of asymmetry are a research target with both theoretical

and managerial implications.

Several candidate predictors are obvious from the framework. Entity-graph presence: brands with canonical Wikipedia entries, Wikidata identifiers, and stable named-entity recognition should experience lower erasure rates. Definitional anchoring: brands that introduced or operationalized terminology that became load-bearing in the category should be cited at rates disproportionate to their corpus share. Information gain: content that is paraphrase-resistant, in the sense that it cannot be cheaply synthesized from other sources in the retrieval set, should survive the synthesis pass with attribution intact.

Less obvious candidates also deserve study. Brand age and category dominance—old, dominant brands may benefit from training-corpus saturation in ways that newer brands cannot match through any current intervention. Brand-name distinctiveness—generic-sounding brand names may be at structural disadvantage in retrieval and reranking, an effect that is mathematically predictable from how embedding models tokenize and compare strings. Visual identity—the role of brand logos in multimodal models is currently underexplored.

The empirical work here is straightforward in design (cross-sectional regression of erasure rates on candidate predictors, controlling for category and corpus presence) and challenging in execution (each candidate predictor needs a defensible operationalization).

4.3 RQ3: How Do Erasure Dynamics Propagate?

Erasure is not a stock; it is a flow. A brand may be present in one assistant’s responses today and absent next month, with the change driven by model updates, reranker shifts, or compounding paraphrase chains in which downstream content reuses upstream syntheses without their original attributions. The dynamics—half-life, recoverability, hysteresis—are an open empirical question with substantial managerial consequences.

Three sub-questions deserve formal investigation. First, what is the typical half-life of brand attribution under continuous model updating, controlling for the brand’s intervention rate? Second, how does paraphrase compounding propagate through multi-hop content chains, and at what point in the chain is original attribution typically lost? Third, are there *recoverable* and *unrecoverable* states—points beyond which a brand’s representation in foundation-model substrates cannot be repaired through any practical intervention?

The dynamics question requires longitudinal designs that the field has not yet committed to. A precedent exists in the digital-shelf and search-rank literature (Reisenbichler, Reutterer, Schweidel, & Dan, 2022), which has begun to track content-marketing effects at scale.

4.4 RQ4: What Is the Relationship Between Latent-Space Representation and Behavioral Outcomes?

A skeptical reader might accept that brand vectors can be measured, that erasure rates can be audited, and that asymmetry can be predicted, while still asking the harder question: does any of it matter for actual consumer behavior?

The connecting variable is **intermediary behavior**. Latent-space representation shapes how AI assistants surface and present brands; assistant behavior shapes consumer exposure, evaluation, and choice. Each link in the chain is amenable to measurement. The composite effect—from latent-space position to ultimate purchase—is what the literature must establish empirically before brand-erasure measurement can be claimed to predict outcomes that managers care about.

Three designs are within reach. First, observational studies that combine brand-erasure audits with category-level revenue or session data, looking for correlations between erasure rate changes and brand-revenue changes over a long enough window to detect lag effects. Second, controlled experiments on simulated consumer–AI delegation tasks, in which the brand’s representation is manipulated experimentally (by, for instance, providing the assistant with different retrieval contexts) and consumer choice is measured. Third, panel studies that track real users’ interactions with AI assistants over time, modeled on the Pew (2025) methodology, extended from search-with-AI-summary to agentic completion of category tasks.

The work is non-trivial—particularly the third design, which raises ethics-board questions about real-user observation that the field has not yet fully resolved. But the questions are tractable, and answers would dramatically tighten the inferential chain from brand-erasure measurement to managerial implications.

4.5 RQ5: How Should Attribution Be Governed in AI-Mediated Markets?

The fifth research question concerns governance. The brand-erasure problem is not only a marketing problem; it is a problem about how attribution should work in systems that synthesize content and execute actions across multiple sources. Marketing scholars have a stake in this conversation because brand outcomes are at issue, but the conversation is broader than marketing.

The EU AI Act (Regulation 2024/1689) created obligations on general-purpose AI systems that include training-data summary requirements (Art. 53) and transparency provisions about model inputs (European Parliament & Council, 2024). Whether these provisions ultimately produce *enforceable* attribution rights—what trade press has begun to call a “right to be cited,” though no such doctrine yet exists in mature form—is a research question with empirical, doctrinal, and normative components. The U.S. patchwork (state-level AI laws, FTC enforcement, sector-specific rules), the U.K. AI Safety Institute’s guidance approach, and the regulatory experiments in Brazil’s CGI.br and ANPD all create a comparative landscape ripe for cross-jurisdictional analysis.

Substantive normative questions accompany the empirical ones. Should attribution be a default obligation of AI assistants, an opt-in for brands that mark content with structured signals, or a market outcome left to private contracting? What welfare consequences follow from each regime, particularly for the long-tail of brands that lack the resources to structure their content for machine readability? Hacker, Engel, and Mauer (2023) and Veale and Zuiderveen Borgesius (2021) have begun mapping the AI Act’s architecture; comparable work that takes brand outcomes specifically as a focal interest has not been done.

5 Methodological Foundations

The five research questions above can be addressed using methods that exist in adjacent disciplines and can be adapted, with care, to brand outcomes. Three deserve explicit mention as foundational.

Embedding probes measure brand position in the latent space of foundation models. The technique is well-established for bias research (Bolukbasi et al., 2016; Caliskan et al., 2017) and applies to brands without conceptual modification, though it requires careful operationalization of attribute sets, sampling protocols across models, and longitudinal tracking. The method’s central limitation is that latent-space proximity is a proxy for retrieval prominence, not a substitute. Triangulation against actual response data is necessary.

Response audits sample AI assistant outputs across category-relevant prompts and code them for brand-attribution metrics. The methodology is closer to algorithmic-accountability work (Diakopoulos, 2016) than to traditional marketing measurement, but the codebook can be made interoperable with brand-equity instruments. Sampling design matters: foundation models are stochastic, model updates produce drift, and personalization affects responses in ways that confound naive single-sample audits. A defensible design samples 3–5 responses per prompt-model cell within a short window and tracks across models and time.

Mixed-methods integration combines latent-space and response-level measurement with traditional survey-based brand-equity instruments to assess convergent and discriminant validity. The integration is the part of the agenda where the field can make distinctive contributions: the survey-based instruments are marketing’s home territory, and triangulation with the AI-substrate measures is the path to a brand-equity construct that captures both consumer cognition and intermediary representation in a coherent way.

The Pew Research Center’s (2025) panel-based audit of AI-summary effects on user click behavior provides a methodological model for an additional family of designs: observational studies of real users’ interactions with AI assistants. Extending the design from search-with-AI to agentic completion is a natural next step that has not yet been taken.

6 Conclusion: Implications for Marketing Science

The construct of brand erasure does not require marketing science to abandon the frameworks the field has developed. It requires the field to extend them into a substrate the frameworks were not built to address. The unit of analysis remains the brand. What changes is the locus of mediation.

The discipline-level move that follows from this article is that brand researchers will need to become bilingual. The native vocabulary of the marketing literature—equity, awareness, image, association—is necessary but not sufficient. The vocabulary of information retrieval, computational linguistics, and AI governance—retrieval, reranking, embeddings, agentic action, attribution—is increasingly load-bearing for explaining brand outcomes. Bilingual researchers will be in short supply for some time, and the field’s productivity in this area will partly reflect its willingness to invest in cross-training, joint appointments, and venues that publish work straddling the disciplines.

A second implication concerns industry-academy relationships. The empirical infrastructure required to study brand erasure at scale—API access to foundation models, cross-platform audit pipelines, longitudinal data on agent behavior—is currently held primarily in industry. The classical model of academic research, in which scholars develop methods and industry applies them, is a poor fit for a domain where industry already holds the data the methods need. Productive arrangements will require explicit data-sharing partnerships, replication-friendly preregistration protocols, and careful management of confidentiality without surrendering the field’s commitment to open science.

A final implication is temporal. The substrate that AI assistants operate in is malleable now, more than it will be in five or ten years. Foundation-model training corpora are still being constituted; entity graphs are still being built; the conventions of attribution that will harden into industry norms are still being negotiated. Brand-theory work that addresses this period will have a chance to influence the conventions that emerge. Brand-theory work that waits for the conventions to settle will arrive after the period in which it could have shaped them.

The agenda we have proposed is an invitation. Brand erasure is not the only construct that will need to be developed for the AI-mediated era, and the five research questions are not exhaustive. But the construct names something real, the questions are tractable, and the methods exist. The work has to start somewhere. We argue it should start here, and soon.

References

- Aaker, D. A. (1991). *Managing brand equity: Capitalizing on the value of a brand name*. Free Press.
- Aaker, D. A. (1996). Measuring brand equity across products and markets. *California Management Review*, 38(3), 102–120.
- Ailawadi, K. L., Lehmann, D. R., & Neslin, S. A. (2003). Revenue premium as an outcome measure of brand equity. *Journal of Marketing*, 67(4), 1–17.
- Anthropic. (2024, October). *Introducing computer use, a new Claude 3.5 Sonnet, and Claude 3.5 Haiku* [Company announcement].
- Bolukbasi, T., Chang, K.-W., Zou, J. Y., Saligrama, V., & Kalai, A. T. (2016). Man is to computer programmer as woman is to homemaker? Debiasing word embeddings. *Advances in Neural Information Processing Systems*, 29.
- Caliskan, A., Bryson, J. J., & Narayanan, A. (2017). Semantics derived automatically from language corpora contain human-like biases. *Science*, 356(6334), 183–186.
- Caplan, R., & boyd, d. (2018). Isomorphism through algorithms: Institutional dependencies in the case of Facebook. *Big Data & Society*, 5(1), 1–12.
- Chung, T. S., Wedel, M., & Rust, R. T. (2016). Adaptive personalization using social networks. *Journal of the Academy of Marketing Science*, 44(1), 66–87.

- Diakopoulos, N. (2016). Accountability in algorithmic decision making. *Communications of the ACM*, 59(2), 56–62.
- European Parliament and Council. (2024). *Regulation (EU) 2024/1689 of the European Parliament and of the Council laying down harmonised rules on artificial intelligence (Artificial Intelligence Act)*. Official Journal of the European Union, L 2024/1689, 12 July 2024.
- Figueira, M. G. (2026). *From information retrieval to agentic action: A framework for brand visibility in AI-mediated markets*. Working paper, Fundação Getulio Vargas (FGV) and Wyse Brand Intelligence.
- Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., Dai, Y., Sun, J., & Wang, H. (2024). *Retrieval-augmented generation for large language models: A survey* [arXiv:2312.10997]. arXiv.
- Hacker, P., Engel, A., & Mauer, M. (2023). Regulating ChatGPT and other large generative AI models. In *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT '23)*.
- John, D. R., Loken, B., Kim, K., & Monga, A. B. (2006). Brand concept maps: A methodology for identifying brand association networks. *Journal of Marketing Research*, 43(4), 549–563.
- Keller, K. L. (1993). Conceptualizing, measuring, and managing customer-based brand equity. *Journal of Marketing*, 57(1), 1–22.
- Keller, K. L. (2001). Building customer-based brand equity: A blueprint for creating strong brands. *Marketing Management*, 10(2), 14–19.
- Khattab, O., & Zaharia, M. (2020). ColBERT: Efficient and effective passage search via contextualized late interaction over BERT. In *Proceedings of SIGIR 2020*.
- Krishnan, H. S. (1996). Characteristics of memory associations: A consumer-based brand equity perspective. *International Journal of Research in Marketing*, 13(4), 389–405.
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474.
- Liu, N. F., Zhang, T., & Liang, P. (2023). Evaluating verifiability in generative search engines. In *Findings of the Association for Computational Linguistics: EMNLP 2023* (pp. 7001–7025).
- Loken, B., & John, D. R. (1993). Diluting brand beliefs: When do brand extensions have a negative impact? *Journal of Marketing*, 57(3), 71–84.
- Menick, J., Trebacz, M., Mikulik, V., Aslanides, J., Song, F., Chadwick, M., Glaese, M., Young, S., Campbell-Gillingham, L., Irving, G., & McAleese, N. (2022). *Teaching language models to support answers with verified quotes* [arXiv:2203.11147]. arXiv.

- OpenAI. (2025, January). *Introducing Operator* [Company announcement].
- Pasquale, F. (2015). *The black box society: The secret algorithms that control money and information*. Harvard University Press.
- Pew Research Center. (2025, July 22). *Google users are less likely to click on links when an AI summary appears in the results*.
- Reisenbichler, M., Reutterer, T., Schweidel, D. A., & Dan, D. (2022). Frontiers: Supporting content marketing with natural language generation. *Marketing Science*, 41(3), 441–452.
- Sarkar, M. B., Butler, B., & Steinfield, C. (1995). Intermediaries and cybermediaries: A continuing role for mediating players in the electronic marketplace. *Journal of Computer-Mediated Communication*, 1(3).
- Schick, T., Dwivedi-Yu, J., Dessì, R., Raileanu, R., Lomeli, M., Zettlemoyer, L., Cancedda, N., & Scialom, T. (2023). Toolformer: Language models can teach themselves to use tools. *Advances in Neural Information Processing Systems*, 36.
- Veale, M., & Zuiderveen Borgesius, F. (2021). Demystifying the draft EU Artificial Intelligence Act. *Computer Law Review International*, 22(4), 97–112.
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Yoo, B., & Donthu, N. (2001). Developing and validating a multidimensional consumer-based brand equity scale. *Journal of Business Research*, 52(1), 1–14.