

NiemanLab

 Adobe Stock open book with sticky notes 2

Generative AI models love to cite Reuters and Axios, study finds

Muck Rack found that major models cite journalism in nearly half of responses that require “recency.”

By **ANDREW DECK** July 24, 2025, 3:28 p.m.

Journalistic content is integral to answers from generative AI tools, a [new report from Generative Pulse by Muck Rack](#) finds.

The authors of the report, published this week, analyzed more than one million citations output by generative AI models. Journalistic content was cited more than 27% of the time across all Muck Rack’s tests, and when it came to queries that implied “a level of recency” (including not just breaking news but timely responses to prompts like “car rental shortages in the U.S.” or “latest advancements in outpatient treatment methods for substance abuse”) that number jumped to 49%, or nearly half of all the links cited.

Muck Rack’s tests were conducted this month and included GPT’s 4o and 4o Mini models, Gemini’s Flash and Pro models, and Claude’s Sonnet and Haiku models.

When ChatGPT and similar generative AI products craft responses, they can only rely on training data up to a point. Muck Rack’s tests confirmed that if you disable citations — and subsequently, the ability to scrape the web in real-time — for these models, they often output inaccurate or out-of-date information.

Muck Rack also found that the type of question asked changed the citation sources significantly. Subjective questions, such as queries asking for advice or step-by-step instructions, pulled far more from “corporate blogs and content.”

RELATED ARTICLE

AI search engines fail to produce accurate citations in over 60% of tests, according to new Tow Center study

March 10, 2025



Journalistic content was cited in 49% of “queries implying a level of recency,” followed by corporate blogs and government/NGO sources.

Muck Rack also found some interesting differences between models.

Overall, Claude appeared to cite journalistic sources the least often. Take Reuters as an example: The news source was cited by Claude 20 times less often than Gemini, and 50 times less often than ChatGPT.

The researchers also noted the top six media outlets cited by each model. Reuters, the Financial Times, Time, Forbes, and Axios all appeared in both the ChatGPT and Gemini lists. (It’s worth noting that [the Financial Times](#), [Time](#), and [Axios](#) all have AI licensing deals with OpenAI, but not with Google.) Again, Claude had less of a hard-news skew in its citations, with outlets like Good Housekeeping, TechRadar, and Harvard Business Review making its top-six list.



The Financial Times, Reuters, Time, Forbes, and Axios were among the most cited by ChatGPT and Gemini.

When it came to citations of journalistic content, the models seemed to favor stories that had been published in the last 12 months. That was especially true of ChatGPT, which saw 56% of its journalism citations published within the past year. For Claude, that number only came out to 36% of all citations.

Muck Rack also broke down queries by industry, including healthcare, energy, and retail. For the media/entertainment, finance/insurance, and government industries, journalistic content was more likely to be cited than in other industries. That is in comparison to corporate blogs, paid media, academic articles, press releases, and other categories.

RELATED ARTICLE

ChatGPT is hallucinating fake links to its news partners' biggest investigations

June 27, 2024

The quality of citations in generative AI products have been a constant point of contention between AI developers and news publishers — particularly as generative AI tools take a [slice out of traditional search traffic](#). ChatGPT and similar products have been found to [hallucinate URLs](#) to news publisher sites, even those with OpenAI licensing deals. They have also been found to [link out to syndications and unauthorized copies](#) of articles instead of the original stories. The frequency of journalistic citations, by the numbers, has been far less documented.

HEADER IMAGE BY [TL_TO_TITO](#) USED UNDER LICENSING VIA ADOBE STOCK.

Andrew Deck is a staff writer covering AI at Nieman Lab. Have tips about how AI is being used in your newsroom? You can reach Andrew via email, Bluesky, or Signal (+1 203-841-6241).