

Association for Information Systems

AIS Electronic Library (AISeL)

PACIS 2026 Proceedings

Pacific Asia Conference on Information
Systems (PACIS)

July 2026

How Should Generative AI Cite Sources? An Empirical Study on Citation Formats and Perceived Transparency

Lufei Liang

Wuhan University, 2025201040003@whu.edu.cn

Bo Yang

Wuhan University, yangboo@whu.edu.cn

Yongqiang Sun

Wuhan University, sunyq@whu.edu.cn

Follow this and additional works at: <https://aisel.aisnet.org/pacis2026>

Recommended Citation

Liang, Lufei; Yang, Bo; and Sun, Yongqiang, "How Should Generative AI Cite Sources? An Empirical Study on Citation Formats and Perceived Transparency" (2026). *PACIS 2026 Proceedings*. 8.

https://aisel.aisnet.org/pacis2026/hci_robotic/hci_robotic/8

This material is brought to you by the Pacific Asia Conference on Information Systems (PACIS) at AIS Electronic Library (AISeL). It has been accepted for inclusion in PACIS 2026 Proceedings by an authorized administrator of AIS Electronic Library (AISeL). For more information, please contact elibrary@aisnet.org.

How Should Generative AI Cite Sources? An Empirical Study on Citation Formats and Perceived Transparency

Completed Research Paper

Lufei Liang

School of Information Management
Wuhan University
299 Bayi Road, Wuhan, Hubei, China
2025201040003@whu.edu.cn

Bo Yang

School of Information Management
Wuhan University
299 Bayi Road, Wuhan, Hubei, China
yangboo@whu.edu.cn

Yongqiang Sun

School of Information Management
Wuhan University
299 Bayi Road, Wuhan, Hubei, China
sunyg@whu.edu.cn

Abstract

Transparency is an important principle for generative AI interface design. To enhance transparency, current generative AI applications have begun to embed citations in generated content, which allows users to access the information sources of the generated content. In practice, three common citation formats are used: list, in-text, and their combined format. However, limited research has examined the effects of different citation formats. To address this gap, this study investigates the impact of citation formats on perceived transparency and the moderating role of task goal specificity. We developed generative AI systems and conducted an online experiment. The results show that both the list and in-text formats increase perceived transparency. Their combined format produces the highest level of perceived transparency, but the additional benefit is limited, suggesting diminishing returns. Furthermore, the list format is more effective in low goal specificity tasks.

Keywords: Generative AI, citation format, information presentation, perceived transparency, interface design

Introduction

With the widespread adoption of Generative AI, hallucinations have become an increasingly pressing concern (Huang et al., 2025). Research has found that in certain high-stakes scenarios (e.g., legal question answering), the hallucination rate can reach 69%-88% (Ho, 2024). However, as users lack knowledge about information sources and generation process of the generated content, they often struggle to identify hallucinations, which thereby exacerbate the spread of misinformation (Dwivedi et al., 2023; Huang et al., 2025). In this context, enhancing the transparency of generative AI, such as disclosing the information sources underlying the generated content, has emerged as an important means of helping users detect hallucinations and make more informed judgments. Consequently, enhancing transparency in generative AI has become a key design principle and an urgent practical necessity (European Union, 2024).

To enhance transparency, an increasing number of Generative AI applications (e.g., ChatGPT, Gemini and Perplexity AI) insert citations in generated content. These citations are embedded hyperlinks that link reference sources. Once users click these hyperlinks, they can trace the supporting information of generated content (Ding, Poudel, et al., 2025; Liu et al., 2023). In current practice, citations are presented in three formats: the basic *list* and *in-text* formats, as well as their *combined format* (Ouyang & Narechania, 2026). These formats differ in their mode of presentation: the list format compiles all references at the end of the content, with metadata (e.g., titles) displayed for each item; while the in-text format places citations directly at the end of the sentences they support; and the combined format presents both a reference list and in-text citations simultaneously. Additionally, these formats also convey different aspects of the same set of references: the list format presents the content cues of cited sources by metadata (i.e., what is cited), while the in-text format reveals the correspondence between citations and generated content (i.e., how it is cited). The combined format integrates both types of information, providing more comprehensive information. Prior research suggests that presentation format is a critical interface design element that shapes how users perceived and process information (Hong et al., 2004; Jiang & Benbasat, 2007). Given that these formats differ in both their mode of presentation and the type of information conveyed, their effectiveness in enhancing users' perceived transparency may vary. However, although these three formats are widely adopted in current applications, existing research has primarily focused on the presence and quality of citations (Ding et al., 2025; Kim et al., 2025), with limited attention paid to how they should be presented. It remains unclear which format is more effective and how citation presentation should be designed to optimize users' transparency perceptions. This gap results in a lack of theoretical guidance for citation presentation design, limiting the potential of citations to enhance transparency in Generative AI interfaces. To fill this gap, this study aims to examine the effects of different citation formats (list, in-text, and their combination) on perceived transparency.

We focus on two research question. First, as two fundamental formats of citation presentation, can both list and in-text format effectively present citations to enhance users' perceived transparency? More importantly, when used in combination, do they lead to the optimal transparency by providing richer information related to citations (Bashirzadeh et al., 2022; Jiang & Benbasat, 2007; Lu et al., 2020), or do they yield a substitutive effect due to their function and information redundancy (Kizilcec, 2016; Ling et al., 2024; Ning et al., 2024)? Accordingly, we propose our first research question: *How do the list and in-text formats influence perceived transparency individually and jointly?*

Furthermore, prior research indicates that the effects of presentation formats depend on task characteristics (Hong et al., 2004; Lim & Benbasat, 2000; Vessey & Galletta, 1991). When users utilize Generative AI for information seeking, tasks vary in goal specificity (Athukorala et al., 2016; Hong, 2006; Tam & Ho, 2006), ranging from high goal specificity tasks that involve searching for specific information, to low goal specificity tasks aimed at broadly gathering information about a topic. Under tasks with high and low goal specificity, do the effects of citation formats alter? Accordingly, we propose the second research question: *How does task goal specificity moderate the effects of citation formats on perceived transparency?*

To address RQ1, this study draws on the literature on information presentation (Deng & Poole, 2010; Gedikli et al., 2014; Jiang & Benbasat, 2007; Lim & Benbasat, 2000; Lu et al., 2020). Research on information presentation suggests that formats of information presentation influence perceived transparency due to different types and amount of information provided (Lu et al., 2020). When presented information is complementary, multiple information cues can strengthen the overall effect (Jiang & Benbasat, 2007). In contrast, redundant information may induce information saturation or overload, weakening its positive impact (Ling et al., 2024). We further draw on Task-technology fit (TTF) theory to address RQ2 (Goodhue & Thompson, 1995). This theory posits that users' perceptions are more positive when presentation formats match task characteristics (Vessey, 1994). TTF thus serves as a complementary theoretical lens, highlighting the boundary condition of task characteristics in explaining when certain presentation formats are more effective.

We conducted an online experiment to investigate the two research questions using Generative AI systems based on the Retrieval Augmented Generation framework (Lewis et al., 2020). These systems differ in citation formats and support real-time, multi-turn interactions between participants and Generative AI. Unlike prior studies relying on questionnaires (e.g., Zhou et al., 2025) or wizard-of-oz approaches (e.g., Shin et al., 2024), our design closely simulates real user interaction, which enhances ecological validity. Our

findings contribute to the literature on citation presentation and AI transparency design by advancing theoretical understanding of the effects of citation formats on perceived transparency. They also offer practical design guidelines for enhancing transparency in Generative AI.

Theoretical Background

Citation Format of Generative AI

With the proliferation of Generative AI, the "black box" issue has emerged as a critical challenge (OECD, 2025). To enhance system transparency, it has become a prevalent practice for Generative AI applications (e.g., ChatGPT, Gemini, Perplexity AI) to embed citations in generated content. These citations typically appear as hyperlinks linking to specific sources (Liu et al., 2023). By following these hyperlinks, users can access and verify the original sources of generated content.

Currently, Generative AI employs two fundamental citation formats: list and in-text. The list format consolidates all reference links as a list at the end of the generated content, with source titles displayed (Figure 1a), while the in-text format inserts references directly after relevant sentences, typically represented by superscript numeric markers (Figure 1b). The two basic formats present the same set of references but differ in both their mode of presentation and the type of information they convey. Specifically, the list format displays cited content through titles and other metadata (i.e., *what* is cited), but due to its separation from the generated content, it fails to clearly indicate which citation supports which sentence (i.e., *how* it is cited). In contrast, the in-text format embeds citations directly at the end of sentences, clearly conveying the correspondence between citations and the generated content. However, because these markers typically appear as numeric superscripts, they cannot adequately convey information about the content of the cited sources.




Figure 1. Different Citation Formats

Further, some applications begin to combine the list and in-text formats (i.e., combined format). Though the three citation formats are widely used in practice, existing research primarily focused on the presence, quantity and relevance of citations (Ding et al., 2025; Kim et al., 2025; Lenz et al., 2025), with limited attention paid to the impact of citation format. Consequently, it remains unclear which is the optimal format to present citations. To fill this gap, the present study attempts to investigate the effects of citation format (list, in-text, and their combination), providing actionable guidelines for designing transparent Generative AI interfaces.

Transparency Design and Information Presentation

Perceived transparency generally refers to the extent to which users can understand the system (Al-Natour et al., 2021; Gedikli et al., 2014). Prior research has demonstrated that information disclosure is an effective approach to enhance transparency, such as by providing explanations or data visualizations (Al-Natour et al., 2021; Ning et al., 2024; You et al., 2022). Through disclosing system-related information, users can gain insight into the system's information sources, operational processes, and underlying rationale, which helps foster a sense that the system is not concealing relevant information and thereby improves their perceived transparency. Moreover, research has found that the effectiveness of transparency enhancement is

influenced by information presentation (Gedikli et al., 2014; Kizilcec, 2016; You et al., 2022). This is mainly because different presentation formats vary in the type, amount, and detail of the information they convey. For example, compared to static format, interactive formats are capable of presenting large amounts of complex information, thereby facilitating deeper user comprehension of the system and enhance users' perceived transparency (Cheng et al., 2019).

However, existing studies also suggests that the relationship between information presentation and perceived transparency is not straightforward, as the amount and type of information disclosed do not always produce consistent effects. On the one hand, a *complementary effect* may occur, where diverse types of information reinforce one another and enhance the overall effectiveness (Jiang & Benbasat, 2007). For instance, Lu et al. (2020) found that presenting local and global explanations together contributes to an improvement compared to only presenting that in one format. This can be explained by the cue-summation theory (Severin, 1967), which posits that the number of relevant cues are summated to facilitate understanding. On the other hand, a *substitutive effect* may arise when different types of information are redundant (Bian et al., 2025). For example, when information about an AI's process is provided, additional information about its purpose does not further enhance (Ning et al., 2024). This is consistent with the redundancy effect in cognitive load theory (Mayer et al., 2001). When presenting the redundant information in multiple formats simultaneously, it forces the brain to process identical information twice, which creates a high cognitive load and thus reduces comprehension.

Applying these theoretical insights to the context of Generative AI, the list and in-text formats represent two different ways of presenting citation information. While both formats disclose the same citations, they emphasize different aspects of information. The list format conveys information about what sources are cited, whereas the in-text format indicates how these citations support the generated content. Consequently, it remains elusive how these formats interact when combined. Specifically, it is unclear whether their combined presentation maximizes users' perceived transparency through information complementarity (Jiang & Benbasat, 2007), or weaken transparency due to information redundancy (Bian et al., 2025). Thus, this study focuses on unpacking the nuanced interplay between list and in-text formats when used in combination.

Task Characteristic: Goal Specificity

Literature on information presentation suggests that the effectiveness of presentation formats is impacted by task characteristics (Lim & Benbasat, 2000; Vessey & Galletta, 1991). In the context of Generative AI, users increasingly use Generative AI as a primary tool for information seeking (Sharma et al., 2024). Prior research classifies search tasks into two types according to their goal specificity (Athukorala et al., 2016; Hong, 2006; Tam & Ho, 2006). Goal specificity refers to how specifically the task goal is defined (Hong, 2006). High goal specificity tasks (or goal-oriented searches) feature clearly defined objectives, such as locating a specific fact (Athukorala et al., 2016), while low goal specificity tasks (or exploratory searches) are characterized by open-ended goals, such as gathering broad information on a topic (Hong, 2006). Prior research has shown that task goal specificity interacts with presentation formats to shape users' perceptions. For example, complex formats can reduce satisfaction in high goal specificity tasks (Nadkarni & Gupta, 2007). However, whether the effectiveness of different citation formats is contingent upon task goal specificity remains underexplored. Addressing this question can inform more adaptive citation designs that better align with specific task contexts. Thus, this study further examines the moderating role of task goal specificity.

Task-Technology Fit Theory

Task-technology fit theory (TTF) (Goodhue & Thompson, 1995) provides a theoretical foundation for understanding how effects of citation formats vary depending on goal specificity. TTF highlights that when technology characteristics (e.g., presentation format) match task characteristics, information processing can be facilitated, thereby contributing to positive outcomes (Vessey & Galletta, 1991; Ziguers & Buckland, 1998). For example, when the list format is provided for the browsing task and the matrix format is provided for the searching task (indicating a high format-task fit), users have a more positive attitude towards the product website (Hong et al., 2004). In this study context, citation format can be conceptualized as a technology characteristic of Generative AI. How well these formats "fit" the task goal specificity may influence users' perceived transparency.

Research Model and Hypotheses

Drawing on literature of information presentation, we first investigated the individual and interaction effect of different citation formats (i.e., list and in-text). Further, we examined the moderating effect of task goal specificity based on Task-Technology Fit theory. The research model is shown in Figure 2.




Figure 2. Research Model

The Effect of Citation Format on Perceived Transparency

In this study, perceived transparency refers to the extent to which users can understand the information sources of generated content (Al-Natour et al., 2021; Gedikli et al., 2014). Prior research indicates that information disclosure reduces users knowledge gap and enhances perceived transparency (Al-Natour et al., 2021; Ning et al., 2024). Since both citation formats disclose source information in different ways, both are expected to enhance perceived transparency relative to no citations. Therefore, we hypothesize that:

H1. Compared to the absence of citations, both list format and in-text format citations enhance users' perceived transparency.

Further, we evaluate the combined effect of the two citation formats through perspective of information provision and functionality. From the perspective of information provision, the list and in-text formats each have strengths and limitations. The list format provides information about citation content, but lacks explicit correspondence with the generated text. Conversely, the in-text format shows how citations support the text but does not display the cited content. When combined, they provide the most comprehensive information cues. Users can understand what sources are cited through the list format, and how those sources correspond with specific parts of the generated content through the in-text format. According to cue-summation theory (Jiang & Benbasat, 2007), the integration of complementary information cues enhances users' understanding. In other words, the combined format outperforms each individual format in improving users' perceived transparency, i.e., $1+1>1$.

However, despite their informational complementarity, the list and in-text formats are fundamentally two different presentations of the same set of citations, both serving the identical function of providing external evidence to generated content. This functional overlap means that the two formats can substitute for one another. Once users are exposed to citations in one format, they can already identify the sources underlying the generated content. Presenting the same citations in an additional format may display more comprehensive citation information on the interface, but does not address any fundamentally new or unmet user need (i.e., the need to understand the information sources). Moreover, users incur additional cognitive costs in processing the more complex presentation, such as understanding how in-text format function as indices to the list citations, which is often non-intuitive for general users (Narayanan Venkit et al., 2025). This may therefore produce a redundancy effect (Mayer et al., 2001). As a result, the marginal benefit brought by the second format diminishes. That is, while the combination of both formats achieves a " $1+1>1$ "

effect through informational complementarity, their functional redundancy makes the overall gain follow a " $1+1<2$ " pattern. Taken together, we hypothesize that:

H2a. The combined format produces higher transparency than either format alone.

H2b. The incremental gain from adding the second format is smaller than the gain from introducing the first.

The Moderating Effect of Task Goal Specificity

According to Task-Technology Fit theory (Goodhue & Thompson, 1995), when the presentation format matches the task characteristic, it can lead to more positive outcomes. In high goal specificity tasks, users focus on searching for information relevant to their goal in the generated text (Athukorala et al., 2016). Once locating the target information, they primarily check whether it is supported by a citation. In-text citations that are adjacent to the specific information make it clear to users that the claim is backed by the corresponding citation. Conversely, in low goal specificity tasks, users engage in broad browsing to build a comprehensive understanding of a topic (Hong, 2006). They focus more on the underlying sources of the whole generated text and seek out supplementary information. By grouping citations together and providing rich content information, the list format serves as a comprehensive overview, which enable users to evaluate all citations. Overall, the in-text format better matches high goal specificity tasks, while the list format better aligns with low goal specificity tasks. Accordingly, the in-text format is expected to be more effective in enhancing perceived transparency under high goal specificity tasks, whereas the list format is expected to be more effective under low goal specificity. Therefore, we hypothesize that:

H3a. The positive effect of the in-text format on perceived transparency is stronger in high goal specificity tasks than in low goal specificity tasks

H3b. The positive effect of the list format on perceived transparency is stronger in low goal specificity tasks than in high goal specificity tasks.

Methodology

We tested our hypotheses by conducting an online 2 (the list format: presence vs absence) $\times$ 2 (the in-text format: presence vs absence) $\times$ 2 (task goal specificity: high vs low) between group experiment.

Experimental System Development

To maximize ecological validity, we developed Generative AI systems based on the Retrieval Augmented Generation (RAG) approach (Lewis et al., 2020), which is also adopted by most popular Generative AI applications like ChatGPT, Gemini and Bing Chat. In the backend, we employed the same GLM-4-long large language model for all conditions to ensure internal validity. For the user interface, these systems mimicked the interface of popular Generative AI tools, including an input box, a chat window, and a conversation history sidebar. More importantly, these systems supported real-time and multi-turn conversations. User could type natural language queries and iteratively refine their information seeking. To implement citations, we prompted the large language model to embed citation links in the generated content. These citations were linked to curated credible sources which were verified in pilot testing. All systems accessed the same source materials, only citation formats varied across conditions. These systems were deployed on Vercel, a cloud server, to ensure stable access throughout data collection. More technical details can be found in Appendix A.

These systems offer advantages over static vignettes or wizard-of-oz approaches used in prior research (Shin et al., 2024; Zhou et al., 2025), as they enable authentic user-GenAI interaction and capture users' natural reactions more accurately.

Experimental Manipulations

We implemented four citation format conditions (Figure 3), including no citation, in-text format, list format, and combined format. In no citation group, generated content contained no citation hyperlinks. In the in-text format, superscript numeric markers (e.g., [1], [2]) were placed right after the supported sentence,

which contain clickable links to the sources. No citation preview or hover card was triggered by mouseover. In the list format, citations were presented as a list at the bottom, with full metadata (title and source) for each citation displayed. The combined format used the in-text superscripts as an index to a comprehensive reference list at the end of generated content.

To manipulate task goal specificity, following prior research (Hong, 2006; Tam & Ho, 2006), participants were randomly assigned to one of two predefined information search tasks. We chose healthcare as task context because it represents a typical high-stakes domain for Generative AI applications (Teddy, 2024). In the high goal specificity task, participants were required to locate the common ways to implement "Intermittent Energy Restriction" (IER), a popular weight loss method, and their differences. In the low goal specificity task, they need to gather broad information related to the topic. We selected "Intermittent Energy Restriction" (IER) for two reasons. First, interest in IER has been growing in recent health discussions (Burger, 2025). Second, IER-related information generally does not require urgent medical intervention, making it suitable for AI-assisted exploration (Claggett et al., 2024)





	
(a) no citation	(b) in-text format
	
(c) list format	(d) combined format

Figure 3. The Screenshots of Different Citation Format Conditions

Procedure and Measures

The experimental procedure is as follows: First, participants were randomly assigned to one of the experimental conditions. Upon assignment, they were asked to carefully read the task description, and subsequently answered the manipulation check regarding goal specificity and measures for control variables. Next, participants clicked a link embedded in the questionnaire to access the designated experimental system. We assigned a unique random number (ranging from 1000 to 9999) to each participant. Access to the Generative AI was granted only after participants inputting this number. The unique number allowed us to accurately match each participant's questionnaire responses with their corresponding interaction logs. Then, participants were free to interact with the system until they believed they can complete the task. After that, participants were asked to complete a manipulation check on citation format and the subsequent measurement scales.

All constructs were measured using items adapted from the existing research, with minor adjustments made to better suit the research context. Specifically, perceived transparency was measured with items adapted from Al-Natour et al. (2021). Prior knowledge of IER and AI knowledge, which were control variables in this study, were measured with scales from Bera et al. (2014) and Shao & Li (2022), respectively. Appendix B lists all measurement items. For all items, we used a seven-point Likert scale ranging from 1 (strongly disagree) to 7 (strongly agree).

Data Collection

Before the main experiment, we conducted a pilot study ($N = 89$) to test system functionality, validate experimental manipulations, and ensure the credibility of the information sources. Results showed that 98.9% of participants found the interaction with the Generative AI system smooth, indicating that the experimental system functioned successfully. In terms of source credibility, 38% rated the information sources as "somewhat credible," 48% as "credible," and 14% as "very credible," suggesting that the selected references were appropriate. Additionally, the manipulations of citation format and task goal specificity were validated as effective.

The main experiment was conducted via Credamo in April 2025. We applied three criteria to remove invalid samples. First, participants who failed to correctly identify the citation format in their assigned condition were excluded. Second, samples with no interaction records in the system were removed, as it indicated a failure to interact with the Generative AI. Third, incomplete questionnaires were considered invalid. After excluding invalid samples, a total of 293 valid responses remained for subsequent data analysis. A post hoc power analysis conducted in G*Power indicated that with a total sample of 293, eight conditions, and $\alpha = 0.05$, the study achieved a statistical power of 0.99 to detect a medium effect size ($f = 0.25$), exceeding the recommended threshold of 0.80 (Cohen, 1977). Chi-square tests on participant demographics and Generative AI experience showed no significant differences across conditions, confirming successful randomization. Additionally, we further conducted ANOVAs to examine whether participants' prior knowledge and AI knowledge differ across conditions. Results indicated that prior knowledge was well-balanced across conditions ($F = 1.562$, $p = 0.146$, partial $\eta^2 = 0.037$), while AI knowledge showed a statistically significant difference across conditions ($F = 3.649$, $p = 0.001$, partial $\eta^2 = 0.082$). Given the significant imbalance in AI knowledge and the theoretical relevance of both variables to perceived transparency, both prior knowledge and AI knowledge were incorporated as covariates in all subsequent analyses. Sample demographics and descriptive statistics of engagement data across experimental conditions are reported in Appendix C.

Data analysis and results

Manipulation Checks

We assessed the success of manipulations. To examine whether participants recognized the citation format correctly, we asked them to report on the citation format they saw (list only, in-text only, both, or none). Only those correctly identifying their assigned condition were considered valid, and the exclusion rate is 4.15%. Additionally, the manipulation of task goal specificity was checked via a three-item perceived goal specificity scale adapted from Tam & Ho (2006). The results of a t-test showed that perceived goal specificity of the high goal specificity task was significantly higher than that of low goal specificity task (M_{high}

goal specificity = 6.24 versus $M_{\text{low goal specificity}} = 5.73$, $t = 5.025$, $p < 0.001$). The manipulation of tasks was deemed to be satisfactory.

Measurement Model

We assessed constructs' reliability, convergent validity and discriminant validity using Partial Least Squares (PLS) structural equation modeling. As shown in Table 1, the composite reliability (CR) and Cronbach's alpha of all constructs exceed 0.7, and the average variance extracted (AVE) was above 0.5, indicating that the reliability is satisfactory. Further, all measurement items loaded above 0.7 on their respective constructs and exceeded their cross-loadings, indicating good convergent validity. The correlations between constructs were lower than the square root of each construct's AVE, confirming satisfactory discriminant validity.

Table 1. Reliability and Correlations								
	Mean	SD	CR	AVE	Alpha	PT	PK	AK
PT	4.91	1.64	0.958	0.884	0.935	0.940		
PK	3.41	1.75	0.961	0.924	0.918	0.287	0.961	
AK	5.45	0.91	0.888	0.727	0.814	0.139	0.328	0.852
Note: PT=perceived transparency, PK=prior knowledge, AK=AI knowledge								


Hypotheses Testing

We conducted an ANCOVA to examine the main effects of list and in-text, their interaction, and their respective interactions with task goal specificity on perceived transparency. Task goal specificity was coded based on the task participants were assigned to (0 = high goal specificity, 1 = low goal specificity). AI knowledge and prior knowledge were included as covariates. As shown in Table 2, both the list format ($F=149.20$, $p < 0.001$) and the in-text format ($F=98.23$, $p < 0.001$) had significant main effects on perceived transparency. That is, citations presented in list and in-text formats can enhance perceived transparency, supporting H1. To further detail the difference between the list only, in-text only, and combined formats, a multiple comparison test was conducted. Note that since perceived transparency in the no citation condition ($M=3.08$) was significantly lower than that in any citation condition, including it in the format comparison would risk confounding the format effect with the presence-versus-absence of citation effect. Thus, the no citation condition was excluded from subsequent analyses. Results (Figure 4) showed that the combined format generated significantly higher perceived transparency than the list format only ($p < 0.001$) and the in-text format only ($p < 0.001$), which supported H2a. The partial η^2 was 0.130, indicating that the effect size approached a large effect ($\eta^2 \approx 0.14$) (Cohen, 1977).


To assess whether exclusions of participants who failed the citation format manipulation check introduced selection bias, we conducted a robustness check by including these participants into the analysis. The results were consistent with the main analysis, with the main effects of list format ($F = 151.62$, $p < 0.001$), in-text format ($F=101.65$, $p < 0.001$), the list and in-text interaction ($F=46.67$, $p < 0.001$), and the list and goal specificity interaction ($F=9.905$, $p=0.002$). These results suggest that the exclusions did not meaningfully bias the findings.

Table 2. Results of ANCOVA					
IV	Sum of square	Df	Mean square	F	Sig.
list format	176.52	1	176.52	149.20	<0.001
in-text format	116.22	1	116.22	98.23	<0.001
task goal specificity	0.37	1	0.37	0.31	0.578
list * in-text	56.03	1	56.03	47.36	<0.001
list * task goal specificity	10.20	1	10.20	8.62	0.004
in-text * task goal specificity	0.02	1	0.02	0.02	0.902
prior knowledge (covariate)	33.28	1	33.28	28.13	<0.001

AI knowledge (covariate)	5.00	1	5.00	4.22	0.041
Note: DV=Perceived Transparency					



Results in Table 2 also showed that the interaction of the list and in-text formats ($F=47.36$, $p<0.001$), and the interaction of the list format and task goal specificity ($F=8.62$, $p=0.004$) were significant, while the interaction of the in-text format and task goal specificity was not significant ($F=0.02$, $p=0.902$). Thus, H3a was not supported. To further examine the significant interaction, we conducted simple main effects analyses (Figure 5). Figure 5a illustrated that the list format significantly enhanced perceived transparency when the in-text format was absent ($\Delta M=2.461$, $p<0.001$), but the effect was weaker when the in-text was present ($\Delta M = 0.662$, $p<0.001$). Similarly, the in-text format had a stronger effect on transparency when list format was absent ($\Delta M=2.169$, $p<0.001$) than when it was present ($\Delta M = 0.370$, $p=0.043$). These results suggested a negative interaction effect between the two citation formats on perceived transparency, supporting H2b. Figure 5b indicated that the positive effect of the list format on perceived transparency was amplified when task goal specificity was low, which supported H3b.



Discussions and Implications

Key Findings

This study investigates how different citation formats (list, in-text and their combination) affect users' perceived transparency, and how the effects vary under different task goal specificity. Key findings are as follows. First, citations presented in both the list and in-text formats effectively enhance perceived

transparency. Moreover, their combined format leads to the highest transparency. This is because the two citation formats provide complementary information (Lu et al., 2020). The list format primarily conveys information about citation content, whereas the in-text format conveys information about the citation correspondence. When combined, they offer more comprehensive information, thereby enhancing the overall effectiveness. This is consistent with cue summation theory (Jiang & Benbasat, 2007).

Second, though the combined format maximizes perceived transparency, the improvement shows a diminishing marginal effect (i.e., $1+1<2$). This is primarily due to their functional redundancy, once one format already satisfies users' needs, the additional format offers little in addressing any unmet needs, and may instead increase interface complexity, thereby producing a redundancy effect that attenuates its marginal benefit (Bian et al., 2025; Geissler et al., 2001).

Third, align with Task-Technology Fit theory (Goodhue & Thompson, 1995), the list format is significantly more effective under low goal specificity tasks. Since these tasks require users to gather broad and comprehensive information, the list format, which aggregates citations and offers content cues, is well suited to this requirement, making it more preferred by users. However, the effect of in-text is not moderated by goal specificity. A possible explanation is that in-text citations are embedded directly at the end of sentences. Accordingly, users can perceive that each claim is supported by a citation during reading, regardless of whether they are searching for specific information or browsing broadly. This makes users' perceived transparency consistent across task conditions. This interpretation is further supported by engagement data. When controlling the number of interaction rounds, the total number of citation clicks differed significantly across goal specificity conditions for the list format ($p=0.038$), but not in the in-text format condition ($p>0.1$), indicating that users interacted with in-text citations consistently regardless of their task goal specificity.

Theoretical Implications

This study makes three key theoretical contributions. First, this research contributes to the emerging literature on citation in Generative AI (Ding et al., 2025; Kim et al., 2025). While prior research has validated the effectiveness of providing citations (Ding et al., 2025), how citations should be presented remains underexplored. This study addresses this gap by investigating the effects of different citation formats (list, in-text, and their combination). Our findings reveal that both list and in-text formats enhance perceived transparency. More importantly, while the combined format maximizes the improvement, it also has a diminishing marginal effect.

Second, while prior studies have discovered complementary and substitutive effects in contexts where multiple information cues are presented (Bashirzadeh et al., 2022; Geissler et al., 2001; Jiang & Benbasat, 2007; Ling et al., 2024), such discussions remain limited in the context of citation presentation, leaving it unclear whether combining different citation formats is beneficial or detrimental. This study extends this stream of literature by unpacking the nuanced interplay between the list and in-text formats. Specifically, we found that combining both formats yields higher perceived transparency than either format alone (i.e., $1+1>1$), yet the two formats are also substitutive, in that the marginal gain of adding the second format is smaller than that of introducing the first.

Third, this study provides empirical evidence for the boundary conditions of TTF in the context of citation presentation for Generative AI (Goodhue & Thompson, 1995). Specifically, we found that the list format is more effective in low goal specificity tasks, while the effect of in-text format remains consistent across different level of goal specificity.

Practical Implication

This study offers several practical implications for citation presentation in Generative AI. First, for contexts that require high transparency, such as medical consultations and academic research, citations should be presented in a combined list and in-text format to achieve the highest perceived transparency. Second, given the diminishing marginal effect of the combined format, designers do not need to force the simultaneous presentation of both citation formats when interface space or resources are constrained (e.g., mobile vs. desktop, short vs. long outputs). A well-designed format can fulfill the vast majority of users' transparency needs. Third, as the effectiveness of citation format is impacted by task type, citation formats can be adapted according to users' task contexts. For example, when users need to gather broad information, the system

should prioritize presenting the list format citations. Additionally, developers can design options for different situations. If the users have a need, they can choose one of those formats.

Limitations and Future Research

Several limitations of this study should be acknowledged. First, perceived transparency was measured using three self-report items adapted from Al-Natour et al. (2021), which primarily capture whether users understand information sources. This operationalization does not fully encompass broader dimensions of transparency, such as process transparency (e.g., why a particular source was cited) and verifiability. Future research should expand the construct by incorporating related dimensions to sharpen construct validity. Additionally, behavioral indicators such as citation click rates, dwell time on cited sources, number of follow-up queries would provide stronger triangulation. Future studies are encouraged to integrate such behavioral data with self-report measures.

Second, the experimental design assumed that all cited sources were credible and of high quality. In practice, however, AI-generated citations vary considerably in source quality. Future research can examine the interaction between citation source credibility and citation format.

Third, regarding task manipulation, although a significant difference in perceived goal specificity was observed between the high goal specificity group ($M = 6.24$) and the low goal specificity group ($M = 5.73$), both groups exhibited relatively high mean scores, suggesting that the low goal specificity task may have retained a goal orientation. This may limit the manipulation validity of the goal specificity distinction. Future studies should consider redesigning low goal specificity tasks or adopting a more granular classification to ensure a clearer contrast between task types. Besides, the experimental scenario was situated specifically in the healthcare information context. While this setting provides ecological validity for health-related information seeking, it may limit the generalizability of the findings to other domains. Future research should replicate and extend these findings across diverse task domains to establish broader applicability.

Fourth, the present study predominantly recruited highly educated, experienced users with more than 12 months of generative AI experience. Future research is encouraged to extend the sample to broader or less experienced user populations to enhance the generalizability of the findings.

Finally, although the present study explains the diminishing marginal returns of the combined format from the perspective of functional redundancy, perceived redundancy and information overload were not directly measured. Future research could further examine the psychological mechanisms underlying the substitution effect between list and in-text citation formats.

Acknowledgement

The work described in this paper was partially supported by the grants from the Major Projects of Ministry of Education Humanities and Social Sciences Key Research Base (Project No. 22JJD870002)

Appendix A. Technical Details of Experimental Systems

The backend is implemented using Dify, leveraging Zhipu AI's GLM-4-long model with both temperature and Top P set to 0.7. The knowledge base comprises 55 documents from authoritative sources (e.g., Wikipedia, Baidu Baike), all centered on the topic of Intermittent Energy Restriction, and vectorized using Zhipu's embedding-3 model. Retrieval combines 90% semantic matching with 10% keyword matching, returning the top 10 most relevant sources per query. To implement citation, we prompt the large language model to strictly base content generation on reference sources, and properly mark these references within the content according to a specific format (i.e., the list, in-text and combined format). Besides, to prevent users' preference bias toward specific AI brands (e.g., ChatGPT vs Gemini), the system prompt also mandates anonymity, strictly prohibiting the model from revealing its identity, role, or developers.

Appendix B. Measures

Table 3. Measurement Constructs
--

Constructs	Items	Sources
Dependent Variables		
Perceived Transparency	- The Generative AI clarifies the information bias behind its generated content.	Al-Natour et al. (2021)
	- The Generative AI justifies the basis behind its generated content.	
	- Overall, the information that the Generative AI relies on to generate content is made known to me.	
Manipulation Checks		
Goal Specificity	- I knew what information I should access	Tam & Ho (2006)
	- I knew what I task target was.	
	- I had ideas on what aspect I should look for.	
Control Variables		
Prior Knowledge	- To what extent do you have knowledge of "intermittent energy restriction".	Bera et al. (2014)
	- To what extent do you have experience in using "intermittent energy restriction".	
AI Knowledge	- I have sufficient knowledge of Generative AI.	Shao & Li (2022)
	- I have expertise in applying Generative AI.	
	- I am skilled at using Generative AI.	

Appendix C. Sample Demographics and Descriptive Statistics

Table 4. Sample Demographics			
Variable	Category	Sample	Proportion
Gender	Male	133	45.4%
	Female	160	54.6%
Age	<20	6	2.0%
	21-25	119	40.6%
	26-35	114	38.9%
	36-45	30	10.2%
	>45	24	8.2%
Education	Below high school	2	0.7%
	High school	6	2.0%
	Bachelor's degree	220	75.1%
	Master's degree	58	19.8%
	Ph D. or equivalent	7	2.4%
Usage experience (months)	0-3	5	1.7%
	3-6	19	6.5%
	6-12	89	30.4%
	>12	180	61.4%

Usage frequency (hours/week)	<4	51	17.4%
	4-8	81	27.6%
	8-16	91	31.1%
	16-32	55	18.8%
	>32	15	5.1%
Note: Report suggests that users with a bachelor's degree or higher make up a large proportion of the GenAI user population (as high as 82%) (Sidoti & McClain, 2025), which indicates that our sample is reasonably representative in terms of education level.			

Table 5. Descriptive Statistics of Engagement Data

	high goal specificity task			low goal specificity task		
	list	in-text	combined	list	In-text	combined
participants (N)	38	34	37	37	31	38
number of interaction turns	2.00 (1.162)	2.47 (1.461)	2.11 (1.220)	2.73 (1.627)	3.0 (1.612)	2.55 (1.201)
number of interaction turns with clicks	0.32 (0.525)	0.65 (0.544)	0.41 (0.551)	0.19 (0.397)	0.77 (0.717)	0.45 (0.724)
total number of clicked links	0.53 (1.033)	0.82 (0.843)	0.51 (0.731)	0.22 (0.479)	1.32 (1.447)	0.55 (1.108)
total number of clicks	0.53 (1.033)	0.94 (1.043)	0.51 (0.731)	0.22 (0.479)	1.45 (1.786)	0.55 (1.108)
dwel time (min)	0.12 (0.459)	0.38 (1.246)	0.21 (0.65)	0.06 (0.312)	0.29 (0.491)	0.23 (0.825)
Note: Values are presented as Mean (SD).						

References

- Al-Natour, S., Benbasat, I., & Cenfetelli, R. (2021). Designing Online Virtual Advisors to Encourage Customer Self-disclosure: A Theoretical Model and an Empirical Test. *Journal of Management Information Systems*, 38(3), 798–827.
- Athukorala, K., Glowacka, D., Jacucci, G., Oulasvirta, A., & Vreeken, J. (2016). Is exploratory search different? A comparison of information search behavior for exploratory and lookup tasks. *Journal of the Association for Information Science and Technology*, 67(11), 2635–2651. <https://doi.org/10.1002/asi.23617>
- Bashirzadeh, Y., Mai, R., & Faure, C. (2022). How rich is too rich? Visual design elements in digital marketing communications. *International Journal of Research in Marketing*, 39(1), 58–76. <https://doi.org/10.1016/j.ijresmar.2021.06.008>
- Bera, P., Burton-Jones, A., & Wand, Y. (2014). How Semantics and Pragmatics Interact in Understanding Conceptual Models. *Information Systems Research*, 25(2), 401–419. <https://doi.org/10.1287/isre.2014.0515>
- Bian, Q., Wang, L., Xin, L., & Ma, B. (2025). Mismatch between warning information and protective behavior: Why experts + AI < 2? *Risk Analysis*, risa.70030. <https://doi.org/10.1111/risa.70030>
- Burger, C. (2025, February 23). *Study Reveals America's Top Health Questions*. Registered Nursing. <https://www.registerednursing.org/articles/americans-top-health-questions/>.
- Cheng, H.-F., Wang, R., Zhang, Z., O'Connell, F., Gray, T., Harper, F. M., & Zhu, H. (2019). Explaining Decision-Making Algorithms through UI: Strategies to Help Non-Expert Stakeholders. *Proceedings of*

- the 2019 CHI Conference on Human Factors in Computing Systems, 1–12. <https://doi.org/10.1145/3290605.3300789>
- Claggett, J. L., Kitchens, B., & Paino, M. (2024). Identifying the peripheral cues in the credibility assessment of online health information. *Information & Management*, 61(8), 104037. <https://doi.org/10.1016/j.im.2024.104037>
- Cohen, J. (1977). Statistical power analysis for the behavioral sciences. *Statistical Power Analysis for the Behavioral Sciences*, Rev. Ed.
- Deng & Poole. (2010). Affect in Web Interfaces: A Study of the Impacts of Web Page Visual Complexity and Order. *MIS Quarterly*, 34(4), 711. <https://doi.org/10.2307/25750702>
- Ding, Y., Facciani, M., Joyce, E., Poudel, A., Bhattacharya, S., Veeramani, B., Aguinaga, S., & Weninger, T. (2025). Citations and Trust in LLM Generated Responses. *Proceedings of the AAAI Conference on Artificial Intelligence*, 39, 23787–23795. <https://doi.org/https://doi.org/10.1609/aaai.v39i22.34550>
- Ding, Y., Facciani, M., Poudel, A., Joyce, E., Aguinaga, S., Veeramani, B., Bhattacharya, S., & Weninger, T. (2025). *Citations and Trust in LLM Generated Responses* (arXiv:2501.01303). arXiv. <https://doi.org/10.48550/arXiv.2501.01303>
- Dwivedi, Y. K., Kshetri, N., Hughes, L., Slade, E. L., Jeyaraj, A., Kar, A. K., Baabdullah, A. M., Koohang, A., Raghavan, V., Ahuja, M., Albanna, H., Albashrawi, M. A., Al-Busaidi, A. S., Balakrishnan, J., Barlette, Y., Basu, S., Bose, I., Brooks, L., Buhalis, D., ... Wright, R. (2023). “So what if ChatGPT wrote it?” Multidisciplinary perspectives on opportunities, challenges and implications of generative conversational AI for research, practice and policy. *International Journal of Information Management*, 71, 102642. <https://doi.org/10.1016/j.ijinfomgt.2023.102642>
- European Union. (2024). *Artificial Intelligence Act*. <https://artificialintelligenceact.eu/chapter/4/>
- Gedikli, F., Jannach, D., & Ge, M. (2014). How should I explain? A comparison of different explanation types for recommender systems. *International Journal of Human-Computer Studies*, 72(4), 367–382. <https://doi.org/10.1016/j.ijhcs.2013.12.007>
- Geissler, G., Zinkhan, G., & Watson, R. (2001). Web Home Page Complexity and Communication Effectiveness. *Journal of the Association for Information Systems*, 2(1), 1–48. <https://doi.org/10.17705/1jais.00014>
- Goodhue, D. L., & Thompson, R. L. (1995). Task-Technology Fit and Individual Performance. *MIS Quarterly*, 19(2), 213. <https://doi.org/10.2307/249689>
- Ho, D. E. (2024, January 11). *Hallucinating Law: Legal Mistakes with Large Language Models are Pervasive*. <https://hai.stanford.edu/news/hallucinating-law-legal-mistakes-large-language-models-are-pervasive>
- Hong, T. (2006). The Internet and Tobacco Cessation: The Roles of Internet Self-Efficacy and Search Task on the Information-Seeking Process. *Journal of Computer-Mediated Communication*, 11(2), 536–556. <https://doi.org/10.1111/j.1083-6101.2006.00026.x>
- Hong, W., Thong, J. Y. L., & Tam, K. Y. (2004). The Effects of Information Format and Shopping Task on Consumers’ Online Shopping Behavior: A Cognitive Fit Perspective. *Journal of Management Information Systems*, 21(3), 149–184. <https://doi.org/10.1080/07421222.2004.11045812>
- Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B., & Liu, T. (2025). A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions. *ACM Transactions on Information Systems*, 43(2), 1–55. <https://doi.org/10.1145/3703155>
- Jiang & Benbasat. (2007). The Effects of Presentation Formats and Task Complexity on Online Consumers’ Product Understanding. *MIS Quarterly*, 31(3), 475. <https://doi.org/10.2307/25148804>
- Kim, S. S. Y., Vaughan, J. W., Liao, Q. V., Lombrozo, T., & Russakovsky, O. (2025). Fostering Appropriate Reliance on Large Language Models: The Role of Explanations, Sources, and Inconsistencies. *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, 1–19. <https://doi.org/10.1145/3706598.3714020>
- Kizilcec, R. F. (2016). How Much Information?: Effects of Transparency on Trust in an Algorithmic Interface. *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems*, 2390–2395. <https://doi.org/10.1145/2858036.2858402>
- Lenz, T., Brackey, R., & Liu, J. (2025, July 15). *The Effects of Citations and Confirmation Bias on Trust in Chatbots*. Proceedings of the Human Factors and Ergonomics Society Annual Meeting.
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-Augmented Generation for Knowledge-

- Intensive NLP Tasks. *Proceedings of the 34th International Conference on Neural Information Processing Systems*, 33, 9459–9474.
- Lim, K. H., & Benbasat, I. (2000). The Effect of Multimedia on Perceived Equivocality and Perceived Usefulness of Information Systems. *MIS Quarterly*, 24(3), 449. <https://doi.org/10.2307/3250969>
- Ling, S., Zhang, Y., & Du, N. (2024). More Is Not Always Better: Impacts of AI-Generated Confidence and Explanations in Human–Automation Interaction. *Human Factors: The Journal of the Human Factors and Ergonomics Society*, 66(12), 2606–2620. <https://doi.org/10.1177/00187208241234810>
- Liu, N., Zhang, T., & Liang, P. (2023). Evaluating Verifiability in Generative Search Engines. *Findings of the Association for Computational Linguistics: EMNLP 2023*, 7001–7025. <https://doi.org/10.18653/v1/2023.findings-emnlp.467>
- Lu, J., Lee, D. (Dk), Kim, T. W., & Danks, D. (2020). Good Explanation for Algorithmic Transparency. *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, 93–93. <https://doi.org/10.1145/3375627.3375821>
- Mayer, R. E., Heiser, J., & Lonn, S. (2001). Cognitive Constraints on Multimedia Learning: When Presenting More Material Results in Less Understanding. *Journal of Educational Psychology*, 93(1), 187–198. <https://doi.org/10.1037/0022-0663.93.1.187>
- Nadkarni, S., & Gupta, R. (2007). A Task-Based Model of Perceived Website Complexity. *MIS Quarterly*, 31(3), 501–524. <https://doi.org/10.2307/25148805>
- Narayanan Venkit, P., Laban, P., Mao, Y., & Wu, C.-S. (2025). Search Engines in the AI Era: A Qualitative Understanding to the False Promise of Factual and Verifiable Source-Cited Responses in LLM-based Search. *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*, 1325–1340. <https://doi.org/10.1145/3715275.3732089>
- Ning, X., Lu, Y., Li, W., & Gupta, S. (2024). How transparency affects algorithmic advice utilization: The mediating roles of trusting beliefs. *Decision Support Systems*, 183, 114273. <https://doi.org/10.1016/j.dss.2024.114273>
- OECD. (2025). *Governing with Artificial Intelligence: The State of Play and Way Forward in Core Government Functions*. OECD Publishing. <https://doi.org/10.1787/795de142-en>
- Ouyang, J., & Narechania, A. (2026). Analyzing the Presentation, Content, and Utilization of References in LLM-powered Conversational AI Systems. *Proceedings of the Extended Abstracts of the 2026 CHI Conference on Human Factors in Computing Systems*, 1–8. <https://doi.org/10.1145/3772363.3798415>
- Severin, W. (1967). Another look at cue summation. *AV Communication Review*, 15(3), 233–245. <https://doi.org/10.1007/BF02768608>
- Shao, Z., & Li, X. (2022). The influences of three task characteristics on innovative use of malleable IT: An extension of adaptive structuration theory for individuals. *Information & Management*, 59(3), 103597. <https://doi.org/10.1016/j.im.2022.103597>
- Sharma, N., Liao, Q. V., & Xiao, Z. (2024). Generative Echo Chamber? Effect of LLM-Powered Search Systems on Diverse Information Seeking. *Proceedings of the CHI Conference on Human Factors in Computing Systems*, 1–17. <https://doi.org/10.1145/3613904.3642459>
- Shin, D., Koerber, A., & Lim, J. S. (2024). Impact of misinformation from generative AI on user information processing: How people understand misinformation from generative AI. *New Media & Society*, 14614448241234040. <https://doi.org/10.1177/14614448241234040>
- Sidoti, O., & McClain, C. (2025, June 25). *34% of U.S. adults have used ChatGPT, about double the share in 2023*. <https://www.pewresearch.org/short-reads/2025/06/25/34-of-us-adults-have-used-chatgpt-about-double-the-share-in-2023/>
- Tam, K. Y., & Ho, S. Y. (2006). Understanding the Impact of Web Personalization on User Information Processing and Decision Outcomes. *MIS Quarterly*, 30(4), 865. <https://doi.org/10.2307/25148757>
- Teddy, R. (2024, September 11). *Dr. Chatbot Will See You Now*. <https://www.nytimes.com/2024/09/11/health/chatbots-health-diagnosis-treatments.html>
- Vessey, I. (1994). The effect of information presentation on decision making: A cost-benefit analysis. *Information & Management*, 27(2), 103–119. [https://doi.org/10.1016/0378-7206\(94\)90010-8](https://doi.org/10.1016/0378-7206(94)90010-8)
- Vessey, I., & Galletta, D. (1991). Cognitive Fit: An Empirical Study of Information Acquisition. *Information Systems Research*, 2(1), 63–84. <https://doi.org/10.1287/isre.2.1.63>
- You, S., Yang, C. L., & Li, X. (2022). Algorithmic versus Human Advice: Does Presenting Prediction Performance Matter for Algorithm Appreciation? *Journal of Management Information Systems*, 39(2), 336–365. <https://doi.org/10.1080/07421222.2022.2063553>

- Zhou, J., Lu, Y., & Chen, Q. (2025). GAI identity threat: When and why do individuals feel threatened? *Information & Management*, 62(2), 104093. <https://doi.org/10.1016/j.im.2024.104093>
- Zigurs, I., & Buckland, B. K. (1998). A Theory of Task/Technology Fit and Group Support Systems Effectiveness. *MIS Quarterly*, 22(3), 313. <https://doi.org/10.2307/249668>