

Optimizing for citation, not clicks: the CITED framework for content strategy in AI-mediated search

Avik Bal

Practitioner-academic paper for peer review · Released: 15th June 2026

Abstract

Generative engines answer a query by retrieving sources, combining them, and citing a few. This moves the unit of visibility from a ranked link to a citation inside an answer. In 2024, 58.5 per cent of United States Google searches ended with no click to the open web; by early 2026 the figure reached 68 per cent, and searches that trigger an AI Overview are zero-click about 83 per cent of the time. Keyword-era content strategy was built for a click economy that is contracting, while roughly 73 per cent of business buyers now use AI tools during research. This paper synthesizes about 50 academic and industry studies from 2024 to 2026 and the mechanism evidence behind them. Citation is decided in three steps: retrieval into a candidate set, selection within that set, and the brand's strength as an entity inside the model. Experimental work found that adding quotations, statistics, and citations raised a source's visibility in answers by 30 to 41 per cent, while keyword stuffing scored below the unedited baseline, and that lower-ranked pages gained most. Industry measurement found brand mentions correlate with AI citation about three times more strongly than backlinks, and that source overlap between two leading engines can be as low as 11 per cent. The paper consolidates these findings into a practitioner framework, CITED: Citation Surface, Information Architecture, Technical Foundation, Evidence-Rich Extraction, and Demand Capture. Each dimension is tied to one mechanism and one measurable signal, with a worked example and an implementation sequence. The paper defines the Citation Flywheel, states six falsifiable propositions, grades the evidence by quality, and sets out where the framework is least certain, including per-engine instability, measurement standards, and the value of a citation that produces no click.

Keywords: generative engine optimization; answer engine optimization; zero-click search; citation visibility; content strategy; large language models

The visibility problem in generative search

A marketing team publishes a 2,000-word guide on a question its buyers ask often. The page reaches the first results screen of Google within three months. Organic clicks barely move. The query now returns an AI Overview that answers the question on the results page, with the guide listed beneath it, cited once in small print or not at all. The page earned its ranking. Almost no one clicked.

Generative engines answer a query by reading several sources, combining them into one response, and attaching a short list of citations. ChatGPT, Google AI Overviews, Perplexity, and Gemini all work this way. SparkToro's 2024 analysis of Datos clickstream data found that 58.5 per cent of United States Google searches and 59.7 per cent of European Union searches ended with no click to the open web.¹ By the first four months of 2026 the United States figure had reached 68 per cent, and searches that surface an AI Overview are zero-click about 83 per cent of the time, with the top organic position losing roughly a third of its remaining clicks when an overview appears.² Visibility and traffic have come apart.

This matters because the audience has moved with the technology. Industry surveys in early 2026 put the share of business buyers using AI tools during purchase research at about 73 per cent.¹⁵ A brand that is invisible inside AI answers is invisible at the start of a growing share of buying

journeys, whatever its Google ranking. The keyword-era chain of rank, then click, then session, then conversion now breaks at the first link for most informational queries.

The argument here is direct. Content now competes for citation inside generated answers, the rules of that competition are partly understood, and they differ from the rules of ranking. The research that describes them is scattered across four labels for one shift, which leaves teams without a usable plan. This paper does three things. It explains the mechanism by which an engine decides what to cite. It consolidates the research into one framework, CITED²², with a mechanism, a measurable signal, and a worked use case behind each part. And it states the framework as falsifiable propositions, so it can be tested rather than believed. CITED is treated throughout as an analytical model, not a product.

Evidence and how it was graded

The source base is about 50 studies on large language models and search published between 2024 and 2026, found across academic databases and preprint servers, with primary industry measurement added. Each study was read for its claim about keyword strategy, content relevance, user intent, traditional versus generative optimization, and the methods it tested. This is a structured synthesis, not a systematic review, and the evidence is uneven, so it is graded in three tiers. The strongest tier is peer-reviewed and preprint research, led by the generative engine optimization experiments of Aggarwal and colleagues at KDD 2024.³ The second tier is primary measurement with a stated method, such as SparkToro's clickstream studies and Profound's citation analysis.¹¹⁷ The third tier is vendor measurement reported through industry compilations, which is directional and is flagged as such wherever it appears. Observations are kept separate from interpretation: a figure states what was measured, and the reading that follows is the author's.

How a generative engine decides what to cite

To know how to be cited, it helps to know how the citation is made. An engine does not rank a page and stop. It runs a sequence, and each step filters the field. Three forces decide the outcome.

Retrieval decides who is eligible

The engine takes the query, often rewrites it into several sub-queries, and pulls a candidate set of passages from an index or a live fetch. Classic search signals still govern this step. An analysis of 432,000 keywords found that 97 per cent of Google AI Overviews cite at least one source from the top 20 organic results, so ranking remains the entry ticket.²⁰ The rewriting changes who qualifies: in Ahrefs data, 86.85 per cent of AI Overviews do not use the exact search query, which is part of why roughly 80 per cent of cited URLs across engines do not rank in the top 100 for the original query.¹⁶ Retrieval also works on passages, not whole pages, so one strong section can carry a page into the set while the rest is ignored.

Selection decides who is quoted

From the retrieved set, the model chooses which passages to quote and attribute. Content design changes this step, and it rewards the opposite of keyword tactics. In the GEO experiments, adding quotations, statistics, and citations raised a source's visibility in answers by 30 to 41 per cent, while keyword stuffing scored 8.7 per cent below the unedited baseline.³ Selection also redistributes visibility toward the underdog. The Cite Sources method raised a fifth-ranked page's visibility by 115 per cent in the same study, while the top-ranked page lost about 30 per cent.³ A page that ranks

poorly but reads well for extraction can take citation share from a page that ranks well but reads poorly.

Entity strength decides who exists in the model

A third force sits beneath both steps. Some engines, ChatGPT most of all, answer partly from knowledge fixed during training rather than from a live fetch. Here the deciding factor is how strongly the brand exists as an entity across the web. Ahrefs' study of 75,000 brands found brand web mentions correlate with AI citation rates at 0.664, about three times the 0.218 correlation for backlinks, with mentions in YouTube titles and transcripts the single strongest factor at 0.737.¹⁶ A brand named often across authoritative sources surfaces in answers even when no page is retrieved.

The three forces do not transfer across engines

These mechanics differ enough by engine that visibility on one says little about another. An analysis of 680 million citations found only 11 per cent domain overlap between ChatGPT and Perplexity. Google's AI Overviews and AI Mode cite the same URLs just 13.7 per cent of the time, and one cross-platform study measured citation-volume variance of up to 615 times for the same brand across engines.¹⁹ A team that measures one engine is blind to most of the field. CITED assigns one dimension to each of the three forces, and two more to the parts the engines leave to the brand: how a topic is structured, and how visibility turns into revenue.

What the research agrees on

Read across the 2024 to 2026 literature, the same patterns recur. They hold across e-commerce, publishing, and general web content, and across controlled experiments and field measurement.

Keyword frequency stops being the lever

The clearest pattern is the fall of exact-match keyword density and the rise of semantic coverage, entities, and intent. Studies of artificial intelligence in search report that models read for meaning, so complete topic coverage surfaces across many related queries rather than one.¹² Real-time systems that generate and update keyword sets from live trend data outperform manual research in fast-moving categories such as e-commerce.¹⁴ The GEO experiments make the point sharply: the keyword-stuffing tactic scored below the unedited baseline, while content-level credibility signals won.³

Structure and entities carry relevance

A second pattern is that machine-readable structure decides whether a model can use a page. Content-centric agent systems that analyze, revise, and re-evaluate a page lift its measured influence in answers by working on entities, claims, and structure rather than word counts.⁶ Field measurement agrees: pages with well-organized headings are about 2.8 times more likely to be cited, and question-and-answer and list formats outperform dense paragraphs.¹⁶

Intent modeling helps, then plateaus

A third pattern is cautious. Role-augmented, intent-driven optimization tailors a page to the informational roles behind a query and improves traction, but the same work reports a ceiling near 70 per cent effectiveness and a risk that overly specific intent modeling overfits one query and loses related ones.⁵ Intent is a strong lever with a known limit.

Earned authority beats owned pages

A fourth pattern compares the playbooks directly. Generative engines lean toward earned media and authoritative third-party content over brand-owned pages making the same claim.⁴ Industry analysis attributed to a University of Toronto study put the gap at roughly 325 per cent more citations for earned media than for owned content.⁴ Studies of what generative engines prefer point the same way, toward credible, well-attributed, semantically clear sources rather than promotional copy.⁸⁷

The metric moves from clicks to citations

The fifth pattern ties the others together. As zero-click behavior rises, the measure that matters moves from click-through to citation share inside answers.¹² Field studies of about 10,000 websites also report that engines prefer content with predictable, semantically consistent phrasing, and that a single synthesized source can serve very different readers, which raises the stakes for being the source the engine picks.⁷

The same shift, four names

Here is the problem for practice. These findings arrive under four labels. Generative engine optimization names the shift from the angle of synthesized answers. Answer engine optimization names it from the angle of featured, citation-backed responses. Semantic SEO names it from the angle of meaning. Entity SEO names it from the angle of knowledge graphs. The four overlap heavily and rarely contradict each other. None of them, alone, tells a content team what to do across all of them. CITED is built to be that single plan.

The CITED framework

CITED consolidates the findings into five dimensions a content team can own and measure.²² The letters read from top to bottom: Citation Surface, Information Architecture, Technical Foundation, Evidence-Rich Extraction, and Demand Capture. Three dimensions map to the three forces behind citation. Two own the work the engines leave to the brand. Figure 1 shows the structure and the loop that connects it.



Figure 1. The CITED framework and the Citation Flywheel. Three dimensions correspond to retrieval, selection, and entity strength; two own topic architecture and conversion. The loop shows how early citations compound into authority that raises future citation probability.

C: Citation Surface (entity strength and earned presence)

Citation Surface is the brand's presence in the sources an engine draws on, and its share of citations within answers for its target queries. It owns the entity-strength force. The evidence is consistent: brand web mentions correlate with citation at 0.664 against 0.218 for backlinks, and earned third-party content earns several times more citations than owned pages.¹⁶⁴ Wikipedia is the most-cited single domain in ChatGPT, Reddit leads on Perplexity, and review and listicle sources such as Forbes, G2, and Nerdwallet recur across engines.¹⁷

Use case. A mid-market software brand earns Citation Surface by being named in the G2 category page, the industry roundups, the comparison listicles, the relevant Reddit threads, and the YouTube explainers that engines pull for its category, and by appearing beside its category terms across the web. Publishing more owned pages does little here. Being talked about elsewhere does most of the work. The signal to track is citation share per engine across a fixed query set, plus brand-mention volume and how often the brand co-occurs with its category terms.

I: Information Architecture (topic and entity structure)

Information Architecture is how content is organized into entities, topics, and internal links so a model can identify what the brand covers and connect its pages. It owns the breadth the engines demand once they rewrite a query into many sub-queries. Because 86.85 per cent of AI Overviews do not use the exact query, coverage of a whole topic cluster, not a single page, decides appearance across the sub-queries an engine generates.¹⁶ Structural revision of entities and claims raises measured influence more than edits to word counts.⁶

Use case. The same software brand defines the entities it wants to own (its product, the problem it solves, the alternatives, the comparisons, the use cases), covers each as a connected cluster with consistent naming, and links them so the relationship is explicit. The brand becomes retrievable across the long tail of reformulated queries and recognizable as the category's source. The signal is breadth of appearance across a cluster of related queries rather than a single query, and consistent entity naming across the site and the web.

T: Technical Foundation (retrieval eligibility)

Technical Foundation is the machine-readability layer that decides whether a page can enter the candidate set at all: crawl access for the right bots, schema, clean and fast rendering, and a healthy organic ranking, since 97 per cent of AI Overviews still cite from the top 20 organic results.²⁰ The detail that teams miss is that training crawlers and search crawlers are now separate. OpenAI runs GPTBot for training and OAI-SearchBot for its search index; Anthropic runs ClaudeBot for training and Claude-SearchBot for search. A brand can opt out of training while staying eligible for citation, but blocking the search crawlers removes it from those answers.²¹

The llms.txt file deserves a clear statement, because it is widely recommended and largely ignored. Google declined to support it on the record in July 2025, no major engine has committed to reading it, and controlled studies found no correlation between publishing it and AI visibility; server logs show the answer bots almost never fetch it.²¹ The lever that works today is a robots.txt that allows OAI-SearchBot, Claude-SearchBot, PerplexityBot, and Googlebot, paired with schema, which one vendor study associated with about 40 per cent more AI Overview appearances, and server-rendered HTML a model can parse. Use case: the software brand confirms search-crawler access in its logs, ships product and FAQ schema, fixes any client-side-only templates, and protects the rankings that gate retrieval. The signal is search-crawler access in server logs, schema coverage, and organic position on priority queries.

E: Evidence-Rich Extraction (citation selection)

Evidence-Rich Extraction is writing so a model can lift a correct, self-contained claim from a single passage. It owns the selection force, and it carries the strongest experimental result in the field. Table 1 shows the GEO tactic effects. The pattern is plain: sourced quotations, sourced statistics, and citations win, keyword stuffing loses, and the best tactic depends on the topic.³ Field measurement agrees that pages with expert quotes earn more citations than pages without.¹⁶

Use case. The software brand opens each section with a one- or two-sentence direct answer, adds at least one sourced statistic and one quotable definition per section, and matches the tactic to the domain, leaning on statistics for finance and legal topics and on expert quotations for strategy and opinion topics. The test is simple: if an engine reads one paragraph of the page, can it quote a correct, sourced claim from it? The signal is extractable-claim density and the appearance of the brand's own phrasing inside answers.

Table 1. What raises citation, from the GEO experiments

Tactic	What it changes	Visibility effect vs baseline	Best-fit domain
Quotation Addition	Adds quotations from credible authorities	Around +40 per cent; the top single tactic	Opinion, culture, strategy
Statistics Addition	Replaces vague claims with sourced numbers	About +30 to +37 per cent	Law, finance, technical
Cite Sources	Adds citations to credible sources	About +28 per cent alone; +115 per cent for a page ranked fifth	Factual, debated topics
Fluency Optimization	Improves readability and coherence	About +29 per cent	General interest
Keyword Stuffing	Repeats query terms, the classic SEO move	About 9 per cent below baseline	None

Results from Aggarwal et al. (2024) on GEO-bench, about 10,000 queries, with a 200-query Perplexity validation. Figures are relative to an unedited baseline and vary by metric and domain.

D: Demand Capture (turning citation into revenue)

Demand Capture converts answer-level visibility into business outcomes when the click may never arrive. This is where CITED reaches past the research frontier, because most studies stop at whether a source is cited and few ask what the citation is worth. The industry signals that exist point to real value: one 2026 report found that the top citation position carried a 23 times conversion premium, and that cited brands earned about 35 per cent more organic clicks and 91 per cent more paid clicks than uncited competitors.¹⁸ The academic e-commerce work shows that content tuned for both relevance and conversion lifts click-through and conversion when a click does happen.⁹

Use case. The software brand makes its name memorable inside answers, so a buyer who saw it cited searches for it directly afterward, holds the few high-intent queries where clicks still occur, and measures branded-search lift and assisted conversions rather than sessions alone. AI citation is treated as the new top of the funnel, which still has to be converted downstream. The signal is branded-search lift, assisted and branded conversions, and share of voice in answers for commercial queries.

The five dimensions run in a dependency order. Technical Foundation makes retrieval possible. Evidence-Rich Extraction wins selection. Information Architecture builds breadth, and Citation Surface builds the entity that the model remembers. Demand Capture turns the result into revenue. A brand can score well on one and fail on another, which is why a single label such as semantic SEO or GEO leaves the gaps that CITED closes.

Table 2. CITED, the mechanism it owns, and the evidence

Dimension	Force it owns	Evidence	Signal to track
C: Citation Surface	Entity strength in the model	Brand mentions correlate 0.664 vs 0.218 for backlinks; earned media around 325 per cent more citations	Citation share per engine; brand-mention volume; co-occurrence with category terms
I: Information Architecture	Breadth across rewritten queries	86.85 per cent of AI Overviews do not use the exact query; structural revision lifts influence	Appearance across a cluster of related queries; entity-naming consistency
T: Technical Foundation	Retrieval eligibility	97 per cent of overviews cite top-20 organic; search crawlers separate from training crawlers; schema around +40 per cent overviews	Search-crawler access in logs; schema coverage; organic rank
E: Evidence-Rich Extraction	Citation selection	Quotations, statistics, citations raise visibility 30 to 41 per cent; keyword stuffing below baseline	Extractable-claim density; brand phrasing inside answers
D: Demand Capture	Conversion of visibility	Top citation position, 23 times conversion premium; cited brands 35 per cent more clicks; 73 per cent of B2B buyers use AI	Branded-search lift; assisted conversions; commercial share of voice

The Citation Flywheel

The five dimensions connect through a feedback loop that Bal terms the Citation Flywheel.²² The mechanism is plain. A citation is itself a signal that other systems observe. More citations raise the brand's mention volume and authority for a topic. Higher authority raises the chance the engine cites the brand again for related queries, and feeds the entity layer that decides parametric recall. An early lead compounds, because brand mentions, the strongest single predictor of citation, accumulate from prior visibility.

The loop has support and a clear limit. Research on authority in AI search argues that real authority still depends on human-driven trust signals and external references that an engine cannot manufacture for a brand on its own, so the loop turns only when the underlying content earns genuine references.¹⁰ The loop also runs in reverse: a brand absent from early citations on a topic grows less likely to be cited as competitors accumulate the mention advantage. There is one opening for challengers. Because selection favors extractable content over rank, a fifth-ranked page gained 115 per cent visibility from one tactic while the top page lost ground.³ A challenger can start the loop with strong extraction before it has the authority, which is why the order of work matters.

Applying CITED in a business

CITED is built to be run, not admired. The five dimensions translate into a sequence, owners, signals, and a measurement habit. The sequence follows the dependency order and the cost-to-impact ratio. Technical Foundation comes first, because a page that cannot be retrieved cannot be cited, and the fix is fast. Evidence-Rich Extraction comes second, because it is the cheapest content change with the largest measured lift, and it works on pages already in the retrieved set. Information Architecture comes third, building breadth across the topic. Citation Surface comes fourth, because earned authority is the slowest to build and the highest ceiling. Demand Capture runs across all of them as the measurement and conversion layer.

A worked example shows the order. A mid-market HR analytics platform serving the Middle East and Southeast Asia finds it is absent from AI answers for queries its buyers ask, such as how to measure attrition risk. Week one, the technical lead confirms OAI-SearchBot, Claude-SearchBot, and PerplexityBot can reach the site, ships FAQ and product schema, and fixes a client-side-rendered docs section, which restores retrieval eligibility. Weeks two to six, editors rewrite the ten priority pages to open with direct answers and add a sourced statistic and a quotable definition to each section, matching statistics to the finance-heavy topics, which lifts the brand's phrasing into answers on several queries. Through the quarter, the content lead builds out the attrition cluster as connected entities, covering the problem, the methods, the benchmarks, and the comparisons, which widens appearance across the reformulated sub-queries. Over two quarters, the partnerships lead earns placements in the analyst roundups, the review sites, and the practitioner communities that the engines cite for the category, which lifts citation share. Throughout, the analyst tracks the result and ties it to branded search and pipeline.

Measurement is the part teams skip and the part that decides whether the work pays. The method is repeatable. Pick 50 to 100 priority queries that buyers actually ask. Run them monthly across ChatGPT, Perplexity, Google AI Overviews, and Google AI Mode, because source overlap between engines runs as low as 11 per cent and a single-engine view misses most of the field. Record whether the brand is cited, in what position, and against which competitors, then compute citation share per engine. Tie that share to branded-search volume in Search Console and to assisted conversions in analytics, so the chain from citation to revenue is visible. Dedicated AI-visibility trackers and brand-radar tools can automate the query runs across engines. Table 3 sets out the first moves by the owner.

Table 3. From framework to first moves

Owner	First action	Signal to track	Visible in
SEO / technical lead	Confirm search-crawler access, ship schema, fix client-side-only rendering, protect priority rankings (T)	Search-crawler hits in logs; schema coverage; rank on priority queries	Weeks
Writer / editor	Open each section with a direct answer; add one sourced statistic and one quotable definition; match tactic to domain (E)	Extractable-claim density; brand phrasing in answers	Weeks
Content lead	Build the topic as connected entities with consistent naming and internal links (I)	Appearance across related queries on each engine	1 to 2 quarters
PR / partnerships	Earn placements in the roundups, review sites, and communities engines cite for the category (C)	Citation share; brand-mention volume	1 to 2 quarters
Analyst	Run 50 to 100 queries monthly across four engines; compute citation share; tie to branded search and assisted conversions (D)	Citation-share baseline; branded-search lift; pipeline	1 quarter

Six falsifiable propositions

A framework that cannot be tested is decoration. The evidence base has a known weakness: few agreed metrics, little independent replication, and almost no longitudinal data. CITED answers that with six propositions, each measurable and disprovable, which together form a test battery and a partial measurement standard.

P1 (Extraction). Content with higher extractable-claim density, meaning direct answers, sourced statistics, and quotable definitions, earns a higher citation share than keyword-matched content of similar length on the same topic. *Test: run two versions across a fixed query set and measure citation share per engine. Grounded in Aggarwal et al. (2024).*

P2 (Earned over owned). For a given query, authoritative third-party sources are cited more often than brand-owned pages making the same claim. *Test: compare citation rates of owned versus earned mentions across a tracked query set. Grounded in Chen et al. (2025) and brand-mention measurement.*

P3 (Architecture). Brands with deeper topic-cluster coverage of an entity appear across a wider set of related queries than brands with single-page coverage, controlling for backlinks. *Test: measure cluster coverage against breadth of query appearance.*

P4 (Flywheel). A brand's citation share in one period predicts its citation share in the next period for related queries, beyond what backlinks predict, because mentions accumulate. *Test: a lagged panel across a query set over six to twelve months.*

P5 (Technical floor). Pages reachable only by training crawlers, blocked to search crawlers, or rendered client-side are cited less than equivalent server-rendered, search-crawlable pages. *Test: use crawler access and render mode as predictors of citation, holding content quality constant.*

P6 (Demand). Citation share for a topic predicts branded-search lift and assisted conversions even when direct clicks from those answers are flat. *Test: correlate citation-share time series with branded-search and assisted-conversion series. This proposition tests whether a citation is worth anything, the open question behind the D dimension.*

Limitations and future research

Five limits deserve a direct statement. First, this is a structured synthesis, and the source pool ranges from peer-reviewed experiments to vendor measurement of uneven quality, so the third-tier figures are directional rather than settled. Second, the strongest content result, that quotations, statistics, and citations win selection, rests heavily on one well-run study and needs independent replication across more engines and a longer horizon. Third, the engines are opaque, change monthly, and disagree with each other, with source overlap as low as 11 per cent, so any single measurement is unstable and the field still lacks a standard citation-share metric reported with confidence intervals. Fourth, most studies are English-language and short-horizon, so multilingual and longitudinal evidence is thin. Fifth, the Demand Capture layer, which decides whether the work pays, is the least studied part of the field and the part CITED leans on most.

The six propositions set the agenda. Each can run as a field study on a fixed query set across several engines, with P6 the most valuable, since it tests whether citation without a click produces demand. Alongside the propositions, the field needs a standard citation-share metric with confidence intervals, coverage of non-Google engines and non-English markets, and panels long enough to confirm or break the Citation Flywheel.

What good looks like

The click economy that funded fifteen years of content strategy is contracting, and the reporting habits built on it now hide more than they show. The research on generative search already points the way, though it speaks in four dialects and the engines keep moving. CITED puts the findings in one language a marketing team can run: be reachable by the search crawlers (T), be quotable with

evidence (E), be the recognized source on a topic (I), be named across the web (C), and turn that visibility into demand (D). The framework's weakest claim is also its most valuable. Getting cited is now reasonably well understood. What a citation is worth, and how to convert it when the click never comes, is the open question that will decide whether AI-era content strategy earns its budget.

Author

Avik Bal is a B2B Digital marketing practitioner with more than 12 years of experience in Web Architecture, SEO, Content Strategy, and marketing analytics for enterprise software, financial technology, and HR technology. He is the author of the book **CITED: The Growth Operating System for AI Search**.

References

1. Fishkin, R. (2025) '2024 zero-click search study', SparkToro, available at <https://sparktoro.com/blog/2024-zero-click-search-study-for-every-1000-us-google-searches-only-374-clicks-go-to-the-open-web-in-the-eu-its-360/> (accessed 14th June, 2026).
2. Fishkin, R. (2026) 'In 2026, less than one third of Google searches still send a click', SparkToro, available at <https://sparktoro.com/blog/in-2026-less-than-one-third-of-google-searches-still-send-a-click/> (accessed 14th June, 2026).
3. Aggarwal, P., Murahari, V., Rajpurohit, T., Kalyan, A., Narasimhan, K. and Deshpande, A. (2024) 'GEO: generative engine optimization', Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, ACM, New York, pp. 5-16, doi: 10.1145/3637528.3671900; preprint arXiv:2311.09735.
4. Chen, M., Wang, X., Chen, K. and Koudas, N. (2025) 'Generative engine optimization: how to dominate AI search', arXiv preprint arXiv:2509.08919.
5. Chen, X., Wu, H., Bao, J., Chen, Z., Liao, Y. and Huang, H. (2025) 'Role-augmented intent-driven generative search engine optimization', arXiv preprint arXiv:2508.11158.
6. Chen, Q., Chen, J., Huang, H., Shao, Q., Chen, J. and Hua, R. et al. (2025) 'Beyond keywords: driving generative search engine optimization with content-centric agents', arXiv preprint arXiv:2509.05607.
7. Ma, L., Qin, J., Xu, X. and Tan, Y. (2025) 'When content is goliath and algorithm is david: the style and semantic effects of generative search engine', arXiv preprint arXiv:2509.14436.
8. Wang, Y., Zhao, S., Kang, Y. and Xiong, C. (2025) 'What generative search engines like and how to optimize web content cooperatively', arXiv preprint arXiv:2510.11438.
9. Yang, H., Lyu, H., Zhang, T., Wang, D. and Zhao, Y. (2025) 'LLM-driven e-commerce marketing content optimization: balancing creativity and conversion', arXiv preprint arXiv:2505.23809.
10. Rosu, M. (2025) 'Authority vs. topical authority (relevancy) in the world of AI', EU Fire, doi: 10.47743/eufire-2025-1-27.
11. Chodak, G. and Blazyczek, K. (2024) 'Large language models for search engine optimization in e-commerce', Communications in Computer and Information Science, Springer, pp. 333-344, doi: 10.1007/978-3-031-56700-1_27.
12. Ziakis, C. and Vlachopoulou, M. (2024) 'Artificial intelligence's revolutionary role in search engine optimization', Springer Proceedings in Business and Economics, pp. 391-399, doi: 10.1007/978-3-031-51038-0_43.
13. Macumber, R. and Pagadala, V. D. E. (2024) 'Google, SEO and helpful content: how artificial intelligence can be helpful for e-commerce websites', Journal of Digital & Social Media Marketing, Vol. 12, No. 3, p. 206.

- 14.** Vadlapati, P. (2024) 'AutoTrendyKeywords: real-time AI-driven trend-based SEO using LLMs', Preprints, doi: 10.20944/preprints202410.1164.v1.
- 15.** Foglift (2026) 'AI search ranking factors' (compiling SE Ranking, Onely, The Digital Bloom and B2B buyer surveys), available at <https://foglift.io/blog/ai-search-ranking-factors> (accessed 14th June, 2026).
- 16.** Ahrefs (2025-2026) AI search citation studies (brand mentions and AI Overview source overlap), reported in Position Digital, 'AI SEO statistics', available at <https://www.position.digital/blog/ai-seo-statistics/> (accessed 14th June, 2026).
- 17.** Profound (2025) 'AI platform citation patterns: how ChatGPT, Google AI Overviews and Perplexity source information', available at <https://www.tryprofound.com/blog/ai-platform-citation-patterns> (accessed 14th June, 2026).
- 18.** The Digital Bloom (2026) 'AI citation position and revenue report', reported in Foglift, 'AI search ranking factors', available at <https://foglift.io/blog/ai-search-ranking-factors> (accessed 14th June, 2026).
- 19.** Averi and Superlines (2026) cross-platform AI citation overlap analyses, reported in Authority Tech, available at <https://authoritytech.io/curated/ai-citation-11-percent-platform-overlap-per-engine-audit-2026> (accessed 14th June, 2026).
- 20.** Leapd (2026) 'How ChatGPT, Google AI Overviews and Perplexity source information in 2026' (compiling seoClarity and Similarweb data), available at <https://www.leapd.ai/blog/ai-visibility/how-chatgpt-google-ai-overviews-and-perplexity-source-information-in-2026> (accessed 14th June, 2026).
- 21.** No Hacks (2026) 'The AI user-agent landscape in 2026: a complete reference', available at <https://nohacks.co/blog/ai-user-agents-landscape-2026> (accessed 14th June, 2026); see also Google's stated non-support for llms.txt, July 2025.
- 22.** Bal, A. (2026) 'CITED: the growth operating system for AI search', June 2026 (accessed 14th June, 2026)