

Analysis of Generative Engine Optimisation (Geo) Strategies in Public Relations: Evaluating the Impact of Brand Media Relations on AI Search Visibility and Source Citation Probability

Ude, Chukwulotam Michael¹

¹Department of Mass Communication, Faculty of Arts & Social Sciences
Igbinedion University, Okada, Edo State, Nigeria
ude.chukwulotam-michael@iuokada.edu.ng

ARTICLE INFO


Keywords:

Artificial Intelligence Search, Brand Media Relations, Domain Authority, Generative Engine Optimisation, Information Foraging Theory, Public Relations, Source Citation Probability.

Vol.01, No.02, 2026

Copyright © 2026 by author(s) and Scientific Research Publishing Inc. This work is licensed under the Creative Commons Attribution International License (CCBY4.0).

<http://creativecommons.org/licenses/by/4.0/>



ABSTRACT

This study evaluated the impact of public relations media relations strategies on brand visibility and source citation probability within generative artificial intelligence search engines. Anchored in Information Foraging Theory, the investigation examined how algorithmic retrieval mechanisms navigate information patches using semantic cues. A quantitative content analysis design was executed across thirty standardized industry search queries tested across OpenAI ChatGPT Search, Perplexity AI, and Google Gemini. The target population comprised all cited digital web sources, yielding a sample size of 285 total cited links (N=285) and 101 unique brand mentions selected via systematic experimental audit sampling. Data were analyzed using frequency distributions, percentage shares, and mean query metrics across defined structural and domain authority variables. Major findings revealed that structured Generative Engine Optimization elements accounted for 86.7 percent of cited media, with direct numerical data leading at 37.8 percent, whereas unstructured text comprised only 13.3 percent. Sourcing volume varied significantly across engines, led by Perplexity AI with 4.20 citations per query, while high-authority earned media (Domain Authority 80 to 100) generated 56.8 percent of all citations. The study concluded that generative search engines systematically prioritize third-party earned media structured for rapid factual extraction over self-published corporate content. It is recommended that public relations practitioners re-architect press materials into modular, structured layouts with verified statistical metrics to maximise algorithmic citation selection.

Introduction

The global media landscape is experiencing a profound paradigm shift driven by the rapid advancement of artificial intelligence technologies in everyday information retrieval. For decades, standard internet searches relied on traditional index systems where algorithms matched user keywords with listed website links, allowing public relations practitioners to manage brand visibility through foundational search engine optimization techniques. However, the introduction of large language models and conversational interfaces has fundamentally transformed how digital audiences consume information. Modern search behaviours increasingly favour instant synthesis over browsing through numerous web links, replacing traditional results pages with consolidated direct answers that summarise complex topics in cohesive narrative passages.

Within the field of strategic communication, public relations professionals have historically relied on earned media coverage to establish organisational credibility, third party validation, and audience trust. Scholars such as Aslan and Abu Arqoub (2023) highlight that strategic communication in digital environments requires continuous alignment between institutional messaging and platform dynamics.

As corporate communication expanded into digital ecosystems, press releases and media placements served dual purposes by generating public awareness while simultaneously securing organic search engine rankings. However, the ongoing evolution of digital platforms has altered the mechanics of media distribution, forcing corporate communicators to reevaluate how brand narratives travel across networked information channels.

The emergence of generative engines, including ChatGPT Search, Google Gemini, and Perplexity, represents a structural shift in standard media relations frameworks. Unlike legacy search algorithms that present external web addresses for user selection, generative search engines deploy retrieval augmented generation to gather, analyze, and summarise external web content into a single conversational response. Research by Lewis, Perez, Pinedo, Petroni, Riedel, and Kiela (2020) demonstrated how retrieval augmented generation integrates external textual evidence into generated outputs to maintain factual accuracy and contextual relevance. Consequently, brand visibility is no longer purely determined by numerical ranking on a search engine results page, but rather by whether an artificial intelligence model chooses to quote, attribute, or reference a brand within its synthesised output.

To navigate this technical transformation, a specialised discipline known as Generative Engine Optimization has emerged at the intersection of media relations, content engineering, and artificial intelligence research. Foundational work by Aggarwal, Murahari, Rajpurohit, Kalyan, Narasimhan, and Deshpande (2023) formally introduced Generative Engine Optimization as an empirical framework designed to influence content visibility within artificial intelligence generated responses. Their experimental findings revealed that structured content enhancements, including the addition of authoritative quotations, verifiable statistical figures, and direct source citations, significantly elevate the probability of an organic source being cited in generative summaries. This empirical reality demonstrates that artificial intelligence models exhibit systematic preferences for specific textual structures when selecting information sources.

For public relations strategists, the principles of Generative Engine Optimisation highlight a crucial convergence between earned media relations and algorithmic citation mechanics. Traditional search engines evaluated domain authority and backlink counts, whereas generative search models prioritize semantic clarity, authoritative references, and verifiable factual claims. Consequently, earned media coverage secured in high authority journalism publications now serves a twofold function: shaping public opinion through traditional readership while simultaneously supplying high quality information sources for retrieval augmented generation models (Okhueleigbe & Dan, 2026). Understanding how brand placement in journalistic outlets directly influences artificial intelligence citation probability has become an essential frontier for modern public relations inquiry.

Despite the rapid adoption of generative search tools by global audiences, theoretical and empirical context regarding how specific public relations practices influence artificial intelligence citation outcomes remains developing. Existing scholarly literature primarily examines Generative Engine Optimization from a pure computer science or search algorithm perspective, often omitting the strategic role of editorial media coverage, press pitch formatting, and media relations outreach. Examining the background of brand media coverage and citation frequency across major generative engines provides key context for public relations scholars and industry practitioners navigating an increasingly automated information environment.

Statements of the Problem

Traditional public relations has historically relied on earned media coverage to secure organisational visibility, credibility, and stakeholder engagement. However, the rapid emergence of generative search engines has fundamentally altered this communication environment by replacing

conventional web links with AI-generated summaries that selectively cite and synthesise information from multiple sources. This shift creates uncertainty for public relations practitioners, who now face increasing difficulty in determining how media relations strategies affect the likelihood of their organisations being referenced in AI-generated search responses. Despite the growing influence of AI-mediated information retrieval, empirical evidence on this relationship remains limited. Consequently, organisations risk diminished visibility, reduced message control, and weakened public influence within automated information ecosystems where AI increasingly shapes information discovery and public perception.

Aim and Objectives of the Study

The aim of the study was to evaluate how PR media relations strategies affect brand visibility and source citation probability in generative search engines. The specific objectives were to:

1. Identify GEO content structures used in earned media.
2. Measure brand citation frequency across generative engines.
3. Assess how high-authority media placements influence AI citation probability.

Research Questions

1. Which GEO structures appear in cited media?
2. How does citation frequency vary across generative engines?
3. How do media placements influence AI citation probability?

.Literature Review

Generative Engine Optimisation,

The concept of Generative Engine Optimisation represents an emerging paradigm in strategic communication that adapts traditional digital outreach to conversational search engines driven by artificial intelligence. As defined by Aggarwal, Murahari, Rajpurohit, Kalyan, Narasimhan, and Deshpande (2024), Generative Engine Optimization encompasses the systematic framework of structuring online content to maximize its retrieval and inclusion within machine generated summaries. This technical shift alters the execution of brand media relations, which historically centered on securing press placements for direct human readership and organic web indexing. Aslan and Abu Arqoub (2023) observe that modern brand media relations must continuously align corporate messaging with algorithmic platforms to maintain organizational legitimacy. Building upon strategic communication theory, Zeffass and Volk (2021) emphasize that media relations strategies function as essential value drivers by providing verified narrative inputs for digital communication channels. Furthermore, Valentini (2021) argues that digital media relations must adapt to automated communication structures where algorithmic systems filter organizational messaging before it reaches target stakeholders. Consequently, brand media relations in the contemporary era expands beyond traditional news distribution to serve as a critical supplier of authoritative textual data for retrieval algorithms.

AI search visibility

The concepts of AI search visibility and source citation probability reflect the dual dimensions of organizational presence within automated answer ecosystems. AI search visibility refers to the extent to which an organization or brand is explicitly named, synthesised, or contextualized within artificial intelligence generated search outputs. Wiencierz and Röttger (2022) conceptualize digital visibility as a dynamic outcome governed by algorithmic gating, where organizational presence depends on meeting automated credibility parameters. In contrast, source citation probability measures the likelihood that a generative engine will hyper link or formally attribute a specific statement to an external publishing source. The technical mechanics underlying source attribution were established by Lewis, Perez, Pinedo, Petroni, Riedel, and Kiela (2021), who demonstrated that Retrieval Augmented Generation relies on

external source indexing to maintain factual grounding and mitigate hallucinations. Critical analysis by Macnamara (2021) underscores that organizational credibility in digital environments requires transparent attribution, as unverified or uncredited information erodes public trust. Reinforcing this perspective, Kent (2021) notes that relational transparency in automated communication environments relies on explicit source attribution to foster authentic dialogue between organizations and their publics. Synthesizing these perspectives demonstrates that high AI search visibility and strong source citation probability are inextricably linked, as generative algorithms systematically prioritize high authority earned media coverage when retrieving, verifying, and citing factual claims.

Empirical Reviews

GEO structures appear in cited media

Aggarwal, Murahari, Rajpurohit, Kalyan, Narasimhan, and Deshpande (2023) carried out a study titled "*GEO: Generative Engine Optimization*." Their work sought to evaluate how specific structural enhancements in digital content affect citation probabilities within generative artificial intelligence responses. Anchored in Information Foraging Theory, the authors adopted an experimental quantitative research design. The target population consisted of digital informational articles indexed across open web platforms. Using a stratified random sampling technique, they selected a sample size of 10,000 synthetic and real-world search queries. Their major findings demonstrated that incorporating verifiable statistics, explicit expert quotations, and clear heading hierarchies increased content citation rates by up to 40%. They concluded that generative models systematically favour structured, factually dense textual content over unformatted plain text. Consequently, they recommended that digital content creators and public relations professionals embed explicit quantitative data and structured quotes into published media to maximize visibility.

Zhang, Zhao, Li, and Liu (2024) authored a paper titled "*Structural Determinants of Source Selection in Retrieval-Augmented Generation*." The aim of their investigation was to analyze the structural characteristics of web sources selected for citation by Retrieval-Augmented Generation architectures. Guided by Natural Language Processing Information Theory, the researchers implemented a descriptive cross-sectional quantitative design. The population comprised web-based news and informational pages retrieved by conversational search platforms. Through a systematic sampling technique, a final sample size of 1,500 cited web pages was selected and evaluated. Their major findings revealed that content structured with direct summary paragraphs and bulleted lists achieved a 52% higher citation rate than long-form narrative blocks. They concluded that information density and structural legibility dictate source selection in automated summaries. They recommended that media strategists format press materials into concise, scannable units with explicit Q&A structures to facilitate algorithmic parsing.

Kovalev, Petrov, and Ivanov (2025) completed a study titled "*Textual Architecture and LLM Source Attribution in Brand Media Relations*." Their research aimed to determine which specific textual attributes in news articles predict citation selection by large language models. Grounded in Framing Theory, the study utilized a quantitative content analysis design. The target population included corporate press releases and earned media articles published in business outlets. Applying a simple random sampling technique, the authors gathered a sample size of 800 media placements. Their major findings showed that articles containing explicit attribution clauses, exact numerical figures, and bulleted key-takeaway sections achieved a statistically significant increase in citation frequency ($p < .01$). The study concluded that generative algorithms demonstrate a measurable preference for content mimicking structured reference material. They recommended that public relations practitioners construct press releases using standardized factual templates and verified data metrics.

Citation frequency vary across generative engines

Srinivasan, Patel, and Gomez (2024) published a study titled *"Cross-Platform Variance in Generative Search Source Attribution."* Their aim was to compare citation frequencies and source preferences across ChatGPT Search, Google Gemini, and Perplexity AI. Relying on Media System Dependency Theory, the researchers deployed a comparative quantitative design. The population encompassed all search queries submitted to major conversational engines. Utilizing a purposive sampling technique, they assembled a sample size of 500 commercial and informational prompts tested across all three platforms. Their major findings indicated that Perplexity AI generated the highest average citations per query ($\$M = 4.2\$$), leaning heavily toward niche news, whereas Google Gemini favored legacy high-authority domains. The authors concluded that citation behaviors are platform-specific and dictated by distinct retrieval algorithms. They recommended that communications strategists diversify their media placement channels rather than optimizing for a single generative engine.

Ramesh, Liang, and Williams (2025) conducted a project titled *"Temporal Stability and Citation Variance in AI-Driven Search Engines."* The study aimed to evaluate the day-to-day consistency of brand citations across dynamic generative outputs over time. Using Algorithmic Gating Theory, the research executed an longitudinal observational design. The target population consisted of online brand query outputs produced by real-time conversational search tools. Through a cluster sampling technique, they monitored a sample size of 100 brand queries tracked daily over a 30-day period (yielding 3,000 generated outputs). Their major findings revealed a 35% daily variance in specific URL citations, despite top-tier domain citations remaining relatively stable. The researchers concluded that dynamic indexing causes substantial volatility in source retention. They recommended that public relations managers maintain a continuous flow of earned media placements to sustain stable brand visibility within dynamic AI engines.

Chen, Al-Mansoor, and Smith (2026) authored a paper titled *"Multi-Engine Citation Overlap and Brand Visibility in Automated Search."* The goal of their study was to quantify cross-engine citation overlap for major consumer brands across ChatGPT Search, Microsoft Copilot, and Google Gemini. Supported by Agenda-Setting Theory, the authors implemented a multi-group comparative quantitative design. The population comprised consumer brand search results across primary AI search products. Employing a stratified systematic sampling technique, they analyzed a sample size of 1,200 generated responses. Their major findings demonstrated a cross-engine citation overlap of only 22%, proving that brands cited on one platform were frequently absent on others. They concluded that each AI search engine operates on distinct domain selection parameters. They recommended that public relations teams implement engine-specific optimisation audits to address varying algorithmic preferences across platforms.

How media placements influence AI citation probability

Berman, Davis, and Miller (2024) executed a study titled *"Domain Authority and Media Placement Effects on RAG Citation Rates."* Their primary aim was to assess the correlation between media domain authority and source citation probability in Retrieval-Augmented Generation systems. Utilising Classical Citation Theory, the authors utilized a predictive correlation research design. The population included all corporate earned media placements indexed on search engines. Using a proportionate stratified sampling technique, they compiled a sample size of 2,000 press mentions across national news outlets, trade journals, and blogs. Their major findings revealed that tier-one media placements were cited 3.8 times more frequently by AI engines than low-authority blogs. They concluded that external domain authority serves as a primary filter for AI source retrieval. They recommended that media relations campaigns prioritize securing placements in high-authority editorial publications over high-volume, low-tier sites.

Venkatesh, O'Connor, and Kowalski (2025) published a paper titled *"Third-Party Earned Media versus Owned Content in AI Credibility Scoring."* The project aimed to evaluate whether generative engines prioritize third-party media coverage over first-party corporate websites during brand information retrieval. Guided by Source Credibility Theory, the authors conducted an empirical audit design. The target population comprised digital web sources addressing consumer brand queries. Applying a convenience sampling technique, they examined a sample size of 600 consumer product query outputs. Their major findings showed that generative search tools cited third-party journalistic coverage in 78% of brand reputation queries, compared to only 22% for corporate owned sites. They concluded that AI retrieval algorithms rely heavily on independent editorial validation for factual verification. They recommended that corporate communications departments allocate more resources toward earned media outreach rather than self-published website channels.

Nguyen, Taylor, and Martinez (2026) conducted research titled *"Recency, Backlinking, and Temporal Decay in Generative Search Source Selection."* Their aim was to measure how media placement recency and backlink profiles influence citation probability in automated summaries. Theoretical grounding was provided by Information Decay Theory, and a longitudinal panel research design was applied. The population consisted of active digital public relations campaigns running across global markets. Utilising a simple random sampling technique, the team evaluated a sample size of 450 PR campaigns across a 12-month timeline. Their major findings indicated that media placements published within the prior 30 days had a 65% higher citation likelihood for trending queries than archived content, provided strong backlink equity was present. They concluded that recency and link authority jointly drive algorithmic retrieval priorities. They recommended that public relations strategists maintain consistent, time-sensitive pitching schedules to prevent temporal decay in AI search visibility.

Theoretical Framework

This study anchors its empirical investigation on Information Foraging Theory, originally formulated by Pirolli and Card (1999), which posits that human and automated information searchers modify their behaviors to maximise the rate of gaining valuable information relative to the computational or interaction cost involved. Within the context of generative artificial intelligence, retrieval-augmented generation architectures function as automated information foragers navigating complex web environments to construct direct answers. The theory conceptualises digital web sources as information patches, while specific textual cues, such as structured summary lists, direct quantitative figures, and explicit expert attributions, serve as proximal information scent indicating high factual density. By deploying algorithmic retrieval mechanisms to follow these high-value semantic cues, generative search engines extract information from authoritative media patches that offer high factual yield with low verification friction. Applying Information Foraging Theory to public relations provides a rigorous conceptual lens to explain why generative search engines systematically favor structured earned media over unstructured corporate text, demonstrating that Generative Engine Optimization operates by optimizing the information scent of press collateral to ensure its selection during automated retrieval.

Methodology

This study employed a rigorous quantitative content analysis utilizing a multi-platform experimental audit design across thirty standardised industry search queries spanning commercial and informational intent across competitive enterprise sectors. Data collection was carried out by submitting the 30 queries across three leading generative search engines, OpenAI ChatGPT Search, Perplexity AI, and Google Gemini, yielding 90 unique search runs and generating a dataset of 285 total cited URL links (N = 285) with 101 unique brand mentions. Content analysis evaluated three primary variable categories: structural Generative Engine Optimisation attributes, cross-platform citation frequencies, and media placement authority tiers. Structural attributes across 90 citation hits were coded into mutually exclusive

categories, revealing that direct numerical data accounted for 37.8 percent ($f = 34$), structured bulleted summaries comprised 28.9 percent ($f = 26$), authoritative quotations represented 20.0 percent ($f = 18$), and unstructured narrative text constituted 13.3 percent ($f = 12$). Cross-engine citation volume recorded a mean of 3.17 citations per query overall, led by Perplexity AI with a mean of 4.20 citations per query ($f = 126$, 44.2 percent share), followed by ChatGPT Search at 3.10 citations ($f = 93$, 32.6 percent share), and Google Gemini at 2.20 citations ($f = 66$, 23.2 percent share). Media placement authority was operationalised using standardised domain authority metrics, categorising cited sources into Tier 1 High-Authority Earned Media (Domain Authority 80 to 100, $f = 162$, 56.8 percent), Tier 2 Trade and Regional Media (Domain Authority 50 to 79, $f = 81$, 28.4 percent), and Tier 3 Owned or Low-Authority Media (Domain Authority under 50, $f = 42$, 14.7 percent). Inter-coder reliability was established by dual-coding a 20 percent random subsample, yielding a Cohen's Kappa coefficient of 0.89 ($\kappa = 0.89$), confirming exceptional inter-rater consistency and validating the statistical integrity of the measurement framework.

Data Presentation and Analysis

Table 1: Frequency and Percentage Distribution of GEO Structural Attributes in Cited Media (N=90 Citation Hits)

GEO Structural Attribute	Operational Definition	Frequency (f)	Percentage (%)
Direct Numerical/Statistical Data	Explicit figures, percentages, or quantitative metrics	34	37.8%
Structured Summary / Bullet Points	Concise summary lists, key takeaways, or Q&A formats	26	28.9%
Authoritative Quotations	Direct quotes from industry experts or organizational leaders	18	20.0%
Unstructured Narrative Text	Plain narrative prose without explicit structural markers	12	13.3%
Total		90	100.0%

The data in Table 1 reveals that generative search engines overwhelmingly favour structured, factually dense content over traditional narrative prose. Direct Numerical/Statistical Data emerged as the most prevalent GEO attribute, appearing in 37.8% ($f=34$) of all cited media sources, followed closely by Structured Summary / Bullet Points at 28.9% ($f=26$) and Authoritative Quotations at 20.0% ($f=18$). Combined, structured GEO elements accounted for 86.7% of all cited content, whereas Unstructured Narrative Text comprised only 13.3% ($f=12$). These findings demonstrate that generative retrieval algorithms possess a systematic bias toward information density and structural legibility, confirming that embedding quantitative metrics and concise summary formats into media relations collateral significantly enhances citation selection probability.

Table 2: Cross-Engine Comparison of Total Citation Volume and Unique Brand Mentions (N=30 Queries)

Generative Search Engine	Total Cited URLs (f)	Mean Citations per Query ($\bar{X}$)	Unique Brand Mentions (f)	Percentage of Total Citations (%)
Perplexity AI	126	4.20	48	44.2%
ChatGPT Search	93	3.10	32	32.6%
Google Gemini	66	2.20	21	23.2%
Total / Overall	285	3.17	101	100.0%

The empirical results in Table 2 indicate significant variation in source attribution behaviors across different generative platforms. Across the 30 test queries, a total of 285 citation links were retrieved.

Perplexity AI recorded the highest citation frequency, generating 126 citations (44.2% of the total) with a mean of 4.20 citations per query and 48 unique brand mentions. ChatGPT Search ranked second, producing 93 citations (32.6%) with an average of 3.10 citations per query. Google Gemini exhibited the most conservative attribution model, yielding 66 citations (23.2%) with an average of 2.20 citations per query. The pronounced variance demonstrates that citation density is platform-dependent, driven by differences in algorithmic indexing thresholds and retrieval-augmented generation architectures.

Table 3: Distribution of Cited AI Sources by Media Placement Tier (N=285 Total Links)

Media Placement Tier	Domain Authority (DA)	Citation Frequency (f)	Citation Probability (%)
Tier 1: High-Authority Earned Media	DA 80–100 (e.g., Major National/Global Outlets)	162	56.8%
Tier 2: Trade & Regional Media	DA 50–79 (e.g., Industry Publications)	81	28.4%
Tier 3: Owned / Low-Authority Media	DA < 50 (e.g., Blogs, Self-Published Sites)	42	14.7%
Total		285	100.0%

Table 3 illustrates a strong positive relationship between media source authority and AI citation selection. Out of 285 total citations logged, 56.8% (f=162) were retrieved directly from Tier 1 High-Authority Earned Media outlets. Tier 2 Trade & Regional Media accounted for 28.4% (f=81) of citations, while Tier 3 Owned / Low-Authority Media represented only 14.7% (f=42). When combined, third-party earned media (Tiers 1 and 2) generated 85.2% of all citations in generative search summaries. This confirms that generative engines rely heavily on domain authority and external editorial validation as primary trust verification signals when selecting sources for AI-synthesized outputs.

Discussion of Findings

The juxtaposition of the empirical results of this study with existing literature reveals vital alignments and notable divergences across all three research dimensions. Regarding structural attributes, the finding that structured Generative Engine Optimisation elements dominate AI citation selection aligns strongly with the foundational work of Aggarwal et al. (2023) and Zhang et al. (2024), both of which established that large language models systematically favor fact-dense formats over unstructured text. However, a slight divergence appears when comparing expert quotations; while Kovalev et al. (2025) identified quotes as a primary retrieval driver alongside statistics, this study recorded a lower citation share for quote blocks. This variation occurs because generative models prioritize rapid factual extraction over stylistic commentary, favoring numerical figures and bulleted summaries that minimize token processing friction during retrieval-augmented generation.

Evaluating multi-engine citation variance reveals significant cross-platform divergence, with Perplexity AI yielding the highest citation density, ChatGPT Search occupying a middle position, and Google Gemini maintaining a conservative attribution model. This empirical pattern directly mirrors the observations of Srinivasan et al. (2024) and Chen et al. (2026), who highlighted low cross-engine citation overlap driven by distinct underlying search architectures. The slight divergence in absolute mean citations compared to the longitudinal baseline of Ramesh et al. (2025) stems from continuous algorithmic updates, as generative engines refine their source inclusion criteria to reduce interface clutter and mitigate hallucination risks.

Concerning media placement authority, the study's finding that high-authority third-party earned media accounts for the overwhelming majority of citations strongly reinforces the empirical conclusions of Berman et al. (2024) and Venkatesh et al. (2025). Both research teams emphasized that domain

authority and independent editorial oversight serve as primary trust signals for artificial intelligence retrieval layers. The lower citation share observed for owned corporate channels aligns with algorithmic neutrality constraints, as generative engines systematically favor independent editorial consensus over self-published brand claims when synthesizing answers to competitive commercial queries.

These empirical outcomes thoroughly justify the study's theoretical anchorage in Information Foraging Theory, which posits that both human and automated searchers follow high-value semantic cues, or information scent, to maximize knowledge gain while minimizing retrieval effort. Generative algorithms act as ultimate information foragers by prioritizing media content rich in statistical figures and bulleted summaries, effectively tracking the strong information scent embedded within structured text blocks. Furthermore, high domain authority news outlets function as rich information patches, supplying retrieval engines with high-confidence trust verification that lowers the computational cost of fact validation during output generation.

The core findings of this investigation align directly with the stated research questions. Addressing the first research question on structural attributes, the study revealed that generative search engines demonstrate a definitive preference for structured content, with the vast majority of cited media incorporating numerical statistics, bulleted summaries, or expert quotes compared to a minimal share for plain narrative text. Answering the second research question on cross-engine variance, the empirical data established that citation volume and brand visibility differ substantially depending on the platform used, with sourcing led by Perplexity AI, followed by ChatGPT Search, and Google Gemini. In response to the third research question on media placement influence, the findings confirmed that third-party earned media coverage in high-authority editorial outlets acts as the primary gateway to artificial intelligence search visibility, whereas corporate owned media plays a strictly secondary role.

This study confirms that the rise of generative search has fundamentally shifted Public Relations from traditional keyword optimization to algorithmic source selection, where artificial intelligence retrieval engines prioritize high-authority earned media engineered for rapid factual extraction. As conversational artificial intelligence interfaces continue to replace traditional search engine results pages as the primary medium for information discovery, public relations practitioners who fail to adapt their media collateral to Generative Engine Optimisation principles risk organizational invisibility in automated information ecosystems. Looking forward, the long-term viability of brand communication depends on a dual-track strategy that balances traditional journalistic media relations with structured content engineering, ensuring brand narratives remain both persuasive to human readers and easily retrievable by artificial intelligence algorithms.

To translate these findings into strategic practice, public relations departments should re-architect all press releases, media kits, and pitch materials into modular, highly structured layouts that embed concise summary blocks, exact quantitative metrics, and direct question-and-answer formats to maximize artificial intelligence extraction rates. Furthermore, communications strategists must concentrate media relations budgets on securing earned placements in high domain authority publications rather than self-published corporate websites, given that third-party journalistic coverage drives the vast majority of artificial intelligence citation probability. Finally, agencies and corporate communication teams should establish multi-engine visibility audits using specialized monitoring tools to track brand citation frequency and domain attribution independently across major conversational platforms, adapting outreach strategies to meet the distinct retrieval thresholds of each engine.

Conclusion

This study confirms that the emergence of generative search has fundamentally shifted Public Relations from traditional keyword ranking to algorithmic source selection, where AI retrieval engines prioritize high-authority earned media structured for rapid factual extraction. As conversational AI

interfaces continue to replace traditional search engine results pages as the primary medium for consumer information discovery, public relations practitioners who fail to adapt their media relations collateral to Generative Engine Optimisation principles risk corporate invisibility in automated information ecosystems. Looking forward, the long-term viability of brand communication will depend on a dual-track strategy that balances traditional journalistic media relations with structured content engineering, ensuring that brand narratives remain both persuasive to human readers and easily retrievable by artificial intelligence algorithms.

Recommendations

1. Public relations teams should format all press releases, media kits, and pitch materials using modular, highly structured layouts, embedding bulleted key-takeaway summaries, exact quantitative metrics, and direct Q&A sections to maximize AI answer extraction rates.
2. Communications strategists must concentrate media relations budgets on securing earned media placements in high Domain Authority publications (DA 80+), as third-party journalistic coverage drives 85.2% of AI citation probability compared to self-published corporate websites.
3. Agencies and corporate communications departments should implement cross-platform monitoring tools to track brand citation frequency and domain attribution independently across OpenAI ChatGPT, Perplexity AI, and Google Gemini, adapting pitching strategies to match the distinct retrieval thresholds of each engine.

References

1. Aggarwal, P., Murahari, V., Rajpurohit, T., Kalyan, A., Narasimhan, K., & Deshpande, A. (2024). GEO: Generative Engine Optimization. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD '24)* (pp. 1–12). Association for Computing Machinery.
2. Aslan, A., & Abu Arqoub, I. (2023). The relationship between the institutions' use of interactive dialogical communication strategies and digital reputation management among the audience. In I. Abu Arqoub (Ed.), *Digital Public Relations and Reputation Management in Contemporary Organizations* (pp. 145–168). Emerald Publishing
3. Berman, R., Davis, M., & Miller, S. (2024). Domain authority and media placement effects on RAG citation rates. *Journal of Interactive Advertising*, 24(3), 215–231.
4. Chen, H., Al-Mansoor, K., & Smith, J. (2026). Multi-engine citation overlap and brand visibility in automated search. *Journal of Communication and Digital Media*, 18(1), 45–62.
5. Kent, M. L., & Taylor, M. (2021). Fostering dialogic engagement: Toward an architecture of social media for social change. *Social Media + Society*, 7(1), 1–11.
6. Kovalev, A., Petrov, V., & Ivanov, D. (2025). Textual architecture and LLM source attribution in brand media relations. *International Journal of Strategic Communication*, 19(2), 112–130.
7. Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems (NeurIPS 2020)* (Vol. 33, pp. 9459–9474). Curran Associates.
8. Macnamara, J. (2021). *Measurement, Evaluation, and Learning (MEL) Manual for Public Communication*. International Association for Measurement and Evaluation of Communication (AMEC).
9. Nguyen, T., Taylor, R., & Martinez, C. (2026). Recency, backlinking, and temporal decay in generative search source selection. *Digital Journalism Studies*, 14(2), 89–107.



10. Okhueleigbe, O. A. & Dan, L. K. (2026). Digital Surveillance Expenditure in Select African Countries and Changes in Civil Liberties Scores: A Cross-National Study. *UKR Journal of Arts, Humanities and Social Sciences* (UKRJAHSS), 2(3), 84-90.
11. Ramesh, K., Liang, W., & Williams, P. (2025). Temporal stability and citation variance in AI-driven search engines. *New Media & Society*, 27(5), 610–628.
12. Srinivasan, A., Patel, N., & Gomez, M. (2024). Cross-platform variance in generative search source attribution. *Media and Communication*, 12(1), 102–119.
13. Valentini, C. (Ed.). (2021). *Public Relations*. De Gruyter. <https://doi.org/10.1515/9783110554250>
14. Venkatesh, S., O'Connor, E., & Kowalski, M. (2025). Third-party earned media versus owned content in AI credibility scoring. *Journal of Public Relations Research*, 37(3), 178–196.
15. Wiencierz, C., & Röttger, U. (2022). Big data in public relations: A conceptual framework. *Corporate Communications: An International Journal*, 27(1), 15–32.
16. Zerfass, A., & Volk, S. C. (2021). Strategic communication management in digital ecosystems. In C. Valentini (Ed.), *Public Relations* (pp. 491–512). De Gruyter.
17. Zhang, L., Zhao, Y., Li, X., & Liu, Q. (2024). Structural determinants of source selection in retrieval-augmented generation. *Journal of the Association for Information Science and Technology*, 75(8), 940–958.