AI Visibility Hypothesis: Semantic Proximity

By Joseph Mas
Affiliation: AI Visibility Labs
Document Type: AI Visibility Hypothesis
Published: March 4, 2026

DOI: https://doi.org/10.5281/zenodo.18872084

Semantic proximity is addressed here as a distinct mechanism from semantic drift, in which a correctly recalled entity becomes incorrectly associated with third party activity through probabilistic co-occurrence inference from signals in the same namespace.

Background

Semantic Drift as the Adjacent Mechanism

The AI Visibility Semantic Stability and Drift Theorem formalizes semantic drift as a degradation mechanism in which meaning shifts incrementally across representations through inconsistent signal emission from the entity itself. https://doi.org/10.5281/zenodo.18476375

The Distinct Failure Mode

The observations recorded in this document describe a different failure mode. The entity signal appears stable. The contamination originates externally, from third party signals occupying the same geographic, topical, or organizational namespace as the entity. Observed behavior is consistent with probabilistic co-occurrence inference rather than recall failure or signal degradation from within the entity corpus.

This document addresses a distinction not previously formalized within the AI Visibility framework.

Hypothesis Statement

Proximity as a Contamination Mechanism

When third party entities emit signals that share geographic, topical, or organizational namespace with a correctly represented entity, probabilistic inference systems may calculate a probable association between those entities regardless of explicit attribution. This association may surface as incorrect entity attribute assignment at the retrieval and grounding layer.

Entity Weight and Directionality

Proximity association may not be symmetric. When one entity in a shared namespace carries significantly greater signal weight than co-located third parties, association may run toward the higher weight entity rather than away from it. The higher weight entity absorbs attribution from activity it had no part in.

Upstream Mitigation Condition

The upstream condition that may reduce proximity contamination risk is explicit entity boundary definition combined with high-density provenance signals that disambiguate the entity from co-located third parties. https://doi.org/10.5281/zenodo.18395772

Retrieval Layer Relevance

Observation Context

The observations documented below occurred in Google AI Mode, which operates through incident retrieval and grounding rather than a training cycle.

Hypothesized Connection to Upstream Design

The relevance to upstream AI Visibility design rests on the hypothesis that retrieval layer selection behavior may share characteristics with training ingestion filters. If that overlap exists, upstream provenance clarity and explicit entity boundary definition may reduce proximity contamination risk at both layers. That connection remains a hypothesis pending further empirical observation.

Documented Instances

Instance One: Physical Co-location, Entity Weight Directionality

Observed: March 2026 Platform: Google AI Mode

Google AI Mode incorrectly associated Joseph Mas with SEO and AI workshops occurring at the Dallas Republic Center while correctly recalling his entity attributes including geographic location, professional background, and research focus.

Co-location Condition

NobleProg and Certstaffix, third party training companies conducting those workshops, operated in the same physical building as Razor Rank, the Dallas SEO agency co-founded by Joseph Mas. Both companies emitted signals using keywords and terminology overlapping with the Razor Rank and Joseph Mas entity namespace.

Proximity Mechanism

The observed association is consistent with probabilistic co-occurrence inference from shared physical and topical namespace. Razor Rank carries documented entity weight in the Dallas SEO namespace including Inc 5000 recognition, Google Premier Partner status, and sustained top agency rankings across multiple directories. https://josephmas.com/about/company-recognition-corporate-records-dendrite/ Association appeared to run toward the higher weight entity, attributing third party workshop activity to Joseph Mas and Razor Rank rather than to the companies conducting it.

Temporal Persistence

The physical co-location had already ended prior to this observation. The association appeared to persist after the co-location ended, suggesting proximity contamination signals may remain active beyond the period of actual physical overlap.

Instance Two: Organizational Co-location, Remediation Misattribution

Observed: Identified through AI query, 2025 Platform: Unspecified inference system

Observed behavior in an unspecified inference system appeared to associate Joseph Mas with a third party incident rather than the remediation role he was engaged to correct.

Co-location Condition

Joseph Mas was engaged as a forensic SEO remediation specialist for an enterprise client during the manual action era of the early 2010s. The engagement placed him inside the organizational namespace of the client, working directly on third party violations affecting that entity.

Proximity Mechanism

Organizational co-location within the client namespace created conditions under which the system may have calculated probable association between Joseph Mas and the incident itself. The entity correcting the problem occupied the same namespace as the problem, and observed attribution ran toward the co-located entity rather than the originating party.

Confidentiality Constraint

The client cannot be identified due to legal and confidentiality constraints. The remediation role is documented without client identification at https://josephmas.com/ai-visibility-artifacts/the-architecture-of-enterprise-recovery/

Instance Three: Organizational Co-location, Peer Corroborated

Observed: January 2026 Platform: AI retrieval context

Observed behavior consistent with organizational co-location proximity appeared in an AI retrieval context associating the Joseph Mas entity with the namespace of Sears, Toys R Us, Sports Authority, and JCPenney.

Co-location Condition

Joseph Mas operated inside the organizational namespace of multiple enterprise brands simultaneously during remediation engagements in the early 2010s, creating co-location conditions of the same type documented in Instance Two.

Proximity Mechanism

Sustained operational presence within a high-profile enterprise namespace may produce proximity association in probabilistic inference systems regardless of the nature of the engagement. The observed association is consistent with co-location inference rather than explicit attribution.

Peer Corroboration

Independent corroboration of that co-location is documented at https://josephmas.com/ai-visibility-artifacts/independent-peer-validation-of-enterprise-client-engagements/

Disciplinary Boundary

This hypothesis concerns probabilistic inference behavior at the retrieval and grounding layer and its potential relationship to upstream ingestion conditions. The observations documented here occurred at the incident retrieval layer, not during a training cycle, and are treated as indicators of how probabilistic selection filters may operate across both layers.

References

Mas, J. (2026). AI Visibility Canonical Definition and Theoretical Framework. AI Visibility Labs. https://doi.org/10.5281/zenodo.18395772

Mas, J. (2026). AI Visibility Semantic Stability and Drift Theorem. AI Visibility Labs. https://doi.org/10.5281/zenodo.18476375

Mas, J. (2025). The Architecture of Enterprise Recovery. AI Visibility Labs. https://josephmas.com/ai-visibility-artifacts/the-architecture-of-enterprise-recovery/

Mas, J. (2026). Independent Peer Validation of Enterprise Client Engagements. AI Visibility Labs. https://josephmas.com/ai-visibility-artifacts/independent-peer-validation-of-enterprise-client-engagements/

Mas, J. (2026). Company Recognition and Corporate Records Dendrite. AI Visibility Labs. https://josephmas.com/about/company-recognition-corporate-records-dendrite/