

Brand Legibility

A Construct for Building Brand Equity in AI-Mediated Markets

Marcos Guimarães Figueira Fundação Getulio Vargas (FGV) Rio de Janeiro, Brazil

Wyse Brand Intelligence Rua do Russell, 804/4º andar, Rio de Janeiro, 22210-010 RJ, Brazil marcos.figueira@fgv.edu.br

May 2026

Abstract

Brand erasure (Figueira, 2026a, 2026b) named the failure mode of brand theory in AI-mediated markets — the systematic suppression of brand identity inside synthesized answers and completed agentic actions. This paper develops the corresponding positive construct: *Brand Legibility*, the degree to which a brand is recognizable, parseable, and attributable by the AI systems that mediate consumer experience. We position Brand Legibility as the latent-substrate counterpart of brand salience (Romaniuk & Sharp, 2004) and the operational property whose absence produces erasure. We introduce the *Semantic Equity Index* (SEI) as the quantitative measure of legibility and the *Brand Legibility Framework* as the operating model that integrates Answer Engine Optimization, Generative Engine Optimization, and Agentic Optimization as three coordinated layers. We close with the *Brand Erasure Audit* as the measurement protocol that operationalizes the framework, and with implications for brand-equity theory, marketing organizations, and the governance of attribution.

Keywords: brand legibility; semantic equity; brand erasure; generative engine optimization; agentic AI; customer-based brand equity; AI-mediated markets

1. Introduction: The Inverse of Erasure

Brand erasure named a problem. This paper names the property whose absence produces it.

When an AI assistant synthesizes a response, paraphrases a source, or executes a task on the user’s behalf, the brand can be present in the substrate the assistant draws on and absent from the experience the user receives (Figueira, 2026a; 2026b). The phenomenon is now measurable. What is not yet operationally complete is the inverse: the property a brand must have for the assistant to surface it, cite it, or transact with it. We propose to call that property *Brand Legibility*, and we argue that the construct is doing work the field’s existing vocabulary cannot.

The argument is structural. Brand awareness presumes a human encounter with the brand. Brand recognition presumes a visual or verbal cue the consumer can react to. Brand salience presumes a mental network of associations that gets activated at the moment of choice. None of these presumes anything about the machine substrate that increasingly precedes the human encounter. The substrate is not a metaphor. It is a different

kind of memory, with different dynamics, different inputs, and different conditions under which a brand exists or fails to exist within it.

This paper does five things. It defines Brand Legibility as a construct distinct from adjacent ideas in the brand-equity tradition. It identifies the three dimensions on which legibility varies. It introduces the Semantic Equity Index as the quantitative measure that converts legibility from a verbal construct into a benchmarkable, monitorable indicator. It presents the Brand Legibility Framework as the operating model that connects measurement to intervention. And it sets out the diagnostic — the Brand Erasure Audit — through which the framework is applied. The agenda from Figueira (2026b) called for measurement; this paper supplies the named instruments and the operating architecture.

The Pew Research Center’s (2025) finding that users click a traditional search result 8% of the time when an AI Overview is present, against 15% without one, has been cited often enough that its weight is in danger of being lost. The number that matters most for this paper is the 1% click rate inside the AI summary itself. That is what life inside the synthesized answer looks like for a brand. Legibility is the property that determines whether the brand was inside the answer in the first place.

The paper unfolds in ten sections after this introduction. Section 2 develops the construct. Section 3 locates legibility in the substrate. Section 4 introduces the Semantic Equity Index. Section 5 presents the framework. Section 6 details the diagnostic audit. Section 7 develops strategic implications. Section 8 discusses the discipline-level shift. Section 9 lays out the research agenda specific to legibility. Section 10 acknowledges limitations. Section 11 concludes.

2. The Construct of Brand Legibility

We define Brand Legibility as follows.

Brand Legibility is the degree to which a brand is recognizable, parseable, and attributable by the AI systems that mediate consumer experience.

Three properties are doing work in the definition. *Recognizable* means the brand exists for the system as a stable, identifiable entity — not a string of tokens that the model resolves differently on different runs. *Parseable* means the brand’s content, products, and affordances are structured in ways the system can extract without loss of identity. *Attributable* means that when the system uses the brand’s content or selects the brand’s surface, the brand’s name survives the operation. A brand that is recognized but not attributed is partially legible; a brand that is parsed but not recognized is partially legible; the construct is the conjunction.

2.1 Distinction from adjacent constructs

Brand Legibility is adjacent to several established constructs in brand theory, but it is not reducible to any of them.

Brand awareness (Keller, 1993) is the consumer’s ability to recall or recognize the brand under category-relevant cues. The substrate is human associative memory. Awareness is a necessary condition for the

consumer’s eventual choice — but in AI-mediated markets, the consumer’s choice is increasingly mediated by an intermediary whose memory is not the consumer’s. A brand can be high in consumer awareness and low in machine legibility simultaneously, because the substrates are different and the conditions that produce stability in each are different.

Brand recognition (Aaker, 1991) is the consumer’s ability to identify the brand from visual or verbal cues. The construct is sensory and human. Brand Legibility is structural and computational. A brand with strong visual recognition can have weak entity-graph presence and weak named-entity stability across models — the conditions that produce machine recognition are not the conditions that produce human recognition.

Brand salience (Romaniuk & Sharp, 2004) is the construct closest in logical structure to Brand Legibility, and the distinction is the most important to articulate. Salience is the breadth of mental availability — how often and how widely the brand comes to mind across category-relevant situations. Legibility is the breadth of machine availability — how often and how widely the brand is surfaced, cited, and transacted with by AI assistants across category-relevant prompts and tasks. The logic structure is parallel. The substrate is not. Salience lives in distributed associative memory in human consumers. Legibility lives in distributed representations in foundation models, in retrieval indexes, and in agent manifests. A brand can be high-salience among human consumers and low-legibility for AI assistants. The reverse is also possible, and is becoming common: brands with strong entity-graph presence that human consumers have never heard of show up disproportionately in AI responses because the substrate that AI assistants draw on overweights what is structured and indexed.

Discoverability in the SEO literature concerns whether content surfaces in a ranked list. Legibility concerns whether the brand survives the post-retrieval synthesis or the agentic action. A brand can be discoverable and still erased. The two constructs operate at different stages of the assistant’s procedure.

The distinctions matter because each adjacent construct has its own measurement instruments and intervention literature. Treating legibility as a synonym for any of them imports the wrong assumptions and obscures what is genuinely new about the AI-mediated substrate.

2.2 Three dimensions

Brand Legibility varies along three dimensions, each tied to a substrate the AI assistant draws on.

Entity legibility is the degree to which the brand exists as a stable, canonical entity in the structured knowledge that grounds retrieval. Operationally: presence in Wikipedia and Wikidata, structured-data coverage (schema.org markup), and stable named-entity recognition across foundation models. Entity legibility is a property of the entity-graph layer that retrievers consult.

Semantic legibility is the degree to which the brand has a well-localized vector representation in the latent space of foundation models, anchored to the attribute vectors that matter for the category. A brand with strong semantic legibility clusters reliably with the attributes the category rewards. A brand with weak semantic legibility either clusters with the wrong attributes (the legacy “discount” association in a “premium” category) or has no stable cluster at all.

Operational legibility is the degree to which the brand’s surfaces are machine-readable in ways that allow

agents to transact. Operationally: published manifests, structured product and pricing data, idempotent endpoints, clear error semantics, and explicit affordance documentation. Operational legibility is what determines whether the agent can complete the task on the brand’s surface or has to substitute.

The three dimensions are partially independent. A brand can be entity-legible (Wikipedia entry, knowledge-graph identifier) and semantically illegible (its vector clusters with the wrong attributes). A brand can be semantically legible and operationally illegible (the model knows what the brand stands for, but the agent cannot complete a purchase on the site). The construct is the conjunction; the dimensions are separately measurable and separately addressable.

2.3 What legibility is for

The construct’s value is not descriptive but operational. Legibility is the property a brand must have for the chain of mediation — retrieval, reranking, synthesis, action — to surface the brand as a brand rather than substitute its category role with a generic answer or a competitor’s surface. Legibility is what survives erasure. Building legibility is the strategic alternative to optimizing for visibility, which addresses the wrong layer of the assistant’s procedure.

3. The Substrate

Brand Legibility lives in three substrates simultaneously, and the conditions that produce legibility in each are different.

The training-corpus substrate is the body of text the foundation model learned from. Brands that were present in that corpus with consistent attribution, distinctive phrasing, and high-authority co-mentions are represented stably in the model’s latent space. Brands that were absent, paraphrased, or syndicated without attribution are represented weakly or with the wrong attribute associations. The training-corpus substrate is non-stationary — every model update changes it — but at any given moment it is the layer that determines whether the model can “think about” the brand at all.

The retrieval-index substrate is the body of structured and unstructured content the assistant consults at inference time, typically through retrieval-augmented generation (Lewis et al., 2020; Gao et al., 2024) or late-interaction retrievers (Khattab & Zaharia, 2020; Santhanam et al., 2022). A brand that is present in this substrate with rich structured data, named methodologies, and paraphrase-resistant phrasing produces retrieval results that the synthesizer has to honor with attribution. A brand whose content is generic, paraphrase-friendly, or unstructured produces retrieval results that the synthesizer can absorb into prose without attribution.

The agent-manifest substrate is the body of machine-readable affordances the agent uses to decide which platform to act on. A brand whose API surface is closer to a well-documented public API than to a website is legible at this layer. A brand whose action surface requires GUI navigation, undocumented forms, or human-only authentication is not.

The implication is that legibility cannot be diagnosed in any single substrate. A brand can be legible in

training corpora and illegible in retrieval indexes; legible in retrieval and illegible in agent manifests. A defensible measurement of Brand Legibility samples all three.

4. The Measure: Semantic Equity Index

The construct as defined is qualitative. The measure that operationalizes it is quantitative. We introduce the *Semantic Equity Index* (SEI) as the headline indicator of Brand Legibility.

4.1 Construction

The SEI is computed from three components, each tied to one of the three substrates identified in Section 3.

The *entity-graph presence score* captures the brand’s standing in structured knowledge. It is computed from the brand’s coverage in canonical public knowledge graphs (Wikipedia, Wikidata, industry-specific graphs), the completeness and consistency of its structured-data markup (schema.org Organization, Product, Service), and the stability of named-entity recognition for the brand across a representative panel of foundation models. The score is bounded and benchmarkable.

The *latent-space proximity score* captures the brand’s position in the semantic substrate. For a target brand, attribute vectors are defined for the category’s relevant dimensions — performance, reliability, premium, value, innovation, and so on, with attribute sets adapted to the category. Cosine similarities between the brand vector and each attribute vector are computed across multiple foundation models. The score aggregates these proximities, weighted by the strategic importance of each attribute to the brand’s intended positioning. The method draws on the embedding-bias literature (Bolukbasi et al., 2016; Caliskan, Bryson, & Narayanan, 2017) but reframes the technique from auditing biases to auditing brand positioning.

The *response-audit citation score* captures the brand’s outcomes in synthesized responses. A representative set of category-relevant prompts is submitted to a population of AI assistants. Responses are coded for brand-mention frequency, share of citations, sentiment of the citing context, and — for tasks rather than queries — agent task-completion rate. The codebook is adapted from algorithmic-accountability methodology (Diakopoulos, 2016) and can be made interoperable with traditional brand-equity instruments (Aaker, 1996; Yoo & Donthu, 2001) for convergent-validity testing.

4.2 Aggregation and weighting

The three components do not enter the SEI with equal weight. The weighting reflects the rate-limiting substrate for the brand’s category. Categories dominated by transactional intent — booking, ordering, scheduling — weight operational legibility more heavily because that substrate determines whether the agent completes the action. Categories dominated by considered purchases — B2B software, financial services, education — weight semantic legibility more heavily because that substrate determines whether the brand survives the synthesis pass. Categories dominated by reference and information needs weight entity legibility more heavily because that substrate determines whether the brand is recognized as the authoritative source.

The weighting is empirically defensible rather than arbitrary. It is derived from category-level patterns of how AI assistants are actually used, and it is revisited as those patterns shift. The SEI is not a fixed formula; it is a structured measurement design that adapts to the category in which it is applied.

4.3 Validation strategy

The SEI requires two validation efforts to be credible as a brand-equity instrument.

Convergent validity with survey-based brand-equity instruments (Aaker, 1996; Yoo & Donthu, 2001) tests whether brands that score high on consumer-perceived equity also score high on SEI, and whether the relationship has the expected category-conditional structure. The expectation is partial convergence — the substrates are different — but with significant correlation in mature categories.

Predictive validity against share-of-citation longitudinal data tests whether changes in SEI predict subsequent changes in observed brand outcomes in AI-mediated channels. The expectation is that SEI changes lead share-of-citation changes by an empirically estimable lag, with the lag reflecting the velocity of training-corpus updates and retrieval-index refreshes in the underlying systems.

Both validation efforts are tractable with audit methods that already exist; the agenda from Figueira (2026b, RQ1 and RQ4) maps the work in more detail.

5. The Brand Legibility Framework

The construct and the measure are necessary but not sufficient. A brand that knows its SEI and its legibility profile still needs an operating model that converts diagnosis into intervention. We propose the *Brand Legibility Framework* as that operating model. The framework presented here is the academic articulation of a practice the author leads at Wyse Brand Intelligence; we name it the Brand Legibility Framework for the purposes of this paper.

The framework re-anchors the three optimization disciplines introduced in Figueira (2026a) — Answer Engine Optimization (AEO), Generative Engine Optimization (GEO), and Agentic Optimization (AgO) — as three coordinated layers, each tied to one of the legibility dimensions and one of the substrates.

The AEO layer targets entity legibility. Its interventions are atomic propositions, schema markup, named-entity binding of facts to the brand, and the disciplined publication of category-defining definitions in extractable form. Its diagnostic indicator is the entity-graph presence score within the SEI. Its failure mode, when not addressed, is the smallest case of brand erasure: factual disintermediation.

The GEO layer targets semantic legibility. Its interventions are information-gain content (proprietary data, original methodologies, named instruments), definitional anchoring of category-load-bearing terms, and the construction of paraphrase-resistant phrasing that survives the synthesis pass. Its diagnostic indicator is the latent-space proximity score within the SEI. Its failure mode is the dominant case of brand erasure: synthetic erasure.

The AgO layer targets operational legibility. Its interventions are infrastructural — published manifests, machine-readable schemas, idempotent endpoints, clear error semantics, and explicit affordance documentation written for an agent reader rather than a human developer. Its diagnostic indicator is the response-audit citation score for task-completion prompts. Its failure mode is the most consequential case of brand erasure: action substitution.

Layer	Legibility Dimension	Substrate	SEI Component	Primary Intervention Lever
AEO (Answer Engine Optimization)	Entity legibility	Knowledge graphs, structured data, named-entity recognition	Entity-graph presence score	Atomic propositions, schema markup, named-entity binding of facts to the brand
GEO (Generative Engine Optimization)	Semantic legibility	Latent space of foundation models	Latent-space proximity score	Information-gain content, definitional anchoring, paraphrase-resistant phrasing
AgO (Agentic Optimization)	Operational legibility	Agent manifests, APIs, structured affordances	Response-audit citation score (task-completion subset)	Published manifests, idempotent endpoints, agent-readable affordance documentation

Table 1. The Brand Legibility Framework. Each layer targets one legibility dimension, operates on one substrate, is diagnosed through one SEI component, and is addressed through one set of intervention levers.

The framework’s central claim is that these three layers are not substitutes. A brand can be strong at AEO and weak at AgO. A brand strong at AgO can still be erased at GEO from the answers that route consumers to its action surface in the first place. The strategic decision the framework forces is which layer is currently rate-limiting for the brand’s category — and the SEI’s component scores are designed to make that decision empirical rather than intuitive.

6. The Diagnostic: Brand Erasure Audit

The framework is applied through the *Brand Erasure Audit*, the measurement protocol that produces the SEI and locates the brand’s position on each of the three layers.

The audit has three parts.

The *embedding probe* measures the brand’s position in the latent space of foundation models. For a target brand, attribute vectors are defined for the category-relevant dimensions identified in Section 4.1. Cosine similarities are computed between the brand vector and each attribute vector across a panel of foundation models (typically four to six). The probe is repeated for direct competitors in the same category. The output is a matrix that locates the brand in semantic space, benchmarks it against competitors, and — when repeated longitudinally — tracks drift across model updates.

The *response audit* samples AI assistant outputs at scale. A representative set of category-relevant prompts is constructed, covering the brand’s strategically important search and task intents. Each prompt is submitted to a population of AI assistants (typically three to five) with three to five sampled responses per prompt-model cell to control for stochasticity. Responses are coded for brand-mention frequency, share of citations, sentiment of the citing context, hallucination rate (where the assistant misrepresents the brand’s facts), and — for AgO prompts — agent task-completion rate.

The *infrastructure audit* assesses operational legibility directly. Manifests are checked for completeness and currency, structured data is validated against schema.org, API endpoints are tested for idempotence and clear error semantics, and authentication flows are evaluated for agent compatibility.

The three parts produce, in combination, the three SEI components, and they produce them in a way that is repeatable, comparable across brands within a category, and trackable across time. The first audit establishes a baseline. Subsequent audits establish the delta — the change in SEI that interventions are accountable to.

SEI Component	Brand under audit	Category leader	Median competitor
Entity-graph presence	0.72	0.89	0.61
Latent-space proximity	0.54	0.78	0.59
Response-audit citation	0.38	0.71	0.42
Composite SEI	0.55	0.79	0.54

Figure 1. Illustrative SEI scorecard output. Scores are bounded [0, 1] and benchmarked against the category leader and the median competitor. The composite SEI is a weighted combination of the three components, with weights determined by the rate-limiting substrate for the category (Section 4.2). The illustrative profile shown — strong entity-graph presence, weak latent-space proximity, weakest response-audit citation — describes a brand whose indexable infrastructure is intact but whose semantic position fails to convert into surfaced citations. The diagnostic points the intervention at the GEO layer.

A complete audit of a single brand in a defined category takes between two and four weeks of work with the right instrumentation. The protocol can be automated for ongoing monitoring; the strategic interpretation of the output cannot.

7. The Asymmetry Property and Strategic Implications

Brand Legibility is asymmetrically distributed across brands. The asymmetry was identified in Figueira (2026b) as a property of brand erasure; it appears here as a property of legibility, and it has direct strategic implications.

Brands with high prior investment in indexed presence — canonical Wikipedia entries, Wikidata identifiers, structured-data coverage, stable named-entity recognition, and a long history of co-mentions with category-load-bearing terms — are legible to AI assistants at substantially higher baseline rates than brands whose distinctiveness was carried in copywriting, visual identity, and broadcast communications. The legacy of brand-equity investment appears in the AI substrate, but only the indexed component of that investment.

This does not invalidate the brand-management orthodoxy of the last three decades. For the substrate that mattered — human associative memory — the advice to build distinctiveness through creative work was, and remains, correct. Memorable advertising, distinctive visual identity, and ownable copywriting still produce the salience, recall, and emotional resonance that human consumers carry with them into the moment of choice. What changes is that the human substrate is no longer the only substrate. The AI-mediated layer operates on a second substrate, and distinctiveness in that substrate is built differently. A brand whose taglines, slogans, and creative campaigns were never embedded in structured-data manifests, defined terms, or knowledge-graph relationships is fully distinctive for human consumers and partially invisible to the machine intermediary that increasingly precedes them. The work the brand did is intact. The work the brand did not do is the new gap.

The strategic implications follow.

For *established brands* with strong indexed presence, the priority is consolidation. The SEI baseline is favorable; the work is to maintain it across the non-stationarity of foundation models, to extend definitional anchoring into emerging category terms, and to harden operational legibility before the agentic transition matures. The risk is complacency: legibility advantages can decay through model updates, paraphrase compounding, and the erosion of original attribution in downstream content.

For *challenger brands* without strong indexed presence, the priority is construction. The SEI baseline is unfavorable; the work is to rebuild the indexed substrate the brand never invested in — entity-graph entries, structured-data coverage, named methodologies, paraphrase-resistant content. This work has a long lead time. It is not a campaign; it is infrastructure. The window in which it can be done at reasonable cost is now, while training corpora are still being assembled and entity graphs are still being filled in. The window will close, in the sense that the cost of forcing inclusion into established AI substrates rises as the substrates harden.

For brands across the spectrum, the framework introduces a question the click economy did not pose: which category-load-bearing terms does the brand own? Definitional anchoring (Figueira, 2026a, Proposition 3) is the lever that converts content into asymmetric citation, and it is available to brands of any size that are willing to operationalize their methodologies under named, distinctive labels. The Semantic Equity Index

makes this lever measurable: brands that own category-load-bearing terms show systematically higher latent-space proximity scores to those terms across foundation models.

8. Brand Legibility as a Discipline

The argument of this paper has been that the brand-equity construct needs an additional pole — the latent-substrate property whose absence produces erasure — and that the pole is operational once the construct, the measure, the framework, and the diagnostic are made concrete. The discipline-level move that follows is that brand management acquires a substrate it did not previously have to design for.

The dual-audience design implication, sketched in Figueira (2026a), follows directly. Every public surface of the brand has two readers: the human who eventually consumes the result, and the machine that mediates access. The machine’s reading happens first and shapes whether the human reading occurs at all. Brand Legibility is the property the brand maintains for the machine reader; it does not displace the work the brand does for the human reader, but it precedes it in the chain of mediation.

The organizational implication is that the functions that have historically operated against different substrates — brand strategy against consumer cognition, SEO against ranking algorithms, content marketing against attention, product engineering against user experience — now share a substrate. The Brand Legibility Framework treats the substrate as the shared design constraint. The traditional handoff between functions produces machine-illegible artifacts at every step; the framework requires the substrate to be designed in at every step.

The professional implication is the emergence of a role whose work is neither marketing nor engineering nor traditional analytics: the AI visibility auditor, whose practice is to maintain the brand’s legibility profile across the non-stationary systems that determine it. The role is closer in spirit to internal audit than to marketing; its work is to keep the brand’s trust calibration accurate as the underlying systems change (Glikson & Woolley, 2020).

9. Research Agenda for Brand Legibility

The construct and the framework open empirical questions that the marketing-science toolkit is positioned to address. The agenda from Figueira (2026b) mapped the broader research program for erasure; this section identifies the questions specific to legibility as a positive construct.

The SEI refinement question. The index as proposed weights its three components by category. The optimal weighting scheme is an empirical question that requires cross-category validation against observed brand outcomes in AI-mediated channels. The agenda includes establishing the convergent and predictive validity of the index, and developing category-specific benchmarks that brands can use to interpret their scores.

The legibility dynamics question. Legibility is non-stationary. The half-life of legibility under continuous foundation-model updating, controlling for the brand’s intervention rate, is an open empirical question with substantial managerial consequences. Related: under what conditions does illegibility become unrecoverable,

in the sense that the brand’s representation in foundation-model substrates cannot be repaired through any practical intervention?

The cross-language legibility question. The SEI as constructed assumes a single-language substrate. Multinational brands operate in multiple language substrates simultaneously, with foundation models exhibiting different legibility profiles across languages. The question is whether legibility transfers across languages within a model, across models within a language, or neither, and what intervention designs work in each case.

The intervention-to-outcome question. The framework prescribes interventions at each of the three layers. The empirical question is the magnitude and lag of the SEI response to specific interventions, and the relationship between SEI changes and downstream brand outcomes (revenue, share, customer acquisition cost) in AI-mediated channels. The question is tractable with longitudinal panel designs adapted from the digital-shelf and content-marketing literature (Reisenbichler, Reutterer, Schweidel, & Dan, 2022).

The governance question. The construct of Brand Legibility intersects the broader governance debate about attribution in AI systems (Hacker, Engel, & Mauer, 2023; Veale & Zuiderveen Borgesius, 2021; European Parliament & Council, 2024). The question is whether legibility, once operationalized, becomes a basis for attribution rights — and whether the welfare consequences of regulating attribution differ when the underlying construct is measurable rather than rhetorical.

10. Limitations

The contribution of this paper is conceptual, and its limitations follow from that decision.

First, the Semantic Equity Index is proposed without empirical validation in this paper. The components are individually well-grounded in adjacent literatures — entity-graph metrics in the structured-data tradition, embedding probes in the bias-audit tradition (Bolukbasi et al., 2016; Caliskan et al., 2017), response audits in the algorithmic-accountability tradition (Diakopoulos, 2016) — but the integrated index, the weighting scheme across components, and the convergent and predictive validity claims await empirical work. The agenda from Figueira (2026b) and the research agenda in Section 9 of this paper map that work; the present paper does not perform it.

Second, foundation models are non-stationary. The latent-space proximity scores that enter the SEI today will shift as models are updated, retraining cycles complete, and retrieval indexes refresh. The framework accommodates this through longitudinal tracking and change-point detection, but the absolute interpretation of any single audit is bounded by the state of the underlying systems at the time of measurement. The construct of legibility is stable; the measurement is not. This is a property of the substrate, not a defect of the instrument.

Third, the framework is developed in a single-language context. Multinational brands operate in multiple language substrates, and foundation models exhibit different legibility profiles across languages — sometimes substantially. The cross-language generalizability of the SEI is an open question, and the weighting scheme proposed in Section 4.2 may require language-specific calibration that the current paper does not address.

Fourth, the framework assumes that AI-mediated channels are a meaningful proportion of the brand’s consumer interaction. For categories in which AI-mediated discovery and transaction remain a small share of consumer behavior, the strategic urgency of the construct is correspondingly lower. The Pew (2025) data document one category — general-information search — in which the shift is material. Other categories may follow on different timelines. The framework does not specify, and the present paper does not predict, the pace of that shift across categories.

Finally, the construct of Brand Legibility intersects domains — information retrieval, computational linguistics, AI governance — in which the conventions are themselves in flux. The framework engages those literatures (Lewis et al., 2020; Liu, Zhang, & Liang, 2023; Hacker, Engel, & Mauer, 2023; European Parliament & Council, 2024) but does not, and cannot, resolve their unresolved questions. The construct is offered as a contribution to a conversation that is still being held, not as a closure of it.

11. Conclusion: From Visibility to Legibility

Marketing has spent thirty years optimizing for visibility. The substrates that determined visibility — search engines, social platforms, advertising inventories — were external systems whose behavior brands could observe and respond to. The substrate that determines whether a brand survives the AI-mediated layer is different in kind. It is internal to the systems that now mediate consumer experience, and it operates on properties — recognizability, parseability, attributability — that the discipline has not historically measured.

Brand Legibility names that property. The Semantic Equity Index makes it measurable. The Brand Legibility Framework organizes the work of building it. The Brand Erasure Audit operationalizes the framework. Together, they offer a discipline for the period in which the substrate is still malleable and the conventions of attribution have not yet hardened.

The brand that survives the AI-mediated layer is not the brand that is most visible. It is the brand the machine can read, parse, and attribute. That is the discipline this paper names — and the construct on which the next generation of brand-equity work will need to be built.

References

- Aaker, D. A. (1991). *Managing brand equity: Capitalizing on the value of a brand name*. Free Press.
- Aaker, D. A. (1996). Measuring brand equity across products and markets. *California Management Review*, 38(3), 102–120.
- Aggarwal, P., Murahari, V., Rajpurohit, T., Kalyan, A., Narasimhan, K., & Deshpande, A. (2024). GEO: Generative engine optimization. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD '24)*. ACM.
- Bolukbasi, T., Chang, K.-W., Zou, J. Y., Saligrama, V., & Kalai, A. T. (2016). Man is to computer programmer as woman is to homemaker? Debiasing word embeddings. *Advances in Neural Information Processing*

Systems, 29.

Caliskan, A., Bryson, J. J., & Narayanan, A. (2017). Semantics derived automatically from language corpora contain human-like biases. *Science*, 356(6334), 183–186.

Diakopoulos, N. (2016). Accountability in algorithmic decision making. *Communications of the ACM*, 59(2), 56–62.

European Parliament and Council. (2024). Regulation (EU) 2024/1689 of the European Parliament and of the Council laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). *Official Journal of the European Union*, L 2024/1689, 12 July 2024.

Figueira, M. G. (2026a). *From information retrieval to agentic action: A framework for brand visibility in AI-mediated markets*. Working paper, Fundação Getulio Vargas (FGV) and Wyse Brand Intelligence. SSRN.

Figueira, M. G. (2026b). *Brand erasure: A research agenda for brand theory in agentic AI markets*. Working paper, Fundação Getulio Vargas (FGV) and Wyse Brand Intelligence. SSRN.

Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., Dai, Y., Sun, J., & Wang, H. (2024). Retrieval-augmented generation for large language models: A survey [arXiv:2312.10997]. arXiv.

Glikson, E., & Woolley, A. W. (2020). Human trust in artificial intelligence: Review of empirical research. *Academy of Management Annals*, 14(2), 627–660.

Hacker, P., Engel, A., & Mauer, M. (2023). Regulating ChatGPT and other large generative AI models. In *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency (FAccT '23)*.

Keller, K. L. (1993). Conceptualizing, measuring, and managing customer-based brand equity. *Journal of Marketing*, 57(1), 1–22.

Khattab, O., & Zaharia, M. (2020). ColBERT: Efficient and effective passage search via contextualized late interaction over BERT. In *Proceedings of SIGIR 2020*.

Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474.

Liu, N. F., Zhang, T., & Liang, P. (2023). Evaluating verifiability in generative search engines. In *Findings of the Association for Computational Linguistics: EMNLP 2023* (pp. 7001–7025).

Pew Research Center. (2025, July 22). *Google users are less likely to click on links when an AI summary appears in the results*.

Reisenbichler, M., Reutterer, T., Schweidel, D. A., & Dan, D. (2022). Frontiers: Supporting content marketing with natural language generation. *Marketing Science*, 41(3), 441–452.

Romaniuk, J., & Sharp, B. (2004). Conceptualizing and measuring brand salience. *Marketing Theory*, 4(4), 327–342.

Santhanam, K., Khattab, O., Saad-Falcon, J., Potts, C., & Zaharia, M. (2022). ColBERTv2: Effective and efficient retrieval via lightweight late interaction. In *Proceedings of NAACL 2022*.

Veale, M., & Zuiderveen Borgesius, F. (2021). Demystifying the draft EU Artificial Intelligence Act. *Computer Law Review International*, 22(4), 97–112.

Yoo, B., & Donthu, N. (2001). Developing and validating a multidimensional consumer-based brand equity scale. *Journal of Business Research*, 52(1), 1–14.