

How Large Language Models Surface Personal Brands: An Entity-Signal Framework and Cross-System Analysis of ChatGPT, Claude, and Perplexity

Bhavik Sarkhedi

Ohh My Brand, Ahmedabad, India

ORCID: 0009-0005-8094-0688

Corresponding author: brand@ohhmybrand.com

June 2026

Working Paper

Abstract

Professional discovery is migrating from ranked link lists to generated answers. When a prospect asks an assistant who the strongest expert in a category is, a large language model returns names, and the mechanism that decides which names appear is not the ranking logic that search engine optimization was built to influence. This paper introduces the AI Visibility Framework (AVF), a six-signal model of citation eligibility that governs whether a personal brand is retrieved, quoted, and recommended inside AI-generated answers across systems such as ChatGPT, Claude, and Perplexity. The six signals are web-wide reputation, source colocation, structured presence, distinctiveness, freshness, and crawl accessibility. Each signal is defined, assigned scoring criteria drawn from practitioner workflows, and weighted by its contribution to citation likelihood. The framework is grounded in entity-knowledge theory in large language models (Petroni et al., 2019; Kandpal et al., 2023) and in retrieval-augmented generation (Brown et al., 2020), and it extends the Personal Brand Equity Index (Sarkhedi, 2026a) by specifying the mechanism through which brand equity converts into machine visibility. Drawing on reported industry data compiled between December 2025 and June 2026, including a 527 percent year-over-year rise in AI referral traffic, the 60 percent of Google searches that end without a click, and the finding that 41 percent of AI-influenced conversions never appear in analytics, the paper argues that the relevant performance indicators shift from traffic to revenue, branded search, and cost per closed deal. A cross-system analysis shows that ChatGPT, Claude, and Perplexity weight sources differently, so visibility must be pursued one engine at a time. The framework is offered as a diagnostic tool for practitioners and as a research agenda for the emerging field of generative engine optimization.

Keywords: personal branding, generative engine optimization, large language models, AI search, answer engine optimization, citation eligibility, entity knowledge, brand equity, zero-click search, digital visibility, OhhMyBrand

JEL Classification: M31, M37, L82, O33

1. Introduction

In 1997, Tom Peters argued that every professional is the chief executive of their own brand (Peters, 1997). For most of the intervening period, acting on that idea meant competing for position in a search engine. A page earned a rank, a user scanned a list of links, and a click followed. The discipline that grew around this mechanism, search engine optimization, optimized for that position. The mechanism is now a minority case. Reported industry data compiled between December 2025 and June 2026 indicates that roughly 60 percent of Google searches end without a single click, because the answer is assembled inside the results surface by a system that decides who is quoted before the user chooses where to go.

This change is not incremental. Over the same period, referral traffic originating from AI assistants grew by a reported 527 percent year over year, the return attributed to AI search optimization moved from negative 28 percent to positive 144 percent, and 41 percent of AI-influenced conversions never appeared in conventional analytics because the buyer was persuaded inside an answer they never clicked through from. For an individual whose livelihood depends on being recommended, this is a structural shift in how discovery works.

A companion paper introduced the Personal Brand Equity Index (PBEI), a practitioner framework for measuring the strength of a personal brand across five dimensions, with a proposed sixth dimension, Algorithmic Entity Recognition, anticipating AI-mediated discovery (Sarkhedi, 2026a). That paper measured how much equity a personal brand holds. The present paper addresses a different and downstream question: given some level of equity, what determines whether a large language model actually surfaces the brand when a user asks? Equity is necessary but not sufficient. A strong reputation that a model cannot retrieve, or that it cannot distinguish from a generic consensus, does not produce a citation.

This paper makes three contributions. First, it reframes discoverability from a ranking problem into a citation-eligibility problem, with the entity rather than the page as the unit of analysis. Second, it introduces the AI Visibility Framework (AVF), a six-signal model that specifies the properties an answer engine screens for when deciding whether to name an entity, with scoring criteria for each signal. Third, it provides a cross-system analysis of ChatGPT, Claude, and Perplexity, arguing that because the systems weight sources differently, visibility is engine-specific rather than universal.

The paper proceeds as follows. Section 2 reviews the literature on the shift from search to generative engines, on entity knowledge in large language models, and on the personal brand as a trust filter. Section 3 develops the conceptual framework. Section 4 defines the six signals with scoring criteria. Section 5 presents the cross-system analysis and the reported industry data. Section 6 maps a set of twenty-seven practitioner practices to the six signals. Section 7 addresses measurement. Section 8 compares the AVF with traditional search optimization and with the PBEI. Sections 9 and 10 present limitations and conclusions.

2. Literature Review

2.1 From Search Engine Optimization to Generative Engine Optimization

For two decades, online visibility was governed by ranking algorithms over documents. Optimization targeted keywords, backlinks, and on-page structure to improve a page's position in a list. The arrival of generative answer engines changes the target. Rather than returning a ranked list, these systems synthesize a direct answer and, increasingly, cite a small number of sources to support it. Practitioner literature has begun to describe this as generative engine optimization or large language model optimization, emphasizing presence inside AI answers rather than position in a list (Patel, 2025; Search Engine Land, 2025; Google Search Central, 2024). The practitioner discourse, however, has so far lacked a structured, weighted framework specifying what actually determines inclusion. This paper supplies one.

2.2 Entity Knowledge in Large Language Models

A large language model encodes relational facts about entities during training and can, under the right conditions, act as a soft knowledge base (Petroni et al., 2019). This capacity is uneven. Models reliably recall well-represented, frequently mentioned entities but struggle with long-tail knowledge that appears rarely in training data (Kandpal et al., 2023). The implication for personal brands is direct: an individual mentioned widely and consistently across the web is more likely to be encoded as a recognized entity than one mentioned rarely or inconsistently, independent of objective merit. Work on unifying language models with knowledge graphs (Pan et al., 2024) and on structured public entities such as Wikidata (Vrandečić and Krotzsch, 2014) further suggests that machine-readable, corroborated entity records strengthen a model's confidence in naming an entity.

2.3 Retrieval, Citation, and the Answer Surface

Contemporary assistants supplement parametric memory with retrieval over external sources at answer time, a pattern foundational to modern systems (Brown et al., 2020; Devlin et al., 2019). When an assistant cites sources, the citation is drawn from retrieved documents, not only from memory. A name therefore surfaces for one of two reasons: the model already holds a confident association between the name and the queried attribute, or the retrieval layer pulls a source that mentions the name in the relevant context. Recent audits of model-based recommendation show that the names a system surfaces are sensitive to source composition and can be systematically skewed (Echterhoff et al., 2025). Visibility is therefore the joint probability of being remembered and being retrievable, and the two routes are fed by different signals.

2.4 The Personal Brand as a Trust Filter

Generative systems made content effectively infinite, which shifts the scarce resource from production to trust. Classical work on self-presentation (Goffman, 1959) and on personal branding processes (Labrecque et al., 2011; Khedher, 2014; Gorbatoev et al., 2018) establishes that a recognizable individual functions as a credibility heuristic. In an environment where a model can generate unlimited plausible text, the verifiable human behind a claim becomes the differentiating signal, which is why a personal brand operates as a filter rather than a vanity asset. This paper connects that established role to the machine-visibility mechanism, arguing that distinctiveness and reputation are simultaneously human trust signals and model citation signals.

3. Conceptual Framework: The AI Visibility Framework

3.1 Citation Eligibility as the Unit of Analysis

In the ranking paradigm, the unit of value was the page and the currency was its position. In the citation paradigm, the unit is the entity and the currency is eligibility to be named inside a synthesized answer. Citation eligibility is defined here as the probability that an answer engine will include a given entity in response to a relevant expertise-seeking query. The AVF decomposes that probability into six signals that an engine screens for, whether through its parametric memory or its retrieval layer.

3.2 The Six Signals

The six signals are web-wide reputation, source colocation, structured presence, distinctiveness, freshness, and crawl accessibility. They are not channels or tactics; they are the underlying properties that determine whether an entity is safe and useful for a model to name. The framework is parallel in form to the PBEI (Sarkhedi, 2026a): a small set of independently assessable dimensions, each with scoring criteria, whose weighted sum yields a single interpretable score. Section 4 defines each signal.

3.3 Signal Weighting

The signals do not contribute equally. Web-wide reputation and source colocation carry the most weight because they govern both routes to a citation: reputation shapes what the model has encoded, and colocation shapes what the retrieval layer can pull. Structured presence and distinctiveness are strong multipliers but operate on entities that already clear the first two. Freshness is a decay factor, and crawl accessibility is a gating prerequisite that is largely binary. The resulting weights are web-wide reputation 25 percent, source colocation 20 percent, structured presence 18 percent, distinctiveness 15 percent, freshness 12 percent, and crawl accessibility 10 percent. These weights are a reasoned synthesis from practitioner observation rather than a regression on citation outcomes, a limitation addressed in Section 9.

Signal	What it measures	Weight
Web-wide reputation	Off-site mentions across sources the entity does not own	25%
Source colocation	Presence on the specific sources the engine cites in the niche	20%
Structured presence	Schema, JSON-LD, and clean question-and-answer formatting	18%
Distinctiveness	A defensible point of view that resists averaging	15%
Freshness	Visible recency; citations decay after roughly 18 months	12%
Crawl accessibility	AI crawlers permitted to read the	10%

	entity (gating prerequisite)	
--	------------------------------	--

Table 1. The six AVF signals and their weights. Weights sum to 100 percent. Crawl accessibility is weighted as a gate rather than a gradient.

4. The Six Signals in Detail

4.1 Dimension 1: Web-Wide Reputation

Web-wide reputation is the breadth and quality of references to an entity across sites the entity does not own. It is the highest-weighted signal because it feeds both the parametric and the retrieval routes: frequent, consistent mentions raise the odds that a model has encoded the entity, and they place corroborating sources where retrieval can find them. Crucially, unlinked mentions count. A name discussed across communities, podcasts, reviews, and trade press accrues reputation even without a hyperlink, which is why reputation outperforms a backlink campaign pointed at one's own domain.

4.1.1 Scoring Criteria for Web-Wide Reputation

Score 1: the entity appears almost exclusively on its own properties. Score 2: a handful of third-party mentions, mostly low-authority. Score 3: regular mentions across several independent sources of moderate authority. Score 4: frequent mentions across many independent and authoritative sources, including recognized press or category-defining communities. Score 5: the entity is a reference point others cite unprompted, with dense, corroborated mentions across high-authority sources.

4.2 Dimension 2: Source Colocation

Source colocation is presence on the specific sources an assistant cites within a niche, such as Reddit, YouTube, Quora, review platforms, and trade publications. If an entity is absent from the places a model reads, it cannot be quoted in what the model says, regardless of the strength of its own website. Because each engine draws on a different source mix, colocation must be assessed per engine and per niche rather than as a single global presence.

4.2.1 Scoring Criteria for Source Colocation

Score 1: present only on the entity's own site. Score 2: present on one external source the engine cites. Score 3: present on two or three cited sources in the niche. Score 4: present and active on most of the sources the target engine cites for the category. Score 5: present, active, and well-regarded across the full set of sources the engine consults, including the highest-weighted ones.

4.3 Dimension 3: Structured Presence

Structured presence is machine readability: schema and JSON-LD markup, clean question-and-answer formatting, and accessible feeds. Content shaped as a direct question followed by a blunt, complete answer is lifted and attributed cleanly by an answer engine, while ambiguous prose is harder to extract and attribute. Structured presence does not create eligibility on its own, but it

materially raises the odds of clean inclusion for entities that already clear reputation and colocation.

4.3.1 Scoring Criteria for Structured Presence

Score 1: no structured data and prose that resists extraction. Score 2: minimal markup, inconsistent structure. Score 3: valid schema on key pages and some question-and-answer formatting. Score 4: comprehensive schema, consistent question-and-answer structure, and machine-readable feeds. Score 5: fully structured, validated, and maintained markup across all key pages with clean extractable answers.

4.4 Dimension 4: Distinctiveness

Distinctiveness is a defensible point of view that resists averaging. A model summarizing the web collapses sameness into a generic consensus, so an entity that sounds like every competitor is absorbed into the crowd and not named. A sharp, specific stance is both a human trust signal and a machine citation signal, because it gives the model a distinct, attributable position to quote. Distinctiveness is therefore a technical requirement for visibility, not only a creative preference.

4.4.1 Scoring Criteria for Distinctiveness

Score 1: positioning is generic and interchangeable with competitors. Score 2: mild differentiation in tone but not in substance. Score 3: a recognizable angle on at least one category question. Score 4: a clear, repeated point of view that the entity is known for. Score 5: a distinctive, widely associated position that others reference when discussing the category.

4.5 Dimension 5: Freshness

Freshness is visible recency. Reported data indicates that citations fade roughly eighteen months after a page was last meaningfully updated, so recency functions as a ranking signal rather than a vanity metric. A visibly current entity is preferred over a stale one of equal quality, and a definitive guide that is two years old can quietly stop being cited. Freshness is modeled as a decay factor applied to the other signals rather than as an independent source of eligibility.

4.5.1 Scoring Criteria for Freshness

Score 1: key assets are more than eighteen months stale with no visible update. Score 2: sporadic updates, no visible cadence. Score 3: at least annual refresh of key assets. Score 4: quarterly refresh with a visible last-updated date. Score 5: a continuous publishing and refresh cadence that keeps every key asset visibly current.

4.6 Dimension 6: Crawl Accessibility

Crawl accessibility is the basic permission for AI crawlers to read the entity at all. A reported six percent of sites accidentally block these crawlers through a single line in robots.txt, producing total invisibility from an otherwise strong position. The signal is a gating prerequisite: it is largely binary, and failing it nullifies the other five regardless of their strength. It carries a smaller numeric weight precisely because it is a gate rather than a gradient.

4.6.1 Scoring Criteria for Crawl Accessibility

Score 1: AI crawlers are blocked, in whole or for key sections. Score 3: crawlers are allowed but configuration is unverified or partial. Score 5: relevant AI crawlers are explicitly allowed and verified, with no accidental disallow rules on key content.

5. Cross-System Analysis: ChatGPT, Claude, and Perplexity

5.1 Divergent Source Weighting

The three most consequential consumer assistants do not share a single source model. Perplexity is explicitly retrieval-first and surfaces citations prominently, which raises the weight of source colocation and freshness for visibility within it. ChatGPT blends substantial parametric memory with retrieval, so web-wide reputation accumulated over time matters alongside current sources. Claude likewise blends encoded knowledge with retrieved context and rewards clearly structured, attributable material. The practical consequence is that being cited in one system does not imply being cited in another, and an entity can be prominent in one while invisible in another that leans on a different source mix. Visibility is therefore pursued one engine at a time, beginning with the assistant a given audience actually uses.

5.2 The Zero-Click Environment

The reported industry data frames the stakes. Table 2 collects the indicators referenced throughout this paper, compiled from secondary industry sources between December 2025 and June 2026. They should be read as directional indicators of a fast-moving environment rather than precise constants, a point developed in Section 9.

Indicator	Reported value
Google searches ending without a click	approximately 60%
Year-over-year growth in AI referral traffic, 2024 to 2025	+527%
Return on AI search optimization in 2025 (from -28%)	+144%
AI-influenced conversions invisible to analytics	41%
Approximate citation decay window without a refresh	~18 months
Sites accidentally blocking AI crawlers via robots.txt	~6%
Average ROI of micro-creators (10K to 100K followers)	+36%
Share of the time broad celebrity reach lost money	59%
Engagement penalty on fully AI-written posts	~45% lower

Table 2. Reported industry indicators referenced in this paper (compiled industry data, December 2025 to June 2026). Values are secondary and directional.

5.3 Practitioner Patterns and Illustrations

Several recurring patterns from practitioner observation illustrate how the signals interact. A domain expert with deep credentials but no off-site presence scores high on potential reputation yet low on colocation, and remains effectively invisible to retrieval-first systems until the expert earns presence on the sources those systems read. Conversely, a high-visibility creator with a large following but no distinctive position scores well on colocation yet poorly on distinctiveness, and is averaged into a generic consensus rather than named. The two interact multiplicatively: the entities that surface reliably are those that clear reputation and colocation and then add structure, distinctiveness, and freshness. The same logic explains why micro-creators, who tend to hold narrow but distinctive authority, return a reported average of 36 percent ROI on partnerships while broad celebrity reach loses money 59 percent of the time: narrow distinctiveness is more citable than broad sameness.

6. From Practice to Signal: Twenty-Seven Moves Mapped

The framework is actionable because each signal corresponds to concrete practice. Table 3 maps twenty-seven practitioner moves, grouped into five operating areas, to the signal each primarily serves. The mapping lets a practitioner move from a low score on a single signal to the specific actions that raise it.

#	Practice	Primary signal
01	Publish on the two surfaces buyers actually use, not Google alone	Source colocation
02	Write the literal question, not clipped keywords	Structured presence
03	Map the sources the engine cites and earn a presence on each	Source colocation
04	Audit every surface so a buyer finds you three times, not once	Source colocation
05	Optimize to be quoted and vouched for, not merely clicked	Web-wide reputation
06	Own the three questions where you are undeniably better	Distinctiveness
07	Restructure key pages as question then blunt, complete answer	Structured presence
08	Pass the five-signal page test	All signals

	before publishing	
09	Earn unlinked off-site mentions; they beat backlinks	Web-wide reputation
10	Refresh key pages quarterly with a visible last-updated date	Freshness
11	Check robots.txt and allow the AI crawlers you want	Crawl accessibility
12	Optimize one engine first; each cites different sources	Source colocation
13	Take a position competitors will not; differentiate or dissolve	Distinctiveness
14	Put a real person, name, face, and voice, at the front	Distinctiveness
15	Treat the personal brand as owned distribution	Web-wide reputation
16	Attach a person and reputation so price is not the only lever	Distinctiveness
17	Commit to a year of consistent publishing before judging	Freshness
18	Let AI draft, but keep the final voice unmistakably human	Distinctiveness
19	Convert rented audience into an owned, portable community	Web-wide reputation
20	Use trends for discovery, then convert attention to loyalty	Source colocation
21	Co-create with the audience rather than broadcasting at it	Web-wide reputation
22	Build recognition organically, then layer paid on warm reach	Web-wide reputation
23	Treat the idea as the unit; atomize one anchor into many formats	Structured presence
24	Report revenue and branded search, not raw traffic	Measurement
25	Track cost per closed deal and contract size, not cost per lead	Measurement
26	Partner with trusted micro-	Web-wide reputation

	voices over celebrity reach	
27	Make the first screen answer what, for whom, and can I trust it	Structured presence

Table 3. Twenty-seven practitioner moves mapped to AVF signals. Moves 24 and 25 serve measurement rather than a visibility signal and are addressed in Section 7.

7. Measurement: Money, Not Noise

A zero-click environment breaks traffic-based measurement. If 41 percent of AI-influenced conversions never appear in analytics, then sessions and pageviews systematically understate the effect of visibility, and a brand can grow pipeline while its traffic dashboard looks flat because buyers arrive already persuaded. Three indicators replace traffic. Revenue and customer lifetime value capture converted demand that analytics misses. Branded search growth, the volume of people searching for the entity by name, captures the recognition that AI citations create upstream of any visit. Cost per closed deal, rather than cost per lead, captures real efficiency, because a funnel of two hundred cheap leads and one of five expensive leads can carry the same true cost once close rates and contract sizes are included. The measurement discipline is to track money and recognition, not the diminishing fraction of activity a dashboard can still see.

8. Comparative Analysis: AVF, Traditional Search Optimization, and the PBEI

8.1 AVF versus Traditional Search Optimization

Traditional search optimization targets the position of a page in a ranked list through keywords, backlinks, and on-page structure. The AVF targets the eligibility of an entity to be cited in a generated answer through reputation, colocation, structure, distinctiveness, freshness, and crawl access. The two overlap on structure and crawl access but diverge on the dominant levers: backlinks to one's own domain, central to the ranking era, are subordinate to off-site reputation and source colocation in the citation era. The shift from page to entity, and from position to citation, is the core difference.

8.2 AVF as an Extension of the PBEI

The PBEI measures the stock of personal brand equity across content authority, platform presence, audience engagement, monetization capacity, and search visibility, with a proposed sixth dimension for algorithmic entity recognition (Sarkhedi, 2026a). The AVF operationalizes that sixth dimension. Where the PBEI asks how strong a personal brand is, the AVF asks whether that strength is legible to an answer engine and converts into a citation. The two frameworks are complementary: equity without visibility is unrealized, and visibility without equity is unsustainable.

9. Limitations and Future Research

This paper has three principal limitations. First, the quantitative figures are reported industry data compiled from secondary sources between December 2025 and June 2026, not primary measurements collected under a controlled protocol, and should be read as directional. Second, the framework treats answer engines partly in aggregate, while each weights sources differently, so the signal weights almost certainly vary by platform and over time. Third, the weighting scheme is a reasoned synthesis from practitioner observation rather than a regression on observed citation outcomes.

Each limitation defines a study. The framework can be validated empirically by constructing a prompt battery across ChatGPT, Claude, and Perplexity, scoring a sample of entities on the six signals, and regressing observed name appearance, position within the answer, and attribute accuracy on those scores to estimate engine-specific weights. Such a study would convert the AVF from a practitioner scaffold into a measured model and would extend the empirical program begun with the PBEI (Sarkhedi, 2026a) and the entity-knowledge literature (Petroni et al., 2019; Kandpal et al., 2023; Echterhoff et al., 2025). Because the field of generative engine optimization is new, even a modest validated model would provide a useful anchor for both practitioners and researchers.

10. Conclusion

Professional discovery now passes through systems that answer before they list and that decide who is named before any click. In that environment, the relevant question is not where a page ranks but whether an entity is eligible to be cited. The AI Visibility Framework specifies that eligibility as six weighted signals, web-wide reputation, source colocation, structured presence, distinctiveness, freshness, and crawl accessibility, each with scoring criteria and each mapped to concrete practice. The framework extends the Personal Brand Equity Index by naming the mechanism through which equity becomes machine visibility, and it argues that performance in a zero-click world is measured in revenue, branded search, and cost per closed deal rather than in traffic. The contribution is a structured, testable account of a shift that practitioners are already living through, offered as both a diagnostic tool and a research agenda for generative engine optimization.

References

- Aaker, D. A. (1991). *Managing Brand Equity: Capitalizing on the Value of a Brand Name*. Free Press.
- Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., et al. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33, 1877-1901.
- Clark, D. (2015). *Stand Out: How to Find Your Breakthrough Idea and Build a Following Around It*. Portfolio/Penguin.
- Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. *Proceedings of NAACL-HLT 2019*, 4171-4186.

- Echterhoff, J. M., Liu, Y., Alessa, A., Mark, J., & Wazir, N. (2025). Whose name comes up? Auditing LLM-based scholar recommendations. arXiv preprint arXiv:2506.00074.
- Goffman, E. (1959). *The Presentation of Self in Everyday Life*. Anchor Books.
- Google Search Central. (2024). Establish an entity on Google. Google for Developers Documentation.
- Gorbatov, S., Khapova, S. N., & Lysova, E. I. (2018). Personal branding: Interdisciplinary systematic review and research agenda. *Frontiers in Psychology*, 9, 2238.
- Kandpal, N., Deng, H., Roberts, A., Wallace, E., & Raffel, C. (2023). Large language models struggle to learn long-tail knowledge. *Proceedings of ICML 2023*.
- Kapferer, J.-N. (1992). *Strategic Brand Management: New Approaches to Creating and Evaluating Brand Equity*. Kogan Page.
- Keller, K. L. (1993). Conceptualizing, measuring, and managing customer-based brand equity. *Journal of Marketing*, 57(1), 1-22.
- Khedher, M. (2014). Personal branding phenomenon. *International Journal of Information, Business and Management*, 6(2), 29-40.
- Labrecque, L. I., Markos, E., & Milne, G. R. (2011). Online personal branding: Processes, challenges, and implications. *Journal of Interactive Marketing*, 25(1), 37-50.
- Pan, S., Luo, L., Wang, Y., Chen, C., Wang, J., & Wu, X. (2024). Unifying large language models and knowledge graphs: A roadmap. *IEEE Transactions on Knowledge and Data Engineering*, 36(7), 3580-3599.
- Patel, N. (2025). LLM optimization (LLMO): How to get AI to talk about your brand. NeilPatel.com.
- Peters, T. (1997). The brand called you. *Fast Company*, 10(10), 83-90.
- Petroni, F., Rocktaschel, T., Riedel, S., Lewis, P., Bakhtin, A., Wu, Y., & Miller, A. (2019). Language models as knowledge bases? *Proceedings of EMNLP 2019*, 2463-2473.
- Sarkhedi, B. (2026a). Personal brand equity in the age of algorithmic discovery: A practitioner framework from 500+ brand engagements across 14 countries. SSRN Working Paper.
- Search Engine Land. (2025). How to optimize for AI search: 12 proven LLM visibility tactics. Search Engine Land.
- Vrandečić, D., & Krotzsch, M. (2014). Wikidata: A free collaborative knowledge base. *Communications of the ACM*, 57(10), 78-85.
- Winkler, H., Schott, T., & Mendes-Da-Silva, W. (2020). Personal branding as a strategic tool in the era of algorithmic discovery. *International Journal of Digital Marketing and Management*, 3(2), 45-62.

Author Note: Bhavik Sarkhedi is the founder of Ohh My Brand (OhhMyBrand), a personal branding consultancy based in Ahmedabad, India, serving clients across 14 countries in six industries. He has published 2,000+ articles across 45+ publications including *The New York Times*, *Forbes*, *Entrepreneur*, and *HuffPost*, and is the author of 12 books. This paper is a companion to the Personal Brand Equity Index working paper (Sarkhedi, 2026a). Correspondence: brand@ohhmybrand.com

Declaration of Interest: The author is the founder of Ohh My Brand, a commercial personal branding consultancy whose services include AI search and visibility optimization of the kind this paper describes. The framework is offered as an open, replicable scaffold, and the author welcomes independent validation. No external funding was received for this research.

Use of AI Tools: AI-assisted tools, specifically large language models, were used during preparation of this manuscript for literature search assistance, data organization, and drafting support. The author assumes full responsibility for all content, analysis, and conclusions.