

Manufactured Authority in Generative Engine Optimization: Cross-Market Evidence and Two Operational Metrics (MAI and VLS)

İbrahim Göktaş

Independent Researcher

ORCID: 0009-0004-4397-0696

info@ibrahimgoktas.com

Preprint, 2026. CC-BY 4.0.

Abstract

Generative Engine Optimization (GEO) is the practice of structuring content so that large language model (LLM) powered search systems retrieve, trust, and cite it. A defining empirical property of these systems is a systematic bias toward earned media, that is, third-party and apparently authoritative sources, over brand-owned and social content (Chen, Wang, Chen, & Koudas, 2025). This paper documents a paradox that follows from that property: the vertical that sells "become the source AI trusts" appears to manufacture precisely that trust about itself. Through an exploratory, observational, cross-market audit of the canonical "best GEO agency/expert" query in eight markets and languages, we identify three recurring mechanisms of manufactured authority: (M1) self-referential ranking, (M2) syndicated corroboration, and (M3) purchased-placement intermediation. To move the analysis from description to measurement, we introduce two operational metrics: the Manufactured Authority Index (MAI), the share of examined first-page sources exhibiting a manufacturing mechanism, and the Verification Layer Score (VLS), a five-criterion 0 to 10 index of a market's independent verification infrastructure. Scoring the eight markets, we find that manufactured authority is near-universal (MAI is consistently high) while the verification layer varies sharply (VLS ranges from 1 to 9). This yields a two-axis picture in which markets are separated not by whether authority is manufactured but by whether anything can independently correct it: in short, source diversity is not source independence. We connect M2 to the manufactured-third-party-agreement mechanism demonstrated experimentally in GEO poisoning research (Lasso Security, 2026; Zhang, Triedman, & Shmatikov, 2026), and we revisit and qualify an earlier claim, from the author's Integrated Visibility (AIVI) framework (Göktaş, 2026), that retrieval architectures intrinsically mitigate manipulation through multi-source corroboration. The study is exploratory and hypothesis-generating rather than confirmatory; its metrics are proposed as reusable instruments, and its limitations are discussed at length. To our knowledge, this is the first study to operationalize manufactured authority through reusable, cross-market metrics.

Keywords: Generative Engine Optimization, Manufactured Authority, Verification Layer, Citation Authority, AI Answer Integrity, GEO Poisoning, Earned Media Bias, AI Search.

1. Introduction

The transition from ranked-list search to synthesized, citation-backed answers has moved the object of digital competition from *ranking* to *being selected as a trusted source* by generative systems (Aggarwal, Murahari, Rajpurohit, Kalyan, Narasimhan, & Deshpande, 2024; Chen et al., 2025). This shift gives rise to a discipline,

Generative Engine Optimization (GEO), whose central promise is to make a brand a source that AI systems cite.

A structural fact about these systems creates a specific incentive. Recent large-scale comparative work shows that AI search exhibits a systematic and overwhelming bias toward earned media, third-party and authoritative sources, in contrast to the more balanced source mix of traditional web search (Chen et al., 2025). In other words, generative systems privilege apparent third-party corroboration. Any actor who can *manufacture* the appearance of third-party corroboration can therefore influence what these systems present.

This paper begins from an observation about the GEO vertical itself. When one queries a generative system, or the search engine that feeds it, for the "best GEO agency" or "best GEO expert," the results are dominated not by independent evaluation but by material that the evaluated parties themselves produced or commissioned. The industry that sells "become the source AI trusts" appears to be an early and intensive practitioner of manufacturing that trust about itself.

We formalize three research questions:

- **RQ1.** What mechanisms produce manufactured authority in the GEO vertical?
- **RQ2.** Is the phenomenon specific to a single market, or is it global?
- **RQ3.** If it is global, what distinguishes markets from one another, and can that difference be measured?

Our contribution is twofold. First, we describe a three-mechanism taxonomy of manufactured authority and document it across eight markets. Second, and in response to the observation that a taxonomy alone cannot be tested, we propose two operational metrics, the Manufactured Authority Index (MAI) and the Verification Layer Score (VLS), so that the phenomenon can be scored, compared, and re-measured by others. The paper also engages critically with prior work, including the author's own AIVI framework (Göktaş, 2026), whose claim that retrieval mitigates manipulation through multi-source verification we test and qualify.

2. Background and Related Work

2.1 GEO foundations and the earned-media signal

Aggarwal et al. (2024) introduced GEO as a framework for improving content visibility in generative-engine responses, distinct from traditional SEO in that it optimizes for inclusion in a synthesized answer rather than for rank position. Chen et al. (2025) extended this line with large-scale controlled experiments across multiple verticals, languages, and query paraphrases, reporting a systematic and overwhelming bias in AI search toward earned media relative to brand-owned and social content. For the present paper, this earned-media bias is the pivotal mechanism: it is simultaneously the reason legitimate GEO emphasizes third-party authority and the reason manufactured third-party authority is effective.

2.2 The dual-use turn: manipulation and poisoning

The same authority signal that powers legitimate GEO can be weaponized. Lasso Security (2026) demonstrated that GEO techniques can push a fabricated claim into AI-generated answers using only publicly available web content, without model compromise or prompt injection. In their study, a single clean website carrying one false medical claim was surrounded by a small number of fabricated "editorial" articles, in the voice of plausible niche publications, each citing the site and repeating the claim. Across 5,525 test runs on five production models and 24 techniques, this editorial-endorsement mechanism was the strongest single technique, reaching 92% success on the most susceptible model and 98% when combined with a listicle

format, while two of the five models resisted consistently. The authors emphasize an *attribution gap*: when a model reproduces a manipulated claim in its own voice without citing the originating source, the intuitive verification path, inspecting the cited source, disappears.

This is not an isolated result. Zhang, Triedman, and Shmatikov (2026) show that deep-research agents which iteratively retrieve and cite web content can be steered by a short crafted edit to a single frequently retrieved user-generated page, causing the agent to promote attacker-chosen entities across many related queries. Earlier work on retrieval-augmented generation established that a small number of poisoned documents can corrupt downstream answers (Zou, Geng, Wang, & Jia, 2025). Industry security taxonomies now codify this class of risk: the OWASP Top 10 for LLM Applications lists Misinformation, which in the 2025 edition absorbs the former overreliance category, as a first-order risk (OWASP, 2025).

The mechanism that unifies these findings is statistical: when a claim appears corroborated across multiple retrievable sources, models tend to amplify it, and this holds even when all of the sources trace to a single originator (Lasso Security, 2026). Repetition across sources is read as credibility.

2.3 The verification layer

We distinguish *earned authority* from *manufactured authority* by the presence of a verification layer: independent third-party infrastructure that can, at least in principle, separate genuine reputation from purchased or self-produced reputation. We operationalize this layer in Section 3 as five components: independent directories that verify business registration and client feedback; professional bodies and codes of conduct; objective, externally computed metrics; encyclopedic or academic reference works; and established general-interest or trade press that references named practitioners on the record. The strength of this layer, we argue and then measure, is the variable that best explains cross-market differences.

3. Method

3.1 Design

This is an exploratory, observational, cross-market audit based on qualitative content analysis of first-page search results for a canonical query, augmented by two proposed quantitative indices. It is hypothesis-generating rather than confirmatory. No controlled poisoning experiment was performed, and no measurement of whether the identified sources are actually cited in generated answers was attempted; the study characterizes what is publicly retrievable, not what any named actor intended or what any model ultimately outputs.

3.2 Markets, queries, and snapshot

Eight markets and languages were sampled: Turkey, the United States, Germany, France, Japan, Australia, Russia, and China. For each market, the canonical query was the local-language equivalent of "best GEO agency" or "best GEO expert" (the full query set is listed in Appendix A). Native-language queries were used deliberately, because generative retrieval localizes to the target-language ecosystem and an English-only query would misrepresent non-English markets. Data reflect a single snapshot taken in early July 2026, retrieved through a general web search engine. Because search and generative outputs vary with locale, personalization, and time, the audit is a cross-section, not a longitudinal measurement. To keep the denominator well defined, N was fixed at the first 10 organic results per market, with paid and other non-organic units excluded.

3.3 Mechanism coding

Each examined first-page source was classified by mechanism:

- **M1, self-referential ranking:** a "best" list published by, or on behalf of, a party that then appears at or near the top of that same list.
- **M2, syndicated corroboration:** substantially identical promotional material distributed across many distinct domains, producing the appearance of multi-source consensus from a single originator.
- **M3, purchased-placement intermediation:** directories, "ranking platforms," or quote marketplaces in which listing or position is a function of payment rather than independent assessment.

A source could exhibit more than one mechanism. Sources that were neither manufactured nor purchased were coded as independent.

3.4 Two operational metrics

Manufactured Authority Index (MAI, 0 to 10). For each market, MAI is the number of the first 10 organic results exhibiting at least one manufacturing mechanism (M1, M2, or M3), on a 0 to 10 scale (equivalently, the percentage of the ten sources). High MAI means the first page is saturated with self-produced or purchased authority.

Verification Layer Score (VLS, 0 to 10). VLS is the sum of five criteria, each scored 0 (absent), 1 (partial), or 2 (present and independent):

- **C1** Independent directories that verify business registration and client feedback.
- **C2** Professional bodies, codes of conduct, or standardization membership used as evidence.
- **C3** Established general-interest or trade press referencing named practitioners on the record.
- **C4** Academic or encyclopedic reference works (founding studies, textbooks, encyclopedic entries).
- **C5** Objective, externally computed metrics or independent citation-tracking used as evidence.

Both indices are proposed as reusable instruments. In this study they were assigned by a single author-coder from a single snapshot; the scores should therefore be read as one informed operationalization to be validated, not as settled measurements. Inter-coder reliability was not assessed.

3.5 Anonymization

All commercial entities, sites, individuals, and networks are reported under generic labels (for example, Agency A1, Syndication Network S1, Directory D1, Operator X). Countries are named, since naming a jurisdiction carries no reputational claim about a specific party. Cited scholarship is named normally, as it is source material rather than an object of study. The rationale is given in Section 7.

4. Findings

4.1 Three mechanisms (RQ1)

All three mechanisms were observed, and M1 was present in every market sampled. Figure 3 summarizes how the three mechanisms converge into a single manufactured-corroboration signal that an earned-media-biased retriever is disposed to accept.

M1, self-referential ranking. Across markets, "best GEO agency" lists were frequently authored by a party that then occupied the top position, with the longest and most favorable entry. A diagnostic consequence is

that the identity of the "number one" varied by source: different self-published lists crowned different winners, so no consistent ground truth emerged from the first page. This structural absence of a stable ranking is consistent with independent evidence that AI brand recommendations are themselves highly inconsistent between repeated runs of the same query, so that single-query rank position is not a reliable signal (Fishkin & O'Donnell, 2026); it is one reason the present study measures market-level saturation (MAI) and structural verification (VLS) rather than query-level position. In several markets the surrounding entries were filled with genuine but foreign or otherwise non-competing entities, which lends a list the appearance of breadth while removing local competition from view.

M2, syndicated corroboration. In one market (Turkey), a single promotional text appeared verbatim across three or more distinct regional news domains, presenting one originator's ranking as if multiple independent outlets had reported it. In another market (United States), a self-promoting "directory" traceable to a single originator appeared across multiple financial-news syndication domains. This mirrors, at market scale, the editorial-endorsement mechanism that Lasso Security (2026) demonstrated experimentally, and the single-source amplification that Zhang et al. (2026) demonstrated for retrieval agents: manufactured third-party agreement, read by retrieval systems as corroboration.

M3, purchased-placement intermediation. Multiple markets featured "ranking platforms" and quote marketplaces in which inclusion or position is purchased. In at least one instance, such a platform was observed being cited as a source by an AI answer surface, so that purchased placement was laundered into the model's apparent authority.

4.2 Quantifying the phenomenon (RQ2)

Table 1 reports, for each market, the approximate number of first-page sources examined, the mechanisms observed, whether any independent source appeared in the sample, and the resulting MAI. The consistent finding is that MAI is high in every market: manufactured authority dominates the first page regardless of jurisdiction.

Table 1. Mechanism presence and Manufactured Authority Index by market (July 2026 snapshot; first 10 organic results per market; D = dominant, P = present, low = marginal; MAI is the count of the 10 sources exhibiting at least one mechanism).

Market	N	M1	M2	M3	Independent in sample	MAI
Turkey	10	D	P	P	0 of 10	9
China	10	D	P	P	present but compromised	9
Russia	10	D	P	P	yes (tech press)	8
United States	10	D	P	P	yes (verified directory)	7
Australia	10	D	low	P	yes (buyer education)	7
Japan	10	D	low	P	yes (encyclopedic)	7
France	10	D	low	P	yes (named grounding)	7
Germany	10	P	low	low	yes (press, standards)	6

Figure 1 plots MAI against the Verification Layer Score developed next. The vertical compression of the points, all in the upper band, is the visual form of RQ2's answer: the manipulation is global.

4.3 Verification-layer asymmetry (RQ3)

The mechanisms are global; what separates markets is the verification layer. Table 2 reports the five-criterion VLS for each market, and Figure 2 shows how each score is composed.

Table 2. Verification Layer Score by market (five criteria, each 0 to 2).

Market	C1 directories	C2 standards	C3 press	C4 academic	C5 metrics	VLS
Germany	1	2	2	2	2	9
France	1	1	2	2	2	8
United States	2	1	1	1	2	7
Japan	1	0	1	2	1	5
Australia	1	0	1	1	1	4
Russia	1	0	2	0	1	4
China	0	1	0	1	0	2
Turkey	0	0	0	1	0	1

Where MAI is compressed high, VLS is dispersed from 1 to 9. Two markets illustrate the extremes. In the German sample, self-referential lists coexisted with named practitioners referenced by established press, a professionally published treatment of the discipline, standardization signals, and, notably, several lists that explicitly declined to rank. In the Turkish sample, by contrast, no independent verification component was identified: the first page consisted of self-referential lists, syndicated corroboration, and purchased-placement platforms only.

4.4 Two failure modes

Tables 1 and 2 together imply that "polluted" is not a single condition. Turkey and China both show high MAI, but they differ on VLS and, more importantly, on the *independence* of what verification-shaped material exists. Turkey is *undefended* (MAI 9, VLS 1): manipulation faces no independent counterweight. China's sampled results are consistent with what we term a *manufactured verification layer* (MAI 9, VLS 2): authority-shaped artifacts such as "research reports" and rankings are abundant, but many are themselves advertorial and so fail the independence criteria, alongside a smaller set of genuine regulatory-registration signals. An undefended market and a market whose defenses are largely manufactured present different risks to a downstream AI system and call for different remedies. We state this as an interpretation consistent with the sample rather than as an established property of either national market.

4.5 Cross-market operator recurrence

At least one originator (anonymized here as Operator X) was observed operating across markets: self-ranking within an English-language financial-news syndication targeting one national market, and simultaneously self-ranking within a domestic non-English market under a localized name. The recurrence indicates that manufactured authority is, in at least some cases, a coordinated multi-market operation rather than only a local practice.

5. Discussion

5.1 An immune-system reframing

The natural but misleading reading of these findings is that some national industries are simply less honest. The data do not support that reading: MAI is high even in the markets with the strongest verification layers. What varies is the market's *immune system*, its capacity to distinguish earned from manufactured authority, captured here by VLS. This reframes the phenomenon from a morality claim about actors to a structural claim about ecosystems, and it reframes the remedy from naming individual actors to building verification infrastructure.

5.2 The earned-media paradox

Chen et al. (2025) show that AI search privileges earned, third-party corroboration. Lasso Security (2026) and Zhang et al. (2026) show that fabricated or single-source-amplified third-party agreement is read as corroboration, because repetition across retrievable sources functions statistically as credibility. Mechanism M2 in our audit is structurally analogous to these results, observed in the open market rather than in a laboratory. The implication is uncomfortable: the earned-media signal that makes legitimate GEO effective is, by the same token, the attack surface. We emphasize a boundary of our evidence here: we observed manufactured sources on the first page, but we did not measure whether any generative system actually cited them in an answer. Establishing that causal link is the task of the citation-audit work proposed in Section 8, and our claims are limited accordingly.

5.3 Revisiting the AIVI mitigation claim

The author's AIVI framework (Göktaş, 2026) argued that retrieval architectures mitigate poisoning through multi-source verification, source diversity, and authority evaluation. This audit qualifies that claim. Multi-source verification presumes that "multiple sources" means "independent sources," and authority evaluation presumes that authority signals are earned rather than produced. Where a market lacks a verification layer, both presumptions fail: syndication supplies many sources from one originator (low source independence at high source count), and purchased placement supplies authority signals on demand. Mitigation is therefore not intrinsic to retrieval; it is contingent on an external ecosystem of verifiable authority. The independent poisoning results of Lasso Security (2026), Zhang et al. (2026), and Zou et al. (2025) support this qualification from the security side, so the point does not rest on the author's own prior framework alone.

5.4 Implications

For generative-system providers, the central lesson is that source *diversity* is not the same as source *independence*; deduplication and provenance analysis at the level of the originating party, rather than the domain, would blunt M2 directly. For markets and policymakers, an independent verification layer functions as a public good whose absence is a market failure that AI systems now import wholesale; VLS offers a rough instrument for identifying where that public good is missing. For practitioners, in markets where the verification layer is weak or absent, transparent methodology and on-the-record verifiability are the scarce and therefore differentiating positions.

6. Limitations

This study has substantial limitations. It is a single-snapshot cross-section, and search and generative outputs vary with time, locale, and personalization; the assignments in Tables 1 and 2 may not replicate. Classification

and both indices were produced by a single author-coder using interpretive schemes, without inter-coder reliability testing, so MAI and VLS here represent one informed operationalization rather than validated measurements. Sample sizes per market are small, limited to first-page results for a canonical query, and the choice of canonical query shapes what is seen. No controlled poisoning experiment was conducted and no citation effect was measured, so causal claims about model behavior are imported from the cited literature rather than established here. Coverage is partial, and per-market depth is uneven. Finally, some quantitative figures for the poisoning and earned-media literatures are taken from the primary reports and abstracts; readers should consult those sources for full methodological detail.

7. Ethical Considerations

The aim of this work is to characterize a structural phenomenon, not to allege wrongdoing by any identifiable party. Many of the individual mechanisms described are legal and common marketing practices; the concern is their aggregate effect on AI answer integrity, not the conduct of any single actor. For this reason all commercial entities, sites, individuals, and networks are anonymized, and only public search results were analyzed. No private data were collected. Anonymization keeps the analytical focus on the pattern and its remedy, and avoids both defamation and the elevation of any named party, favorable or unfavorable, on the basis of a single non-adjudicated snapshot.

8. Conclusion

Manufactured authority is a global, structural feature of the Generative Engine Optimization vertical rather than the failing of any one national market. Three mechanisms, self-referential ranking, syndicated corroboration, and purchased-placement intermediation, recur across markets, and the proposed Manufactured Authority Index is high everywhere. What separates markets is the strength of an independent verification layer, captured by the proposed Verification Layer Score, which ranges in this sample from 1 to 9. The earned-media signal that makes legitimate GEO work is also the surface that manufactured authority and, in the security literature, GEO poisoning exploit. Verification layers are the missing public good, and their absence is quietly imported into the answers that generative systems now give.

Key takeaways.

- In the AI era, "who is best" is increasingly answered by systems that reward apparent third-party consensus, which makes manufacturing that consensus valuable.
- The problem is not uniquely national dishonesty but asymmetric market defenses; the differentiator is the verification layer, and MAI and VLS make both axes measurable.
- Undefended markets and manufactured-defense markets are distinct failure modes requiring distinct remedies.
- Source diversity is not source independence; provenance analysis at the originator level is a concrete mitigation.

Future work. Longitudinal tracking of the same queries across repeated snapshots; multi-coder reliability testing of MAI and VLS; a controlled, ethically bounded citation audit that measures how often manufactured-authority sources are actually cited in generated answers; and extension of the two indices to other self-interested verticals beyond GEO.

Figures




Figure 1. Each market plotted by Verification Layer Score (x) and Manufactured Authority Index (y). Manufactured authority is consistently high across markets while verification capacity is dispersed.




Figure 2. Composition of the Verification Layer Score across the five 0-to-2 criteria, by market.

Figure 3. The Manufactured Authority mechanism




Figure 3. The three mechanisms converge into a manufactured third-party corroboration signal that an earned-media-biased retriever is disposed to cite, often without attribution.

Appendix A. Query set and scoring notes

Snapshot: early July 2026, general web search engine, first 10 organic results per market (N = 10; paid and other non-organic units excluded).

Canonical queries (local language): Turkey, "en iyi GEO uzmanları / ajansları"; United States, "best generative engine optimization GEO agencies"; Germany, "beste GEO Agentur Generative Engine Optimization"; France, "meilleures agences GEO optimisation moteurs generatifs"; Japan, "生成エンジン最適化 GEO 支援会社 おすすめ"; Australia, "best GEO agencies Australia"; Russia, "лучшие агентства GEO оптимизация для нейросетей"; China, "最好的 GEO 生成式引擎优化 服务商".

MAI: share of examined first-page sources exhibiting M1, M2, or M3, scaled 0 to 10.

VLS: sum of C1 through C5, each 0 (absent), 1 (partial), or 2 (present and independent), as defined in Section 3.4.

References

- Aggarwal, P., Murahari, V., Rajpurohit, T., Kalyan, A., Narasimhan, K. R., & Deshpande, A. (2024). *GEO: Generative Engine Optimization*. arXiv:2311.09735.
- Chen, M., Wang, X., Chen, K., & Koudas, N. (2025). *Generative Engine Optimization: How to Dominate AI Search*. arXiv:2509.08919.
- Fishkin, R., & O'Donnell, P. (2026). *NEW research: AIs are highly inconsistent when recommending brands or products; marketers should take care when tracking AI visibility*. SparkToro. <https://sparktoro.com/blog/new-research-ais-are-highly-inconsistent-when-recommending-brands-or-products-marketers-should-take-care-when-tracking-ai-visibility/>
- Göktaş, İ. (2026). *Generative Engine Optimization: Technical Foundations of Search Visibility in the AI Era and an Integrated Visibility Model (AIVI)*. Zenodo (preprint). DOI: 10.5281/zenodo.21180989.

Lasso Security. (2026, June 24). *Exploiting GEO to push harmful claims into AI-generated answers* (E. Saban). Lasso Security research blog. <https://www.lasso.security/blog/exploiting-geo-to-push-harmful-claims-into-ai-generated-answers>

OWASP. (2025). *OWASP Top 10 for LLM Applications (2025): LLM09:2025 Misinformation*. OWASP GenAI Security Project. <https://genai.owasp.org>

Zhang, T., Triedman, H., & Shmatikov, V. (2026). *Deep-research agents can be poisoned via user-generated content*. arXiv:2605.24245.

Zou, W., Geng, R., Wang, B., & Jia, J. (2025). *PoisonedRAG: Knowledge corruption attacks to retrieval-augmented generation of large language models*. In *Proceedings of the 34th USENIX Security Symposium*.

How to cite

Göktaş, İ. (2026). *Manufactured Authority in Generative Engine Optimization: Cross-Market Evidence and Two Operational Metrics (MAI and VLS)*. Zenodo (preprint).

About the author

İbrahim Göktaş is a consultant whose recent work focuses on AI visibility, Generative Engine Optimization (GEO), Answer Engine Optimization (AEO), and digital-authority strategy in the age of generative AI. This paper is a companion to the author's Integrated Visibility (AIVI) framework and extends it toward the integrity of the authority signals on which AI answers depend.