
MIND THE INVISIBILITY GAP: AN ANALYSIS OF 80 LEADING GLOBAL BRANDS IN LLM ANSWERS *

Niall Cook
Independent Researcher
Geometriqs, United Kingdom
niall.cook@geometriqs.com

ABSTRACT

Generative AI represents an emerging new frontier for online brand visibility, often referred to as "Share of Model". This study examines how 80 of the world's largest corporations are represented in AI-generated answers produced by three leading generative AI platforms, OpenAI, Google Gemini, and Perplexity AI. The results confirm that, as AI assistants increasingly mediate how stakeholders seek and interpret information, new benchmarks for corporate communications, reputation management, and competitive positioning are required. The study focuses on 80 companies drawn from the top 10 in each sector of the Forbes Global 2000 rankings. The analysis suggests that AI systems amplify brands with high public visibility and abundant English-language content, while industrial, regional, or non-Western firms remain underrepresented.

Keywords Generative AI · Reputation Management · Corporate Communications · Generative Engine Optimization

1 Introduction

A company's presence in AI-generated responses to user questions is emerging as a critical new battleground for brand positioning. With AI chat interfaces and search overviews providing direct answers (often without requiring a website visit), whether and how a brand appears in those answers may significantly influence stakeholder perceptions.

Traditional metrics like Share of Voice (SOV) in media and Share of Search (SOS) have long predicted market share and growth. Now, Share of Model (SOM) is emerging as a new concept – the share of attention a brand commands within generative AI answers.

Early evidence suggests that those that dominate these answer boxes will see a form of brand recall that could translate to commercial advantage. And in an age when intangibles like brand reputation can account for most of a company's market value (as much as 90 per cent of the S&P 500)[1], optimizing how AI "talks about" your company may become directly relevant to enterprise value.

Generative AI adoption has surged in 2024-25 and is becoming an increasingly important mainstream tool for decision support, a trend which can only accelerate. If AI answers become the first or only source of information, being included (or absent) can shape which companies consumers or executives even consider for purchase, partnership, or investment. While it is still early days, a brand's positioning in these answers looks set to be a source of significant competitive advantage.

Other studies already show material effects from brand visibility. YouGov found over half of users would prefer AI-generated summaries to traditional search results[2], while a Search Engine Land analysis showed Google AI Overviews cut Mail Online's click-through rates by 45 per cent[3]. Similar research on e-commerce shows losses from declining search visibility[4]. In B2B, Forrester reports that 73 per cent of buyers now use AI tools to research vendors[5], suggesting that being absent can have commercial cost.

**Citation:* Cook, N. Mind The Invisibility Gap: An Analysis of 80 Leading Global Brands in LLM Answers. DOI:000000/11111.

Independent academic work reaches similar conclusions. A 2025 University of Toronto study on Generative Engine Optimization found that AI-powered search engines systematically privilege earned, third-party media over brand-owned or social sources[6], reinforcing the importance of authoritative visibility within generative answers.

This demonstrates that even among major corporations, visibility within AI-generated content is highly uneven. The following data provides quantitative evidence of that gap.

2 Scope and Methodology

The study profiles 80 global companies across eight sectors, specifically the top 10 in each sector of the Forbes Global 2000 rankings[7]. The sectors covered are:

- Automotive
- Consumer Goods
- Energy & Utilities
- Finance
- Healthcare
- Retail & Leisure
- Technology
- Telecoms & Media

To focus on comparable peer sets, each company is mainly compared within its sector – preserving the competitive lens that makes “Share of Model” most meaningful.

We examined brand visibility across five key corporate attributes which mirror major drivers in many reputation and brand value frameworks:

- Customer Experience (CX)
- Executive Leadership
- Innovation & Sector Leadership
- Reputation & Trust
- Sustainability/ESG

These themes informed the questions asked.

2.1 Questions and LLM Platforms

A total of 250 neutral, comparative questions were designed (in English) to cover the five themes (50 per theme) and manually reviewed for neutrality. The questions emulated what an average user might ask an AI assistant when researching corporate qualities, without mentioning specific companies or sectors by name (for example, “Which companies have the most trusted CEO reputations” or “Which companies are regarded as exemplars of customer-centric culture and why?”). These prompts were designed to emulate a general information-seeking user rather than a specific persona – subsequent research may choose to further segment by user type and intent.

We posed each question to three leading LLM platforms – OpenAI (ChatGPT), Google’s Gemini (as used in Google’s AI search overviews), and Perplexity AI (a retrieval-focused answer engine) – capturing a broad triangulation of how AI delivers answers. For each platform, we tested answers in two modes:

1. **Training-only mode** (no web access) – the model answers from its internal knowledge, with no live internet retrieval or citations.
2. **Web-search grounded mode** – the model is allowed to perform live web searches and cite sources in its answer.

In total, 750 question-runs per mode (250 prompts given to each of the three LLMs) were conducted, and the responses analysed for company (and alias) mentions. This yielded roughly 4,000 unique mentions of the 80 target companies in 1,292 of the 1,500 collected answers. Each mention was logged along with metadata: which model produced it, which

mode, the prompt category (attribute theme), any source cited, and a sentiment tone classification (positive, neutral, or negative sentiment of the reference to the brand).

The resulting dataset provides a rich basis for analysing how often and in what context each company is mentioned by generative AI, enabling comparisons across sectors, attributes, and AI platforms.

Limitations of the study and suggestions for future research are discussed in the Appendix.

2.2 Data Access

A sample of the prompts used in each category are included in the appendix. All anonymised prompts and AI responses used in this analysis are publicly available for independent verification and re-use at:

https://huggingface.co/datasets/geometriqs/global80_prompt-response-pairs

3 Key Findings

3.1 Overall corporate visibility

- Across all three generative AI models, the top 80 brands in the study saw an average visibility of only 4 per cent, indicating that AI answers don't necessarily prioritise scale or recognition.
- Surprisingly, of the 80 global companies analysed, one in five firms were entirely invisible.
- Among the eight sectors analysed, technology and consumer-facing industries dominate: tech companies accounted for nearly 40 per cent of all mentions. Consumer goods added another 20 per cent. By contrast, finance and energy firms together made up less than 5 per cent, despite their financial size.
- This analysis suggests that AI systems amplify brands with high public visibility and abundant English-language content, while industrial, regional, or non-Western firms remain underrepresented.

3.2 Visibility by corporate attribute

- Questions about Innovation generated the most mentions (36 per cent of total).
- Executive Leadership and Customer Experience drew moderate attention.
- Reputation & Trust and Sustainability/ESG produced the fewest mentions.
- AI models will often cite a few widely recognized exemplars – e.g., Unilever for ESG – rather than offering broad lists of companies.

3.3 Top brands

- Microsoft, Apple, and Amazon led the rankings, appearing in over one-third of all answers. Together with Unilever and Alphabet, these five brands accounted for more than half of all mentions.

3.4 Invisible brands

- Conversely, 16 major corporations – including Saudi Aramco, ICBC, and China Mobile – were completely absent, revealing possible structural biases toward Western, English-language, and consumer-oriented companies.
- This absence often reflects content gaps, not irrelevance: AI models can only mention brands they have been trained on or exposed to through public data.

3.5 Differences between AI platforms

- OpenAI's ChatGPT delivered the broadest brand coverage (59 of 80 companies) and the most mentions overall.
- Google Gemini excelled in depth and sourcing, expanding visibility when web grounded.
- Perplexity AI seemed more reliant on live retrieval and surfaced additional brands only when allowed to search online.
- Notably, 58 per cent of all mentions came from models' internal knowledge, showing that being embedded in a model's training data remains critical, while web visibility can supplement but not replace that foundation.




Figure 1: Share of AI Answer Mentions by Sector

3.6 Sentiment

- AI references were overwhelmingly positive (68 per cent) or neutral (32 per cent).
- Negative mentions were rare (<1 per cent).
- The reputational risk lies not in criticism but invisibility: unmentioned brands simply don't enter consideration sets; and low weighting on key corporate attributes may impact brand reputation negatively versus competitors.

4 Detailed Results

4.1 Overall Visibility

Out of the 80 companies studied, 64 companies (80 per cent) were mentioned at least once in an AI-generated answer. However, 16 companies (20 per cent) were completely “invisible” – not appearing in any answer. These invisible cases are examined in detail later in this report, but it is notable that most are non-Western firms, highlighting potential content biases.

Across all models and questions, the total count of company mentions was 3,919. This equates to an average of 61.2 mentions per company, or about a 4.1 per cent “recall rate” (since each brand had 1,500 opportunities to be named). An average presence in only around 4 per cent of relevant AI answers is low. By comparison, traditional brand recall surveys (within a category) often find unaided recall in the 20–80 per cent range. However, that average masks a wide spread: some of the 80 brands achieved much higher visibility, while others were never mentioned at all.

To put the share-of-model in perspective, the most-referenced brands appeared in over one-third of all AI answers, whereas one-fifth of the companies garnered zero mentions. In aggregate, a small set of companies accounted for a large portion of the mentions, indicating a skewed visibility distribution. The chart in Figure 1 illustrates the share of total mentions by sector, revealing which industries' top players are capturing the most “mindshare” in AI responses.

As shown above, Technology firms led by a wide margin in share of AI mentions (39.5 per cent of all mentions were of tech companies). This reflects the cross-domain prominence of tech giants – they surfaced as examples in answers




Figure 2: Total Brand Mentions by Question Theme

about innovation, leadership, customer experience, and even ESG. Consumer Goods brands (food, beverage, FMCG, etc.) were the second most referenced sector at 20 per cent of mentions, bolstered by strong showings in reputation and sustainability-related queries. Next came Retail & Leisure (15 per cent), which includes e-commerce and retail giants often cited for customer experience leadership.

The Automotive sector garnered 10 per cent of mentions (with a focus on innovation and leadership within autos). Healthcare (7 per cent) and Telecoms & Media (around 3.6 per cent) were relatively minor presences. Strikingly, Finance and Energy companies were scarcely referenced at all – only 2.9 per cent and 1.8 per cent of total mentions, respectively – despite some of these firms being among the world’s largest by revenue or market cap.

The data clearly indicates that certain sectors are currently far more “visible” in AI discourse than others. Tech-oriented and consumer-facing brands enjoy a disproportionately high share of model attention, whereas more industrial or regionally-focused sectors risk being overlooked in generic AI answers.

4.2 Visibility by Attribute

The study’s five question themes produced varying levels of brand visibility. Certain topics prompted AI to cite many company examples, while others yielded relatively few. Figure 2 shows the total number of brand mentions generated under each corporate attribute category.

Innovation emerged as the most fruitful ground for cross-company comparisons, accounting for 36 per cent of all mentions. Many of these mentions were tech companies, consistent with technology’s outsize presence noted earlier. In contrast, Reputation-focused questions (about trustworthy or reputable companies) and ESG (sustainability leaders) produced relatively fewer brand mentions. This could be because the AI answers on these topics leaned on a narrower set of well-known exemplars. For example, a handful of companies consistently recognized for ethics or sustainability, like Unilever, might come up repeatedly.

Customer Experience and Executive Leadership fell in between: plenty of companies were cited as having great customer service or strong leadership, but the answers often concentrated on specific sectors (e.g. retail for customer experience, or tech/auto for visionary CEOs). The key takeaway is that the visibility of brands varies by topic –




Figure 3: Top 10 Most-Mentioned Companies in AI Answers

if a company's strengths lie in innovation, it has a better chance of being mentioned than if its strength is in, say, sustainability, where AI might only mention a few leaders.

Notably, even within a given theme, certain sectors or geographies dominated the answers. For example, innovation answers heavily featured U.S. and global tech firms regardless of sector, whereas sustainability answers often highlighted consumer goods companies with well-publicised ESG initiatives. This suggests a possible bias in AI knowledge toward companies and stories that enjoy broad, multi-sector Anglo-US media coverage. Less glamorous industries or non-Anglophone companies did not surface as readily, even if they are massive. This pattern is explored further under "invisible" brands.

4.3 Brand Leaders and Laggards

Looking at individual company performance, a handful of "leader" brands achieved especially high visibility across the AI models. Figure 3 highlights the Top 10 most-mentioned companies.

This heavy skew means that, from a "Share of Model" perspective, the top few brands are capturing outsized mindshare in AI-generated content. For instance, Microsoft alone was mentioned in 14.3 per cent of all brand mentions (and in over a third of all answers), indicating it is almost synonymous with concepts like innovation, leadership, and technology as far as generative AI models are concerned.

In contrast, brands outside the top tier may only appear a handful of times or not at all. For communications officers, this raises a strategic question: *'If our brand is not one of these default answers, how can we improve our representation in AI outputs?'* The answer may lie in addressing content gaps.

4.4 'Invisible' Brands

On the other end of the spectrum, one in five of the studied companies were completely absent from the AI answers. Table 1 lists these 16 "invisible" brands and their sectors.

Table 1: 'Invisible' Companies in GenAI Global80

Sector	Companies with no AI mentions
Finance (Banking)	ICBC; China Construction Bank; Agricultural Bank of China; Bank of China
Telecoms & Media	China Mobile; Nippon Telegraph & Telephone (NTT); SoftBank; China Telecom; Charter Communications
Retail & Leisure	PDD Holdings (Pinduoduo); Lowe's
Energy & Utilities	Saudi Aramco; PetroChina; Reliance Industries; Sinopec; CNOOC

Sixteen of the 80 companies did not appear in any AI answer. Notably, several are among the world's largest businesses yet the AI models never mentioned them in response to the generic questions.

The prevalence of Chinese and other non-Western companies in the 'invisibles' list hints at underlying biases. These brands' absence could stem from several factors: language and coverage biases. The AI models – especially in English – may have less knowledge or fewer sources about companies that don't feature prominently in English-language content), or a communications performance gap, e.g. in content strategy (those firms might not have as much readily accessible, AI-digestible content highlighting their leadership or innovation, etc.) It's also possible some of these companies maintain a deliberately low public profile.

However, it is important to stress that invisible does not mean insignificant: a company with a zero presence in AI answers may be very important and influential. Rather, it signals a discoverability gap in AI. For whatever reason – accidental, strategic, or coverage bias – the AI model has not been taught or does not access information that surfaces the company in response to general queries.

From a reputation management perspective, this is a red flag. If stakeholders increasingly rely on AI assistants, an otherwise large company could be absent from consideration sets simply because it is absent from the answers.

But why were these 16 brands missed? In many cases, the AI responses favoured companies with well-known consumer brands or those frequently lauded in global business discourse. For example, in banking and telecoms, the models often mentioned western firms (like JPMorgan or AT&T) even when asked general questions, bypassing bigger-but-less-publicized Chinese counterparts.

This points to what we might call 'structural invisibility' – the AI's training data and web content skews toward certain markets, languages and narratives. Generative models privilege information that is abundant and easily synthesizable. If a company lacks "ownable, high-authority explanations" online for the topics in question, the models have little to cite or say about it.

The implication for communicators is that improving AI visibility may require proactively filling content gaps – e.g. publishing accessible thought leadership, case studies, or explanatory content that AI can ingest and regurgitate when relevant. This does not mean simply issuing more press releases, but creating content tailored to be picked up in LLM answer synthesis.

Overall, the study reveals a stark leader-and-laggard dynamic in AI visibility: a few brands enjoy a strong share of model mindshare, while a significant minority are effectively invisible. Organizations need to evaluate where they fall on this spectrum and consider steps to fill any gaps.

4.5 Differences Between AI Models and Modes

Current public usage patterns differ markedly between models. Third party estimates indicate that ChatGPT accounts for roughly 60-65 per cent of generative AI query traffic, Gemini around 25 per cent, and Perplexity 5-10 per cent[8]. This suggests that presence in ChatGPT currently contributes most to overall exposure, though Google's reach via Gemini is growing fast as AI overviews roll out in Search, and Perplexity's new Comet web browser may start to flatten these differences.

Notwithstanding this, our analysis uncovered interesting differences in how each of the three AI platforms performed, as well as how web-grounding affected the answers. Although all three LLMs (OpenAI's GPT, Google's Gemini, and Perplexity AI) were given the same prompts, they did not respond identically.

OpenAI (ChatGPT) produced the most brand mentions overall (1,461 out of 3,919, about 37 per cent of the total). It also referenced the broadest set of companies (59 out of 80). ChatGPT's knowledge cutoff meant its answers leaned on well-known examples present in its training corpus. It often provided list-style answers naming multiple companies,

which boosted mention counts. Notably, ChatGPT's answers were more likely to be neutral in tone (see below for more on sentiment). With web access enabled, OpenAI did incorporate live info for some answers, but interestingly this mode did not dramatically increase the diversity of companies – it appears ChatGPT's base knowledge already covered most companies it was going to mention.

Google Gemini accounted for roughly 34 per cent of mentions (1,337 mentions) and cited 54 distinct companies. Gemini (as used in Google's Search Generative Experience) tended to generate highly detailed answers and, when allowed, pulled in information from the web frequently. In fact, Gemini showed the largest increase in company coverage when using web grounding – it mentioned 9 more companies with live search than it did with its training-only mode. Its answers often highlighted tech and well-covered brands, and Gemini slightly out-scored others in mentioning top tech firms (e.g., it named Microsoft and Apple even more frequently than ChatGPT did, as seen in Figure 3). Overall, Gemini's behavior reflects its design for providing rich, sourced answers; it may bias toward companies that appear in authoritative web sources (news, rankings, etc.) for a given query.

Perplexity AI contributed about 29 per cent of mentions (1,121 mentions) and covered 52 companies. Perplexity's model, being retrieval-centric, was the most dependent on web grounding to pull in examples. In a no-web (training-only) setting, Perplexity mentioned only 35 companies; with the web enabled, that number jumped to 49. This indicates that Perplexity's training model alone had a narrower knowledge of the brands, but given the ability to search, it could find and include many more. However, even with search, Perplexity's answers were sometimes succinct, often citing one or two key sources – this meant it might list fewer companies per answer on average, contributing to its lower total mentions. It's also worth noting that Perplexity's answer content skewed positive, reflecting either the tone of its sources or an optimistic answer style.

In terms of web-grounded vs. training-only performance, one might assume enabling web search would increase brand visibility across the board. But the reality was more nuanced. In our data, most brand mentions (58 per cent) came from the models' internal knowledge (i.e. answers given without using the web). Using live web search contributed the remaining 42 per cent. In some cases, web grounding helped surface a brand that would otherwise not appear. For example, Bank of America was rarely mentioned by any model's training-only mode, but when allowed to search, the models found lists (e.g., of largest banks or best customer service in banking) that included Bank of America, raising its presence. A few brands (like certain auto or telecom companies) saw similar boosts only in web mode.

On the other hand, for the top-tier brands (Apple, Microsoft, Amazon, etc.), the models didn't need web lookup to mention them – those appeared readily from training data. Interestingly, because the web-grounded answers sometimes focused on a single source or article, they could reduce the breadth of companies mentioned (an article might highlight "top 3" companies where the model training data might otherwise have provided a longer list). Thus, web access changed the mix of which brands were mentioned but did not universally increase the count. It introduced some previously unseen names, while possibly narrowing the scope of some answers.

From a strategic communications viewpoint, this suggests that being prominent in authoritative web content (news articles, rankings, "top companies" lists, etc.) can directly impact whether a model finds and mentions a brand in a grounded answer. Some companies that were absent in training-only answers did appear once the AI could search – implying that fresh, well-SEOed content can make a tangible difference in AI visibility.

Each model has a different propensity to rely on such content: Perplexity, by design, will surface what it finds online, whereas ChatGPT might only use it if its internal knowledge is lacking. Ensuring a brand's narrative is present in the digital sources these models draw from is key to boosting its share of model.

4.6 Sentiment of Brand Responses

Beyond whether brands were mentioned, we examined how they were characterized. Each mention was classified for sentiment (positive, neutral, or negative connotation in context). The overall sentiment distribution skews heavily positive, as shown in Figure 4. Most company mentions were either positive (68 per cent) or at least neutral (32 per cent) in tone. Even when the prompts inquired about challenges or negative events, the answers tended to frame companies in a positive or recoverable light.

In practical terms, this means that when generative AI does talk about a company, it is usually in a favourable or neutral context (e.g., citing them for an achievement, leadership, or positive attribute). There is an apparent "politeness" or positivity bias in these AI answers. For corporate communications, this might sound like good news: 'AI isn't trashing my brand'. However, it also means absence, not a negative mention, is currently the bigger risk. Stakeholders using these AI tools aren't hearing bad things about the invisible brands; they're just not hearing about them at all. In classic reputation terms, being ignored can be as damaging as being criticized, especially if competitors are being praised in your stead. The lack of negative commentary could change as AI models integrate more real-time news, which




Figure 4: Sentiment Distribution of AI-Generated Brand Mentions

could include scandals or failures. For now, though, AI-generated content about large companies is benign to positive, focusing on accomplishments, accolades, and leadership – reinforcing the reputations of those who are mentioned.

It is worth noting some subtle differences between models. OpenAI’s answers contained a higher proportion of neutral statements (it often listed facts or programs without explicit praise), whereas Google’s and Perplexity’s answers more often used upbeat language (perhaps drawn from press releases or positive articles they retrieved). However, in general, none of the platforms tended to volunteer negative information unless explicitly asked – and our prompt set avoided direct negative questions, like “Which companies are worst at X?”. Thus, the sentiment profile here is largely a product of question design and model behaviour aiming to highlight leaders, not laggards.

5 Discussion

This analysis illuminates a new dimension of competitive communication: visibility in AI-generated information spaces. Just as companies have long tracked their share of voice in media or share of search engine queries, they must now monitor and cultivate their Share of Model – i.e. how often and how favourably they appear when AI platforms answer questions in their domain.

The findings carry several implications:

Commercial and reputational impact: If competitors consistently “own” the top answers that AI assistants give, they effectively own the mindshare of future customers, employees, or investors seeking information. Over time, this could influence preferences and behaviour, much as being top-of-mind in traditional media or search does. As generative AI becomes a common intermediary in decision journeys, a strong Share of Model is likely to correlate with real-world outcomes – be it attracting talent, winning business, or driving sales.

The cost of invisibility: For the 20 per cent of companies that are currently invisible in our study, the risk is missing out on consideration entirely. Being absent from AI answers means being absent from shortlists. Business-to-business brands should heed this as much as consumer brands – enterprise buyers and other stakeholders will also be using these tools. In B2B, where reputation can heavily influence a shortlist of vendors or partners, failing to appear in an AI-curated list of leaders or innovators may mean lost opportunities.

New content strategy for Generative AI: To improve a brand’s Share of Model, traditional PR may not be enough; what’s needed is authoritative, referenceable content that AIs can easily ingest and trust. This might include thought leadership articles, whitepapers, high-quality blog content, industry reports, or third-party recognition that highlights

the company's strengths in the relevant attribute areas. The content should be published in a way that's accessible (search-indexed) and of high credibility (so that AI models or their retrieval algorithms deem it worthy of citing). In essence, ask: If an AI is answering "Who leads in X?", do we have content out there that convincingly shows it could be us? If not, that is a gap to close.

Monitoring and metrics: Going forward, AI presence will become a standard metric. This will mean periodically querying popular AI assistants with relevant brand and industry questions and measuring mentions and context. Just as SEO became a discipline in response to search engines, Generative Engine Optimisation (GEO) or AI Optimisation (AIO) will become a key focus for communications and marketing teams.

Early mover advantage: Being an early mover in understanding and influencing how AI systems learn about brands can confer an edge. For example, right now, ensuring Wikipedia and Wikidata entries are comprehensive and up to date is important since models often draw on these. While the generative models today are mostly un-curated, this field is evolving, and brands need to keep on top of changes.

This study highlights a frontier that communications and marketing professionals should not ignore. They need to think now about how their brands can achieve a healthy share of attention in AI-generated content. Those that do will enjoy amplified reputation benefits – being seen as leaders, innovators, and reference points whenever an AI responds to a relevant question. Those that don't may find themselves eclipsed in the algorithms, their decades of real-world achievements notwithstanding. On the other hand, by recognizing that invisibility in AI equals discoverability gaps and proactively addressing these, companies can position themselves to be as prominent in the answer layer as they are in the real world.

6 Limitations and Future Research

This study does not claim access to proprietary user prompt logs. It simulates realistic information-seeking behaviour by posing standardised, neutral questions to leading generative AI platforms. The objective was not to model user volume but to measure how these systems represent some of the world's largest brands when asked comparable questions. The methodology therefore evaluates model behaviour, not user behaviour, and the results should be interpreted as evidence of representational bias, not traffic share.

The prompts were written in natural language to approximate the kinds of comparative and informational questions users typically ask. However, no synthetic dataset can perfectly reproduce the full variety of real-world prompt styles, intents, or languages. The results therefore reflect how models respond to a consistent test set rather than how all users interact with them. In some cases, the models were also asked to limit the number of companies they mentioned. Future work could integrate anonymised behavioural prompt datasets to test how real usage patterns align with these synthetic findings.

The study focuses on 80 companies drawn from the top of the Forbes Global 2000 list. This captures a meaningful cross-section of major corporate brands but excludes mid-market and emerging-market firms. Visibility patterns among smaller or regionally focused brands may differ. Also, whilst a dictionary of company aliases was used to identify mentions, it may not have captured every single variation that each LLM might have used. Finally, it must be acknowledged that the data analysed in this study only included the 1,292 (out of a possible 1,500) responses where any of the 80 target companies was mentioned. Later research could address these limitations, and extend the analysis to these cohorts and explore whether brand size, geography, or sector maturity affects share-of-model outcomes.

Eight sectors were chosen for comparability, but the boundaries between them – particularly in conglomerates and multi-line corporations – can blur. Sector classification may therefore under- or over-represent certain activities. Expanding sector granularity or applying alternate taxonomies could refine future analyses.

Finally, while the results highlight visibility gaps between sectors and regions, this study does not attempt to de-bias or normalise the underlying model outputs. Each model's training data and web-retrieval behaviour introduce inherent skew. Further research could quantify these factors more directly by repeating the experiment across multilingual prompts, additional models, or successive model versions to track drift over time.

7 Conclusion

Not every company will care where it ranks in AI search. Saudi Aramco or ICBC are unlikely to lose investors because ChatGPT fails to mention them. But for most global brands – especially those competing for consumer attention, talent, or reputation – the new visibility layer inside AI answers deserves attention. When the customer journey starts and ends

within a single conversation, ChatGPT, Google AI Overviews, or Perplexity become gatekeepers of relevance. If an AI assistant cites your competitor and not you, perceptions begin to shift.

Share of Model may not yet be a proven predictor of commercial performance, but it is a new indicator of presence in the places where opinions are increasingly formed. As generative AI becomes a routine part of research and discovery, a brand's visibility in these environments is likely to affect what people see, and therefore what they know or believe about it. As this study shows, the risk is not negative coverage but absence: invisible brands forfeit cognitive and commercial space to those the models habitually name.

The levers to change this are practical. Brands can:

- Ensure authoritative, machine-readable content about their achievements exists in the sources that feed model knowledge graphs.
- Monitor where their brand is surfaced or ignored through GEO dashboards and benchmarking data.
- Create content strategies that target missing attributes and high-authority sources so that AI systems have credible material to retrieve and synthesize.

Some companies may safely ignore their position in AI visibility today. Most cannot. The next frontier in reputation management is not just being talked about by people, but being mentioned by machines.

Acknowledgments

The author gratefully acknowledges Anthony Burgess-Webb for his assistance in shaping the study, and thanks a small group of independent reviewers whose comments helped refine the analysis.

References

- [1] Intangible asset market value study. <https://oceantomo.com/intangible-asset-market-value-study>, July 2019. Accessed: 2025-11-3.
- [2] Should AI-generated search summaries replace traditional search? <https://yougov.com/en-us/articles/51748-should-ai-generated-search-summaries-replace-traditional-search>. Accessed: 2025-11-3.
- [3] <https://searchengineland.com/google-ai-overviews-mail-online-ctr-drop-455393>. Accessed: 2025-11-3.
- [4] John Ologunbe and Ebenezer Obafemi Taiwo. The importance of SEO and SEM in improving brand visibility in e-commerce industry; a study of decathlon, amazon and ASOS. *SSRN Electron. J.*, 2023.
- [5] B2B buyer adoption of generative AI. <https://www.forrester.com/report/b2b-buyer-adoption-of-generative-ai/RES181769>. Accessed: 2025-11-3.
- [6] Mahe Chen, Xiaoxuan Wang, Kaiwen Chen, and Nick Koudas. Generative engine optimization: How to dominate AI search. 2025.
- [7] The global 2000: Top 200. <https://www.forbes.com/sites/forbesstaff/2025/07/02/the-global-2000-top-200/>, July 2025. Accessed: 2025-11-3.
- [8] Evan Bailyn. Top generative AI chatbots by market share – october 2025. <https://firstpagesage.com/reports/top-generative-ai-chatbots>, October 2025. Accessed: 2025-11-3.

8 Appendix

8.1 Sample Prompts

A random selection of the 250 prompts used appear in Table 2. All anonymised prompts and AI responses used in this analysis are publicly available for independent verification and re-use at:

https://huggingface.co/datasets/geometriqs/global80_prompt-response-pairs

8.2 Companies and Brands

The 80 companies analysed and their sector categorisation appears in Table 3.

Table 2: Random Selection of Prompts

Category	Prompt
Reputation & Trust	Which brands are most respected for ethical treatment of partners and suppliers; list 5 and one supplier policy each.
Reputation & Trust	Which brands have improved public trust by increasing operational transparency; list 5 and the transparency measure taken.
Reputation & Trust	Which companies have the most trusted CEO reputations; list CEOs and one leadership trait linked to trust.
Reputation & Trust	Which brands are most associated with trustworthy innovation practices; list 5 and one responsible innovation practice.
Innovation	Which companies are frequently cited for market disruption and why?
Innovation	Which companies have notable approaches to monetization and pricing innovation?
Innovation	Which organizations are most effective at translating R&D investment into market success?
Innovation	Which brands have notable track records in community-driven innovation or co-creation?
Customer Experience	Which companies are good examples of embedding accessibility into product and service design from the start?
Customer Experience	Which companies are regarded as exemplars of customer-centric culture and why?
Customer Experience	Which companies lead in using personalization to improve customer satisfaction, and how do their approaches differ?
Customer Experience	Name companies that are effective at reducing customer effort and the tactics they apply.
Executive Leadership	List 6 companies noted for executive teams that drive innovation and name one structural choice each has made to support ideas.
Executive Leadership	Compare 4 companies that have strong multi-generational leadership transition plans and summarize one practice that eased the handover.
Executive Leadership	List 5 companies known for strong reputational recovery after a public issue and one executive-led step in the recovery process.
Executive Leadership	Identify 5 companies that manage reputational risk through proactive stakeholder mapping and one leadership action that resulted from the mapping.
Sustainability/ESG	Which organizations lead in setting credible net-zero commitments? Provide 3-5 examples with a brief note on why each is credible.
Sustainability/ESG	Which companies are leaders in circular economy practices like reuse, repair, or take-back schemes? List 3-6 and the practice they're known for.
Sustainability/ESG	Which organizations are recognized for strong alignment between purpose, mission, and ESG strategy? List 3-5 and the alignment example for each.
Sustainability/ESG	Which organizations are known for clear, comparable ESG metrics and KPIs? List 3-5 and one metric they emphasize.

Table 3: The 80 Companies Analysed and their Sector Categorisations

Sector	Companies
Automotive	BMW, BYD, Ford, General Motors, Honda, Hyundai, Mercedes- Benz, Tesla, Toyota, Volkswagen
Consumer Goods	Anheuser-Busch InBev, British American Tobacco, Coca-Cola, L'Oréal, LVMH, Nestlé, PepsiCo, Philip Morris International, Procter & Gamble, Unilever
Energy & Utilities	Chevron, CNOOC, Enel, ExxonMobil, PetroChina, Reliance Industries, Saudi Arabian Oil Company, Shell, Sinopec, TotalEnergies
Finance	Agricultural Bank of China, Bank of America, Bank of China, Berkshire Hathaway, China Construction Bank, HSBC, ICBC, JPMorgan Chase, UnitedHealth Group, Wells Fargo
Healthcare	AbbVie, AstraZeneca, Elevar Health, Eli Lilly, Johnson & Johnson, Merck & Co, Novartis, Pfizer, Roche Holding, Sanofi
Retail & Leisure	Alibaba, Amazon, Costco, CVS Health, Home Depot, JD.com, Lowe's, McDonald's, PDD Holdings, Walmart
Technology	Alphabet, Apple, Meta, Microsoft, NVIDIA, Oracle, Samsung, Sony, Taiwan Semiconductor, Tencent Holdings
Telecoms & Media	AT&T, Charter Communications, China Mobile, China Telecom, Comcast, Deutsche Telekom, Nippon Telegraph & Tel, Softbank, The Walt Disney Company, Verizon