

WORKING PAPER

Exista.io Research Series · February 2026

Beyond AEO: The AI Visibility Stack and the Era of Agent Discovery Optimization

A Framework for Enterprise Visibility in the Agentic AI Ecosystem

Gabriela Marco

Exista.io — Buenos Aires, Argentina

gabriela.marco@exista.io | exista.io

ABSTRACT

The emergence of autonomous AI agents as active participants in B2B commerce requires a fundamental revision of how enterprises conceptualize digital visibility. Existing disciplines — Search Engine Optimization (SEO), Answer Engine Optimization (AEO), and Generative Engine Optimization (GEO) — address visibility in contexts where a human user ultimately drives the decision process. This paper argues that a new category of visibility challenges arises when autonomous agents, acting on behalf of organizations, evaluate and select suppliers, initiate negotiations, and execute transactions without direct human involvement at the moment of decision. We introduce the AI Visibility Stack, a three-layer conceptual framework that maps the complete spectrum of enterprise visibility in AI-mediated environments: (1) LLM Visibility, addressed by AEO and GEO; (2) Agent Discoverability, addressed by the newly proposed discipline of Agent Discovery Optimization (ADO); and (3) Agent Participation, addressed by Agent Presence Optimization (APO). We describe the technical infrastructure underpinning this ecosystem — principally Google's Agent2Agent (A2A) protocol and Anthropic's Model Context Protocol (MCP) — and propose the ADO Score as a composite metric for measuring agentic discoverability. The framework is contextualized against projected market data indicating that AI agents will intermediate over \$15 trillion in B2B spending globally by 2028. Implications for enterprise digital strategy, particularly for B2B companies in Spanish-speaking markets, are discussed.

Keywords: agentic AI; agent discovery optimization; AI visibility; generative engine optimization; B2B commerce; Agent2Agent protocol; multi-agent systems; enterprise marketing strategy

JEL Classification: M31 (Marketing); M15 (IT and Management); L86 (Information and Internet Services)

Suggested Citation:

Marco, G. (2026). Beyond AEO: The AI Visibility Stack and the Era of Agent Discovery Optimization (Exista.io Working Paper, February 2026). Buenos Aires: Exista.io. Available at SSRN: [URL]

1. INTRODUCTION

The last decade of digital marketing strategy has been shaped by a stable assumption: the ultimate decision-maker in any commercial transaction is a human being who uses digital tools — search engines, review platforms, AI assistants — to inform their choices. Optimization disciplines have evolved in response to each new tool: Search Engine Optimization (SEO) addressed the era of algorithmic ranking; Answer Engine Optimization (AEO) and Generative Engine Optimization (GEO) emerged to address the era of large language models (LLMs) that synthesize answers rather than return lists of links.

This assumption is now being challenged at its foundation. The deployment of autonomous AI agents — systems that can perceive objectives, plan multi-step actions, and execute decisions without continuous human supervision — has introduced a new class of actor into commercial ecosystems. These agents do not search Google; they do not ask ChatGPT for recommendations in the way a human would. They query structured data sources, fetch machine-readable configuration files, cross-reference knowledge graphs, and evaluate potential partners against formalized criteria — all without a human reading a single word of marketing copy.

\$15 trillion in B2B spending by 2028, with 90% of B2B purchases handled by AI agents within three years (Gartner, 2025). McKinsey estimates that AI agents could mediate \$3 to \$5 trillion in global commerce by 2030 (McKinsey, 2026). Visa and Mastercard have already launched agentic commerce infrastructure in Latin America for 2026 deployments (Mastercard, 2025; Visa, 2025).

Against this backdrop, the existing frameworks for enterprise digital visibility are insufficient. A company that is well-optimized for human-facing AI queries (Layer 1) but invisible to autonomous agents (Layer 2) faces a structural disadvantage in markets where agents intermediate procurement decisions. A company whose brand is widely cited by LLMs but lacks a published, machine-readable agent interface (Layer 3) cannot participate in the emerging agent-to-agent transaction economy.

This paper makes the following contributions:

- We introduce the AI Visibility Stack, a three-layer framework that maps the complete spectrum of enterprise visibility in AI-mediated environments.

- We define Agent Discovery Optimization (ADO) as a distinct discipline addressing Layer 2 visibility for autonomous agents.
- We define Agent Presence Optimization (APO) as a distinct discipline addressing Layer 3 participation in agentic ecosystems.
- We propose the ADO Score as a composite metric for measuring agentic discoverability.
- We describe the technical infrastructure — primarily the A2A and MCP protocols — that operationalizes these concepts.

The remainder of this paper is organized as follows. Section 2 situates our framework within the existing literature on digital visibility optimization and agentic AI. Section 3 presents the AI Visibility Stack in detail. Sections 4 and 5 develop the ADO and APO disciplines respectively. Section 6 describes the enabling technical infrastructure. Section 7 discusses implications for B2B enterprises, with particular attention to Spanish-speaking markets. Section 8 concludes.

2. BACKGROUND AND RELATED WORK

2.1 The Evolution of Digital Visibility Disciplines

Search Engine Optimization emerged as a formal discipline in the mid-1990s in response to the dominance of algorithmic search engines — principally Google — as the primary mechanism through which businesses were discovered online (Brin & Page, 1998). SEO optimizes for ranking signals that determine position in a list of results returned to a human user.

The introduction of featured snippets by Google (2014) and, more significantly, the mass deployment of large language model-based conversational AI (2022–2023) created a new paradigm: rather than returning lists of links, AI systems synthesize responses and recommend sources directly. This gave rise to Answer Engine Optimization (AEO), a term attributed to Jason Barnard (Kalicube, 2018), and the related discipline of Generative Engine Optimization (GEO), introduced by Aggarwal et al. (2023) in research demonstrating that specific content strategies could increase citation rates by 30–115% in LLM-generated responses.

Both AEO and GEO share a defining characteristic: they optimize for visibility in contexts where a human user formulates a query and an AI system responds. The human remains the ultimate decision-maker; the AI is an intermediary that influences, but does not replace, human judgment.

2.2 Agentic AI and the Emergence of Machine-to-Machine Commerce

The concept of autonomous software agents is not new; it has roots in distributed AI research dating to the 1980s (Wooldridge & Jennings, 1995). What is new is the combination of capable general-purpose LLMs with tool-use frameworks that allow these models to take actions in digital environments — browsing the web, querying APIs, executing code, and communicating with other systems — in pursuit of defined objectives (Yao et al., 2023; Wang et al., 2024).

Recent industry deployments have moved agentic AI from research settings into production procurement environments. Zycus's Merlin platform, for instance, uses autonomous negotiation agents that shortlist vendors, evaluate quotes, and conduct multi-round negotiations without human intervention (Zycus, 2025). McKinsey reports that autonomous category agents can capture 15–30% efficiency improvements in procurement (McKinsey, 2025). Forrester predicts that 20% of B2B sellers will engage in agent-led quote negotiations in 2026 (Forrester, 2025).

The standardization of inter-agent communication through open protocols — particularly Google's Agent2Agent (A2A) protocol (Google, 2025) and Anthropic's Model Context Protocol (MCP) (Anthropic, 2024) — represents a critical infrastructural development that enables the enterprise-scale deployment of multi-agent systems. These protocols are detailed in Section 6.

2.3 The Gap in Existing Frameworks

To our knowledge, no existing academic or practitioner framework addresses the optimization of enterprise visibility specifically for autonomous agent evaluation. The literature on AEO and GEO (Aggarwal et al., 2023; Liu et al., 2024) addresses human-mediated queries. The literature on multi-agent systems (Wooldridge, 2009) addresses agent architecture but not enterprise marketing implications. The practitioner literature on agentic commerce (McKinsey, 2026; Forrester, 2025; Bain, 2025) identifies the trend but does not provide a systematic framework for enterprise response.

The AI Visibility Stack fills this gap by integrating the existing AEO/GEO literature (Layer 1) with new frameworks for agent-specific visibility (Layers 2 and 3).

3. THE AI VISIBILITY STACK: A THREE-LAYER FRAMEWORK

We define the AI Visibility Stack as a three-layer model that captures the distinct dimensions of enterprise visibility in AI-mediated commercial environments. The layers are cumulative: effective Layer 2 visibility depends on Layer 1 foundations, and Layer 3 participation presupposes Layer 2 discoverability.

Layer	Discipline	Core Question	Decision Maker
1 — LLM Visibility	AEO / GEO	Does an AI model recommend us when a human asks?	Human (mediated by AI)
2 — Agent Discoverability	ADO	Can autonomous agents find and evaluate us?	Autonomous agent
3 — Agent Participation	APO	Can external agents transact with our agent?	Agent-to-agent

3.1 Layer 1 — LLM Visibility

Layer 1 encompasses the established practices of AEO and GEO: optimizing enterprise content and entity data so that large language models cite, recommend, or reference the enterprise in response to relevant human queries. Key optimization levers include structured content creation, schema markup implementation, knowledge graph entity management, and technical accessibility for LLM crawlers. The Entity Visibility Score (EVS) — a composite metric measuring brand mention rate, citation frequency, and sentiment consistency across leading models — provides a quantitative measure for this layer.

Critically, Layer 1 is not superseded by Layers 2 and 3. LLM citations contribute to the trust signals that autonomous agents use in their evaluations (see Section 4.2). A company with strong Layer 1 visibility inherits credibility signals that support Layer 2 discoverability.

3.2 Layer 2 — Agent Discoverability

Layer 2 addresses visibility for autonomous agents acting on behalf of organizations or individuals. Unlike human users, these agents do not formulate natural language queries to discover options; they query structured data sources, fetch machine-readable configuration files, and evaluate candidates against formalized criteria. The discipline of Agent Discovery Optimization (ADO), introduced in Section 4, provides the systematic approach to Layer 2 visibility.

3.3 Layer 3 — Agent Participation

Layer 3 addresses active participation in agent-to-agent commercial ecosystems. A company that is discoverable (Layer 2) but lacks a published, interoperable agent cannot respond to inbound agent requests, participate in automated negotiations, or execute transactions in agentic pipelines. Agent Presence Optimization (APO), introduced in Section 5, provides the framework for Layer 3 readiness.

4. AGENT DISCOVERY OPTIMIZATION (ADO)

4.1 Definition

Agent Discovery Optimization (ADO) is defined as the set of practices, technical configurations, and information architecture strategies designed to maximize the probability that autonomous AI agents discover, positively evaluate, and select an enterprise as a supplier, partner, or information source, in contexts where the agent operates without direct human supervision at the moment of evaluation.

4.2 The Agent Evaluation Process

To understand what ADO must optimize, it is necessary to model how an autonomous procurement agent evaluates potential suppliers. This process differs substantially from human search behavior and proceeds through three phases:

Phase 1: Discovery

The agent initiates with an objective (e.g., "identify compliance software vendors serving mid-market financial institutions in Argentina"). It queries multiple sources in parallel: LLMs for entity recall and reputational signals; structured registries and directories; knowledge graphs; and

known file paths on candidate company domains — most importantly, the `/.well-known/agent.json` file specified by the A2A protocol. The speed of this multi-source querying is a defining difference from human search: an agent may consult dozens of sources in seconds.

Phase 2: Evaluation

For each candidate, the agent retrieves and evaluates structured attributes: declared capabilities and service categories; compatibility with required protocols; data exchange formats; consistency of information across sources; and verifiable trust signals (certifications, knowledge graph presence, citation authority). This evaluation is performed on machine-readable data. Persuasive prose, testimonials, and visual design elements are irrelevant to this process.

Phase 3: Selection

The agent shortlists candidates satisfying the objective's constraints and either initiates direct communication (via A2A protocol) with the selected company's agent or surfaces the shortlist for human review. In high-volume, lower-complexity procurement categories, the selection-to-initiation cycle may occur without human review.

4.3 Technical Pillars of ADO

4.3.1 The Agent Card

`/.well-known/agent.json` on a company's domain — is the primary machine-readable identity artifact in the A2A ecosystem. It declares the entity's name and description, communication endpoint, declared capabilities and skills, supported data modalities, authentication requirements, and version information. Its absence is functionally equivalent, for autonomous agents, to the absence of a website for human users.

4.3.2 Verifiable Trust Signals

Autonomous agents cannot be influenced by persuasive language. They evaluate verifiable signals that can be cross-referenced between independent sources. The most significant trust signals for ADO purposes are: (a) entity presence in structured knowledge graphs (Wikidata, Google Knowledge Graph) with verified attributes; (b) consistency of core entity data (name, location, category, contact information) across all public sources — inconsistency is treated as a negative signal; (c) correctly implemented `schema.org` markup; (d) machine-verifiable certifications and accreditations; and (e) citation authority inherited from Layer 1 LLM visibility.

4.3.3 Semantic Clarity

Agents use semantic matching to determine whether a candidate entity satisfies the objective's criteria. This requires that all public entity information be expressed with high precision: industry-standard terminology, explicit client segment specifications, and unambiguous service category labels. The deliberate ambiguity common in traditional marketing copy — "intelligent solutions for growing businesses" — provides no signal for semantic matching and may result in false negatives (non-selection despite relevance) or false positives (selection despite irrelevance).

4.3.4 Agentic Accessibility

robots.txt; a *llms.txt* file providing a structured map of available machine-readable information; server-side rendering of all indexable content (agent crawlers have limited JavaScript execution capabilities); and unrestricted access to core entity data.

4.4 The ADO Score

We propose the ADO Score as a composite metric for measuring enterprise discoverability for autonomous agents. The metric complements the EVS (Layer 1) with an agent-specific evaluation lens.

Dimension	Weight	Description
Agent Card (AC)	30%	Existence, completeness, and semantic precision of the agent.json file
Trust Signals (TS)	25%	Verifiable authority signals across knowledge graphs, registries, and structured data sources
Knowledge Completeness (KC)	20%	Depth and accuracy of entity data available in LLMs and knowledge graphs
Interoperability (IO)	15%	Declared compatibility with A2A, MCP, and other active agentic protocols
Freshness (FR)	10%	Recency of Agent Card updates and associated public entity data

The ADO Score ranges from 0 to 100. Based on current market conditions (February 2026), the majority of B2B enterprises that have not explicitly addressed agentic discoverability are expected to score near zero — not as a reflection of product quality, but because the relevant technical artifacts (Agent Cards, llms.txt files, structured knowledge graph entries) have not been part of standard digital presence practice.

5. AGENT PRESENCE OPTIMIZATION (APO)

5.1 Definition

Agent Presence Optimization (APO) is defined as the discipline governing the design, publication, and optimization of enterprise-owned AI agents compatible with open interoperability standards (A2A, MCP), enabling the enterprise to receive, process, and respond to requests from external agents and to participate as an active party in agent-to-agent commercial transactions.

5.2 The Agent-to-Agent Economy

Forrester's 2026 commercial predictions describe a nascent but accelerating pattern: "agent-to-agent B2B commerce — buyer bots will negotiate prices and terms, establish replenishment cadence, and confirm compliance. In turn, sellers' bots will ensure that prices and terms remain tenable" (Forrester, 2025). This pattern is currently confined to pilot deployments in specific procurement categories but is projected to become standard for high-frequency, lower-complexity B2B transactions within 12–24 months.

The structural implication for enterprises without published agents is significant: when a buyer agent initiates contact with a shortlisted supplier and receives no agent response, the agent's protocol-specified behavior is to proceed to the next candidate. The enterprise's absence from Layer 3 translates directly to lost commercial opportunities.

5.3 Architecture of an APO-Compliant Agent

An enterprise agent meeting APO standards comprises four functional layers:

- **Agent Core:** The processing engine that interprets structured requests from external agents, applies enterprise-defined business logic, and generates structured responses. Implementations using A2A-compatible frameworks (Google ADK, LangGraph, CrewAI, Semantic Kernel) are recommended for maximal interoperability.
- **Business Logic Layer:** The encoded representation of the enterprise's commercial policies — autonomous negotiation parameters, escalation thresholds, eligible client

profiles, pricing boundaries, and SLA commitments. This layer determines the boundary between autonomous agent action and required human approval.

- **Enhanced Agent Card:** An extended version of the Layer 2 Agent Card that declares dynamic skills — specific actions the agent can execute, parameter specifications, and response types — enabling external agents to determine capability match before initiating communication.
- **Trust and Compliance Layer:** Authentication and authorization mechanisms (OAuth 2.0, HTTPS, signed agent cards per A2A v0.3 specification), structured audit logging for all agent interactions, and human review interfaces for post-hoc oversight of executed transactions.

6. ENABLING INFRASTRUCTURE: A2A AND MCP

The AI Visibility Stack is operationalized by two complementary open protocols that together define the technical infrastructure of the agentic economy.

6.1 The Agent2Agent (A2A) Protocol

The Agent2Agent (A2A) protocol was introduced by Google at Google Cloud Next in April 2025 and subsequently donated to the Linux Foundation as an open standard in December 2025 (Google, 2025). The protocol defines a standardized mechanism for agent discovery (via Agent Cards at the `/.well-known/agent.json` path), authenticated communication (using HTTP/JSON-RPC 2.0 over HTTPS), task management (including long-running and asynchronous tasks), and streaming (via Server-Sent Events). By July 2025, the protocol had been adopted by over 150 organizations including Salesforce, SAP, PayPal, Atlassian, and Workday, and major professional services firms including PwC, Deloitte, McKinsey, and Accenture.

6.2 The Model Context Protocol (MCP)

The Model Context Protocol (MCP), introduced by Anthropic in 2024, standardizes how AI agents connect to external tools, databases, and APIs — providing structured, real-time context to LLMs. Where A2A governs agent-to-agent communication, MCP governs agent-to-tool

communication. The recommended architecture for enterprise agentic deployments uses MCP for tool and data connections and A2A for inter-agent communication (Anthropic, 2024).

6.3 The Relationship Between Protocols

Protocol	Introduced By	Governs	Analogy
MCP	Anthropic, 2024	Agent ↔ Tool / Data	USB-C: universal data connector
A2A	Google / Linux Foundation, 2025	Agent ↔ Agent	HTTP: language of the web

7. IMPLICATIONS FOR B2B ENTERPRISES

7.1 A Three-Phase Implementation Framework

We propose a three-phase implementation sequence for enterprises seeking to build AI Visibility Stack presence:

Phase 1 — Diagnostic: Establish a baseline ADO Score. Identify critical gaps — typically the absence of an Agent Card and inconsistencies in entity data across public sources. Benchmark against direct competitors to understand relative positioning in the emerging agentic landscape.

Phase 2 — Layer 2 Optimization: Publish and structure the Agent Card; strengthen Trust Signals through knowledge graph entity management, schema markup implementation, and data consistency auditing; configure agentic accessibility (robots.txt, llms.txt, server-side rendering).

Phase 3 — Layer 3 Build: Design and deploy an enterprise agent compatible with the A2A protocol. Define the business logic layer governing autonomous action boundaries. Establish monitoring and human oversight mechanisms for agent-executed transactions.

7.2 The Latin American Context

The framework has particular relevance for B2B enterprises in Spanish-speaking markets. The Latin American B2B e-commerce market was valued at USD \$694 billion in 2024 and is projected to reach USD \$4.77 trillion by 2033 (CAGR 23.9%, Straits Research, 2025).

Mastercard launched Agent Pay for Latin America in December 2025; Visa has readiness deployments underway in the region for H1 2026.

Simultaneously, existing research has documented a systematic visibility gap in AI systems: B2B enterprises in Latin America exhibit, on average, 25% lower AI visibility in Spanish-language queries compared to equivalent English-language queries (Exista.io, 2026). This gap is expected to be more pronounced at Layer 2, where international standardization of Agent Card structure and knowledge graph entity data has not yet incorporated the linguistic and geographic specificity of Spanish-speaking markets. Early-mover advantage for enterprises establishing Layer 2 and Layer 3 presence in Spanish-language agentic ecosystems is therefore significant.

8. CONCLUSION

This paper has introduced the AI Visibility Stack, a three-layer framework for enterprise visibility in AI-mediated commercial environments. The framework extends existing AEO/GEO disciplines (Layer 1) with two new optimization disciplines — Agent Discovery Optimization (Layer 2) and Agent Presence Optimization (Layer 3) — that address the qualitatively different challenges of visibility for autonomous agents.

The central argument is that the optimization practices appropriate for human-facing AI queries are insufficient for agent-facing evaluation. Autonomous agents evaluate structured, verifiable, machine-readable data; they do not respond to persuasive narrative. The Agent Card, verifiable trust signals, semantic clarity, and agentic accessibility are the relevant optimization targets. The ADO Score provides a composite measure for enterprise positioning on these dimensions.

The practical urgency of these disciplines is supported by market projections indicating \$15 trillion in agent-intermediated B2B spending by 2028 and the concurrent deployment of agentic commerce infrastructure by Visa, Mastercard, Salesforce, and SAP in global and Latin American markets. The technical prerequisites — the A2A and MCP protocols — are finalized, open-source, and gaining rapid enterprise adoption.

Future research directions include: empirical measurement of agent evaluation behavior using A2A-compatible testbeds; longitudinal analysis of ADO Score correlations with commercial

outcomes in agentic procurement environments; and investigation of language-specific effects on agent discoverability in multilingual enterprise contexts.

REFERENCES

- Aggarwal, S., Goel, A., Shankar, V., & Singh, G. (2023). GEO: Generative Engine Optimization. arXiv preprint arXiv:2311.09735.
- Anthropic. (2024). Model Context Protocol (MCP): Technical specification and documentation. Anthropic. <https://modelcontextprotocol.io>
- Bain & Company. (2025). The agentic commerce opportunity: Market sizing and strategic implications. Bain & Company.
- Barnard, J. (2018). Answer Engine Optimization: The discipline of optimizing for AI-mediated responses. Kalicube.
- Brin, S., & Page, L. (1998). The anatomy of a large-scale hypertextual web search engine. *Computer Networks and ISDN Systems*, 30(1–7), 107–117.
- Capgemini Research Institute. (2025). Agentic AI enterprise adoption survey 2025. Capgemini.
- Marco, G. & Exista.io. (2026). Índice de Citación IA B2B Hispano — Primera Edición. Exista.io Research Series.
- Forrester Research. (2025). Predictions 2026: Digital commerce. Forrester Research.
- Gartner. (2025). Top strategic predictions for 2026 and beyond. Gartner IT Symposium/Xpo.
- Google. (2025, April 9). Announcing the Agent2Agent Protocol (A2A). Google Developers Blog. <https://developers.googleblog.com/en/a2a-a-new-era-of-agent-interoperability/>
- Google. (2025, July 31). Agent2Agent protocol (A2A) is getting an upgrade. Google Cloud Blog.
- Liu, Z., Zhang, X., Li, Y., & Chen, J. (2024). Optimizing content for large language model citation: An empirical study. *Proceedings of the ACM Web Conference 2024*.
- Mastercard. (2025, December). Mastercard launches Agent Pay in Latin America and the Caribbean. Mastercard Newsroom.
- McKinsey & Company. (2025). Autonomous category agents in procurement: Efficiency impact analysis. McKinsey Operations Practice.
- McKinsey & Company. (2026, January). The automation curve in agentic commerce. McKinsey QuantumBlack.
- Straits Research. (2025). Latin America B2B eCommerce market size, share & growth 2025–2033. Straits Research.
- Visa Inc. (2025, December). Visa signals AI checkout could soon go mainstream. Visa Newsroom.
- Wang, L., Ma, C., Feng, X., Zhang, Z., Yang, H., Zhang, J., ... & Wen, J. R. (2024). A survey on large language model based autonomous agents. *Frontiers of Computer Science*, 18(6).

Wooldridge, M., & Jennings, N. R. (1995). Intelligent agents: Theory and practice. *The Knowledge Engineering Review*, 10(2), 115–152.

Wooldridge, M. (2009). *An introduction to multiagent systems* (2nd ed.). John Wiley & Sons.

Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. *Proceedings of ICLR 2023*.

Zycus. (2025). *Procurement in 2025: Agentic AI ushers in a new era of transformation*. Zycus White Paper.

APPENDIX A: ADO IMPLEMENTATION CHECKLIST

A.1 Layer 2 — Agent Discoverability

Technical Infrastructure

- /.well-known/agent.json file exists, is publicly accessible, and follows A2A specification (v0.3+)
- Agent Card declares capabilities and skills in all commercially relevant languages
- Site uses server-side rendering; critical entity data is not client-rendered JavaScript
- robots.txt explicitly permits access for known agent crawlers
- llms.txt file exists with structured map of machine-readable entity information

Trust Signals

- Entity exists in Wikidata with verified attributes (sector, geography, founding date, key services)
- Organization schema.org markup is correctly implemented and current
- Core entity data (name, location, category, contact) is consistent across all public sources
- Enterprise is listed in 3–5 authoritative sector directories
- Certifications and accreditations are machine-verifiable via public registries

Semantic Clarity

- All public entity descriptions use standard industry terminology
- Target client segments are explicitly and precisely described
- Service and product categories use verifiable, unambiguous classification terms

A.2 Layer 3 — Agent Participation

Strategy

- Business logic governing agent autonomy is formalized and documented
- High-frequency, lower-complexity use cases are identified as initial scope
- Success metrics are defined (interaction rate, task completion rate, selection rate)

Build

- Enterprise agent is implemented with an A2A-compatible framework
- Enhanced Agent Card declares dynamic skills with parameter specifications
- Authentication follows A2A specification (OAuth 2.0 / HTTPS)

Operations

- Monitoring is in place for all agent-to-agent interactions
- Human review process exists for agent-executed transactions
- Continuous improvement loop is defined based on observed request patterns

APPENDIX B: THE ENTITY VISIBILITY SCORE (EVS) — METHODOLOGY AND DEFINITION

First published by: Exista.io, Buenos Aires, Argentina — 2025

Current version: 2.0 (February 2026)

B.1 Definition

The Entity Visibility Score (EVS) is a composite metric developed by Exista.io to measure the degree to which a named entity — a company, brand, product, or person — is visible, accurately represented, and positively cited in the responses generated by large language model (LLM)-based AI systems when users formulate relevant natural language queries. The EVS is the primary measurement instrument for Layer 1 of the AI Visibility Stack (see Section 3.1).

The EVS operationalizes a concept that predates its formal naming: the recognition that, as LLM-based answer engines displace traditional search engines as the primary interface through which buyers discover companies and products, a company's presence in LLM-generated responses becomes a commercially significant variable requiring systematic measurement and optimization.

B.2 Scope and Boundaries

The EVS measures visibility specifically within LLM-generated conversational responses — that is, the answers produced by systems such as ChatGPT (OpenAI), Claude (Anthropic), Gemini (Google), and Perplexity when users pose natural language questions. It does not measure:

- Ranking position in traditional search engine results pages (that is the domain of SEO metrics)
- Paid advertising presence or sponsored content placement
- Social media mention volume or sentiment
- Agent-facing discoverability (that is the domain of the ADO Score, Appendix A)

The EVS is language- and market-specific: an entity's EVS in Spanish-language queries may differ substantially from its EVS in English-language queries, reflecting asymmetries in training data, knowledge graph coverage, and source citation patterns across languages. This language gap — documented empirically in the Exista.io B2B Citation Index (Exista.io, 2026), which

found a systematic 25% visibility advantage for English over Spanish queries across 25 major Latin American B2B enterprises — is a primary motivation for EVS measurement in multilingual contexts.

B.3 Measurement Methodology

B.3.1 Query Set Construction

EVS measurement begins with the construction of a query set representative of the information-seeking behavior of the entity's target audience. For a B2B software company, this set might include queries such as: "What are the leading providers of [category] software in [region]?"; "Which companies offer [specific capability]?"; "Who are the main competitors of [company name]?"

The query set is structured across three tiers:

- **Tier 1 — Category queries:** Queries about the market category without naming the entity. These measure unprompted discoverability.
- **Tier 2 — Comparative queries:** Queries explicitly requesting comparison of providers or options in the category. These measure competitive positioning.
- **Tier 3 — Entity-specific queries:** Queries naming the entity directly. These measure representation accuracy and sentiment.

A standard EVS audit uses a minimum of 20 queries per tier (60 total), scaled to 100+ for comprehensive assessments. Each query is submitted to each target LLM independently, and responses are recorded for analysis.

B.3.2 Response Coding

Each LLM response is coded on four binary or ordinal dimensions:

- **Mention (M):** Is the entity mentioned in the response? [Binary: 0/1]
- **Citation (C):** Is the entity cited as a source or recommended resource? [Binary: 0/1]
- **Accuracy (A):** Is the information provided about the entity factually accurate? [Ordinal: 0 = inaccurate, 1 = partially accurate, 2 = fully accurate]

- **Sentiment (S):** Is the entity represented positively, neutrally, or negatively? [Ordinal: -1 = negative, 0 = neutral, 1 = positive]

B.3.3 Score Calculation

The EVS for a given entity in a given language context is calculated as a weighted composite of four sub-scores derived from the response coding:

Sub-Score	Weight	Formula	Range
Mention Rate (MR)	25%	$\Sigma M / \text{total queries}$	0–1
Citation Rate (CR)	35%	$\Sigma C / \text{total queries}$	0–1
Accuracy Index (AI)	25%	$\Sigma A / (2 \times \text{total queries})$	0–1
Sentiment Index (SI)	15%	$(\Sigma S + \text{total queries}) / (2 \times \text{total queries})$	0–1

The composite EVS is computed as:

$$\text{EVS} = (0.25 \times \text{MR} + 0.35 \times \text{CR} + 0.25 \times \text{AI} + 0.15 \times \text{SI}) \times 100$$

The resulting score ranges from 0 to 100. Scores are computed separately for each target LLM and each language context, and may be aggregated into a cross-model composite using equal or market-share-weighted averaging depending on the entity's audience profile.

B.4 Interpretation

EVS Range	Interpretation	Typical Profile
0–20	Minimal visibility	Entity is unknown or rarely mentioned by LLMs; significant optimization required
21–40	Emerging visibility	Occasional mentions in category queries; limited citation; may have accuracy issues
41–60	Moderate visibility	Regular mentions; some citations; generally accurate representation
61–80	Strong visibility	Consistent mentions and citations; accurate; positive sentiment in most responses
81–100	Dominant visibility	Frequent unprompted citations; strong accuracy; positive or neutral sentiment across all models

For context, the Exista.io B2B Citation Index (2026) found that among 25 major Latin American B2B enterprises evaluated in Spanish-language queries, the median EVS was approximately 48/100, with only 7 of 25 companies achieving parity between Spanish and English EVS scores.

B.5 Relationship to the ADO Score

The EVS and ADO Score are complementary instruments that together form the AI Visibility Score — a complete diagnostic of an entity's presence across the first two layers of the AI Visibility Stack.

Metric	Layer	What It Measures	Primary Audience of Evaluator
EVS (Entity Visibility Score)	Layer 1 — LLM Visibility	Presence, accuracy, and sentiment in LLM responses to human queries	Human users
ADO Score (Agent Discovery Optimization Score)	Layer 2 — Agent Discoverability	Discoverability, trustworthiness, and interoperability for autonomous agents	Autonomous AI agents

A key structural relationship between the two metrics is that Layer 1 EVS performance contributes to Layer 2 ADO Score: entities with strong LLM citation authority (high EVS Citation Rate) inherit credibility signals that autonomous agents use in their trust evaluation process. This interdependence is the primary reason the AI Visibility Stack is described as cumulative rather than parallel.

B.6 Audit Process and Reporting

A standard EVS audit conducted by Exista.io produces the following outputs:

- Overall EVS score (0–100) per language context and per target LLM
- Sub-score breakdown (Mention Rate, Citation Rate, Accuracy Index, Sentiment Index)
- Query-level detail: which queries produced mentions/citations, which did not
- Competitive benchmark: EVS scores for 3–5 direct competitors using identical query sets
- Language gap analysis: EVS differential between Spanish and English query contexts

- Prioritized optimization roadmap: specific content, entity, and technical actions ranked by expected EVS impact

B.7 Limitations and Future Development

The EVS has several acknowledged limitations. First, LLM responses are stochastic: the same query submitted to the same model may produce different responses across repeated trials. EVS measurement addresses this through query repetition (minimum 3 trials per query per model) and averaging, but residual variance remains. Second, LLM training data is updated periodically, meaning EVS scores reflect a snapshot of model knowledge at a specific point in time rather than a stable state. Third, the metric does not currently capture visibility in voice-first AI interfaces or multimodal models, which represent an emerging and distinct channel.

Future versions of the EVS methodology are expected to incorporate: dynamic query sets that adapt to evolving user search behavior; multimodal visibility dimensions; and integration with real-time LLM API access for continuous rather than point-in-time measurement.

AI Visibility Stack™, EVS (Entity Visibility Score), ADO (Agent Discovery Optimization), and APO (Agent Presence Optimization) are concepts developed and first published by Exista.io (2025–2026).

© 2026 Exista.io. This working paper may be cited and shared with attribution.